

A High Throughput Paxos Variant That Mere Mortals Can Understand

Neil Giridharan, Joseph M. Hellerstein, Ion Stoica,
Adriana Szekeres, Michael Whittaker


C1

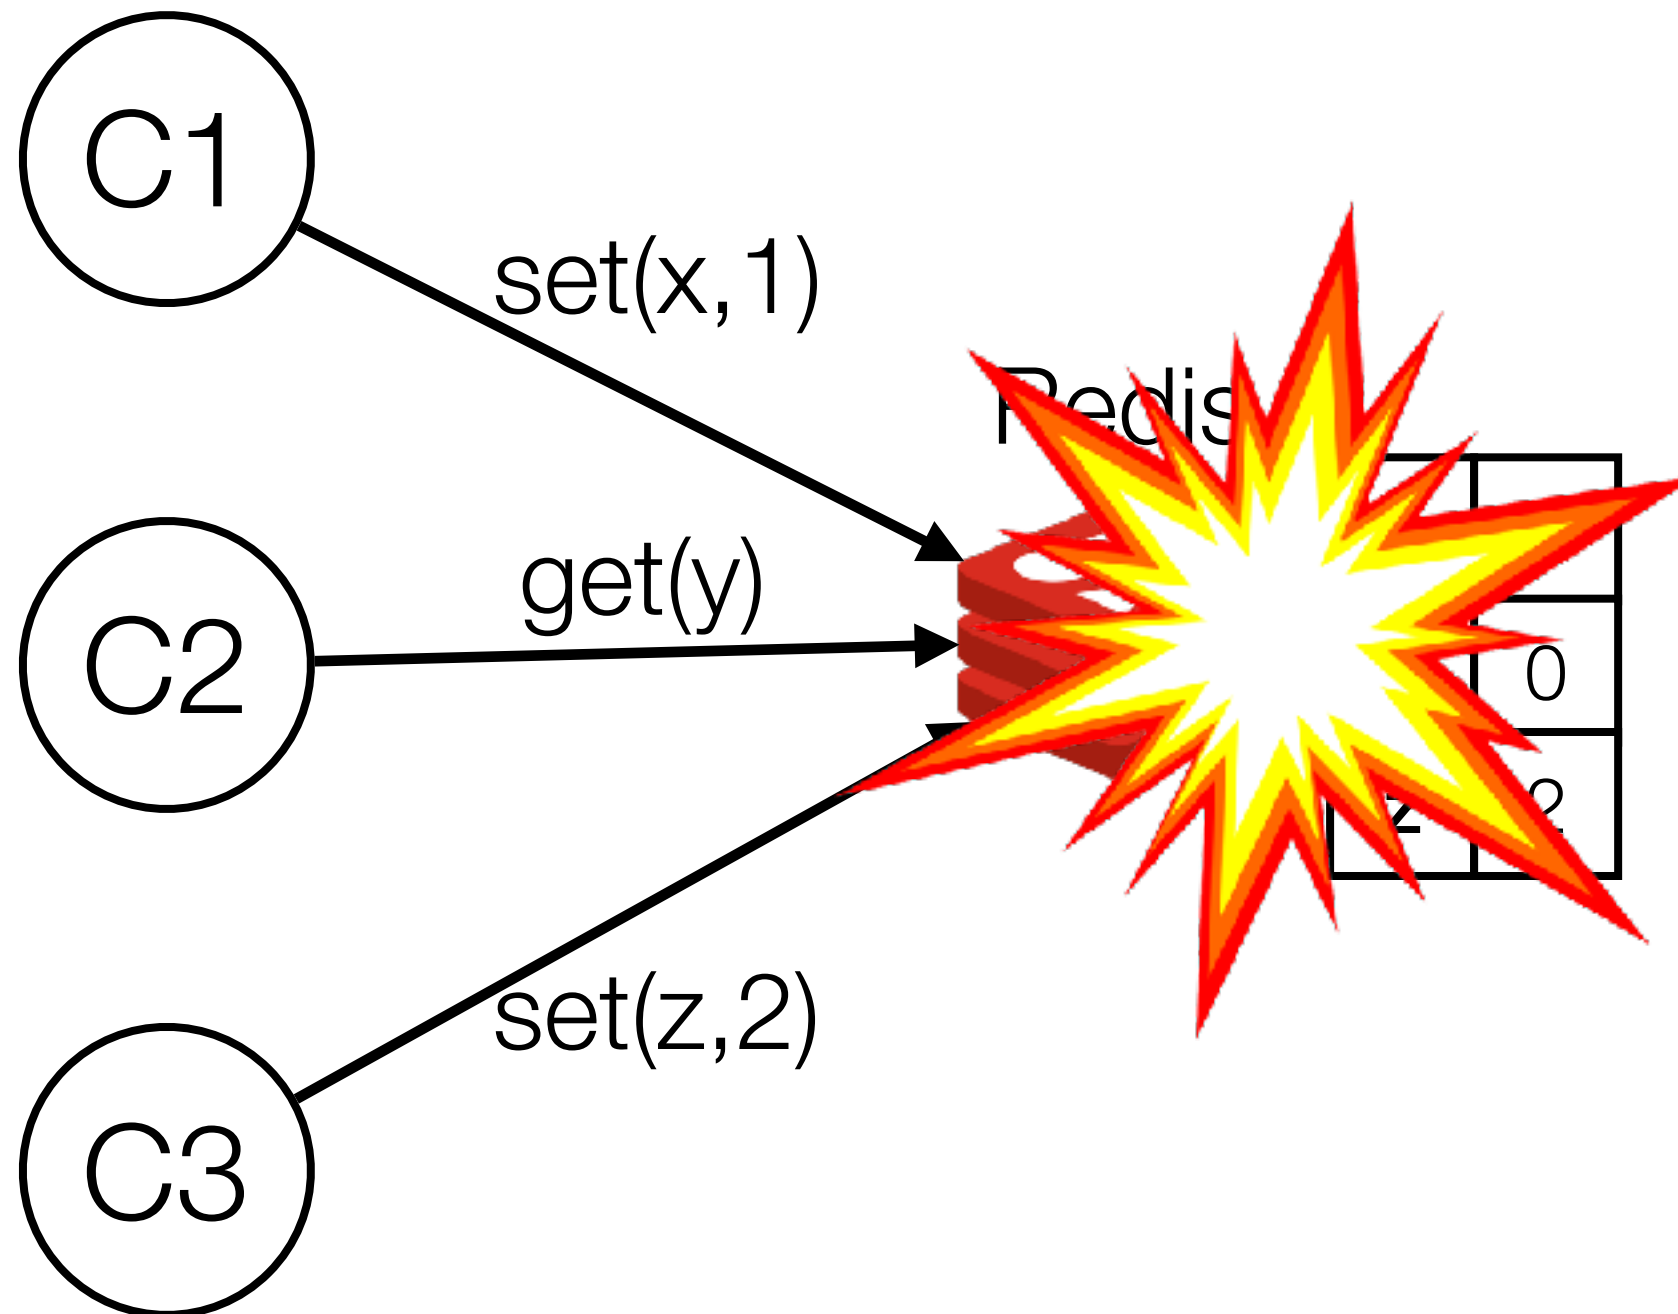
C2

C3

Redis







C1

C2

C3

Redis



x	0
y	0
z	0



x	0
y	0
z	0



x	0
y	0
z	0

C1

C2

C3


Redis




x	0
y	0
z	0







x	0
y	0
z	0



x	0
y	0
z	0









State
Machine
Replication















MultiPaxos

clients



clients








$f+1$
proposers



clients

$f+1$
proposers

$2f+1$
acceptors













clients

$f+1$
proposers

$2f+1$
acceptors

$f+1$
replicas











clients

$f+1$
proposers

$2f+1$
acceptors

$f+1$
replicas

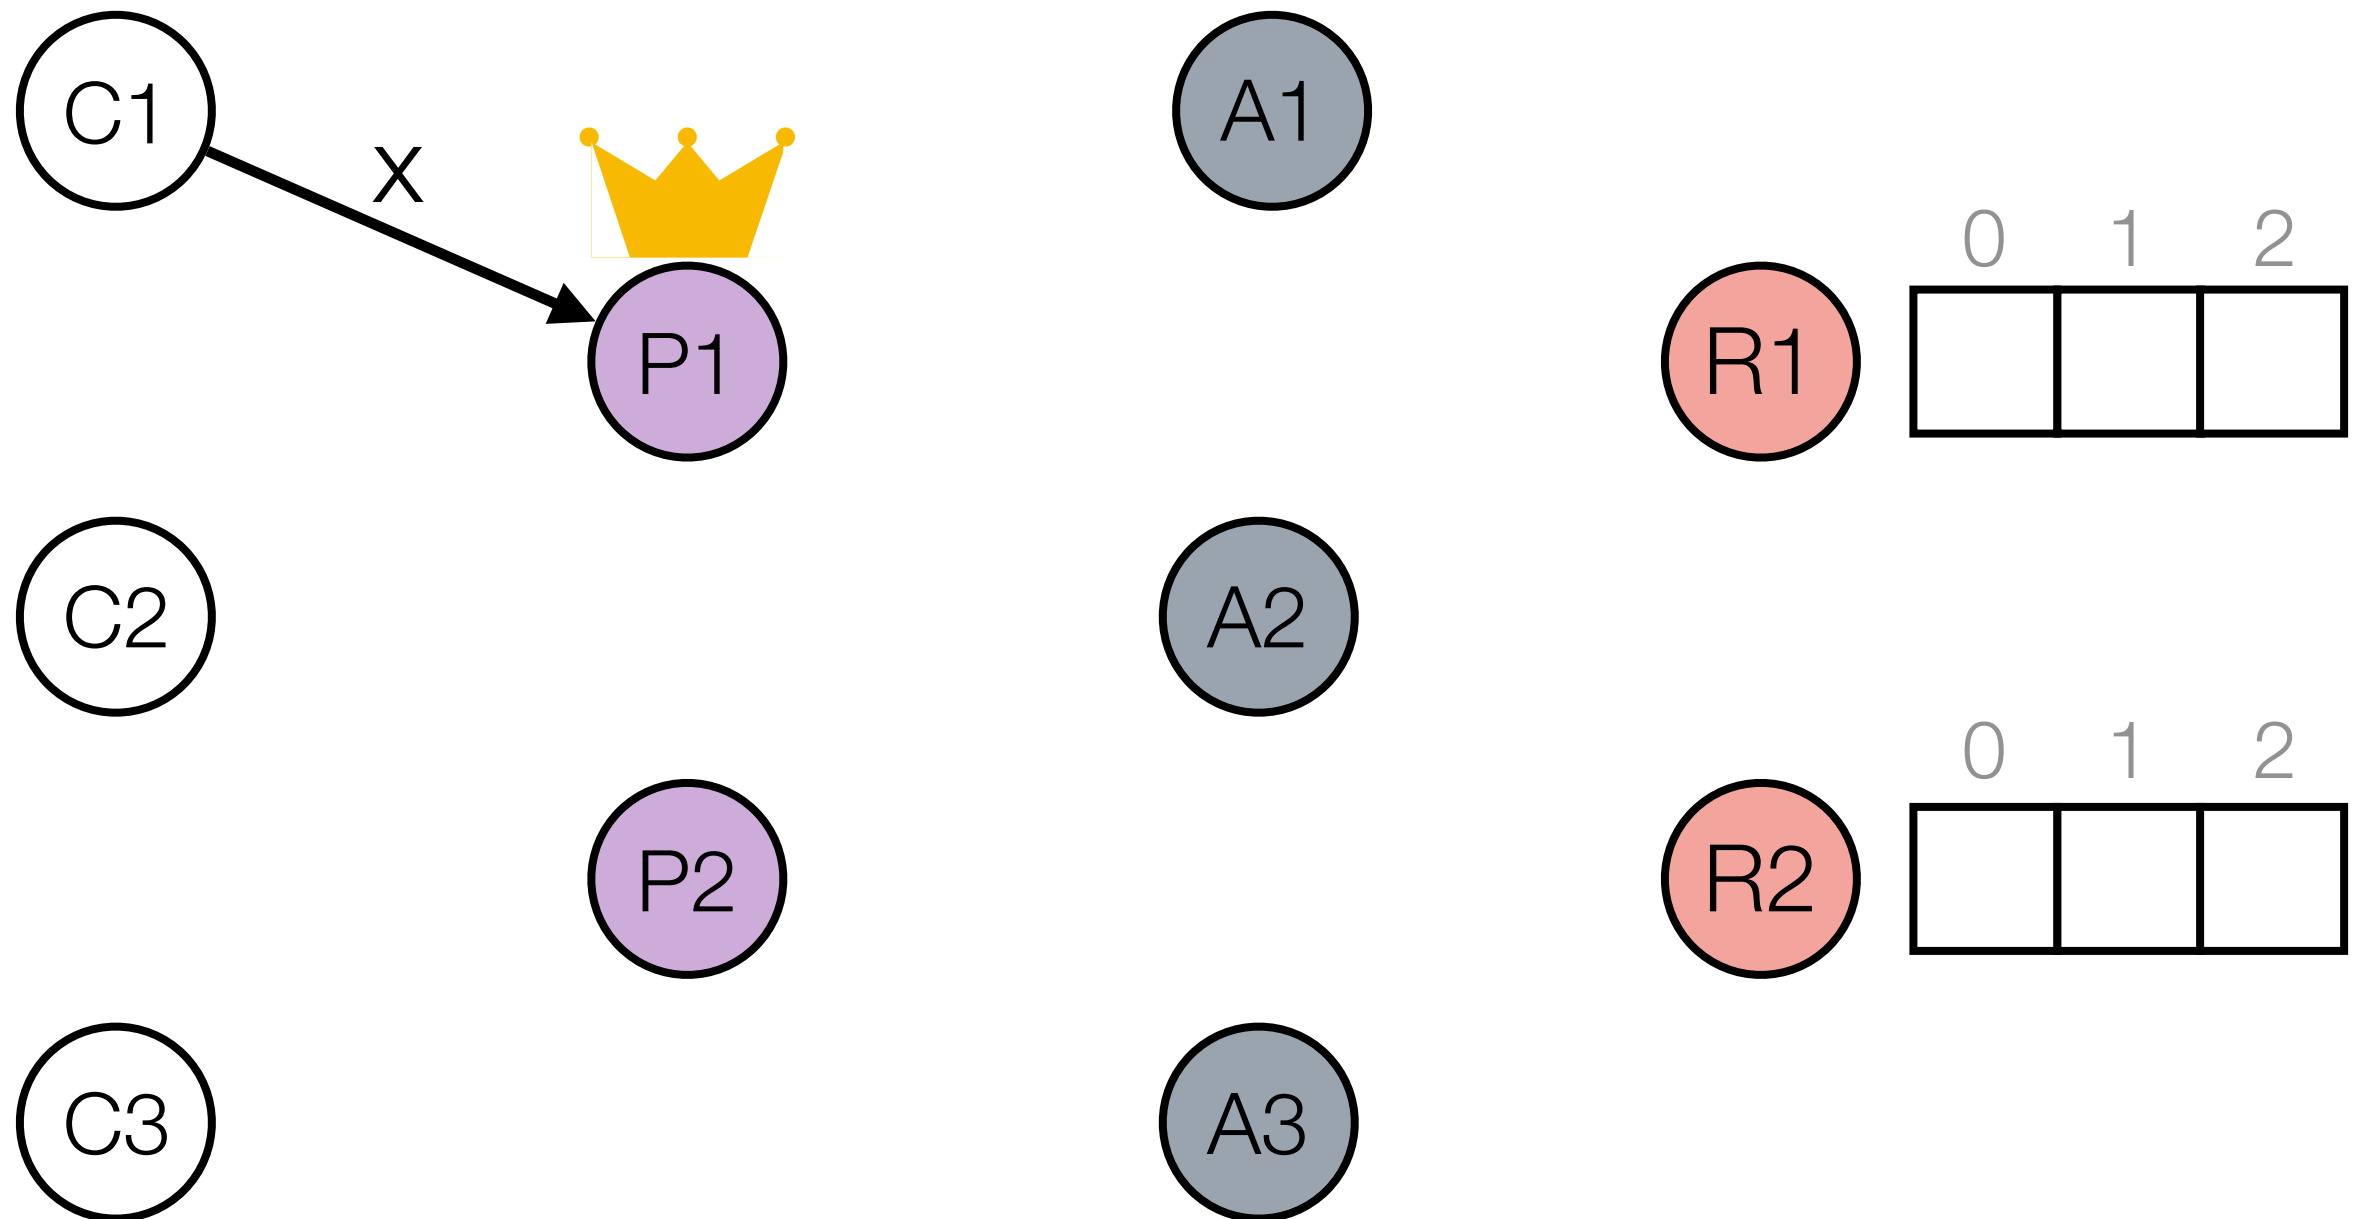


clients

$f+1$
proposers

$2f+1$
acceptors

$f+1$
replicas




clients

$f+1$
proposers

$2f+1$
acceptors

$f+1$
replicas

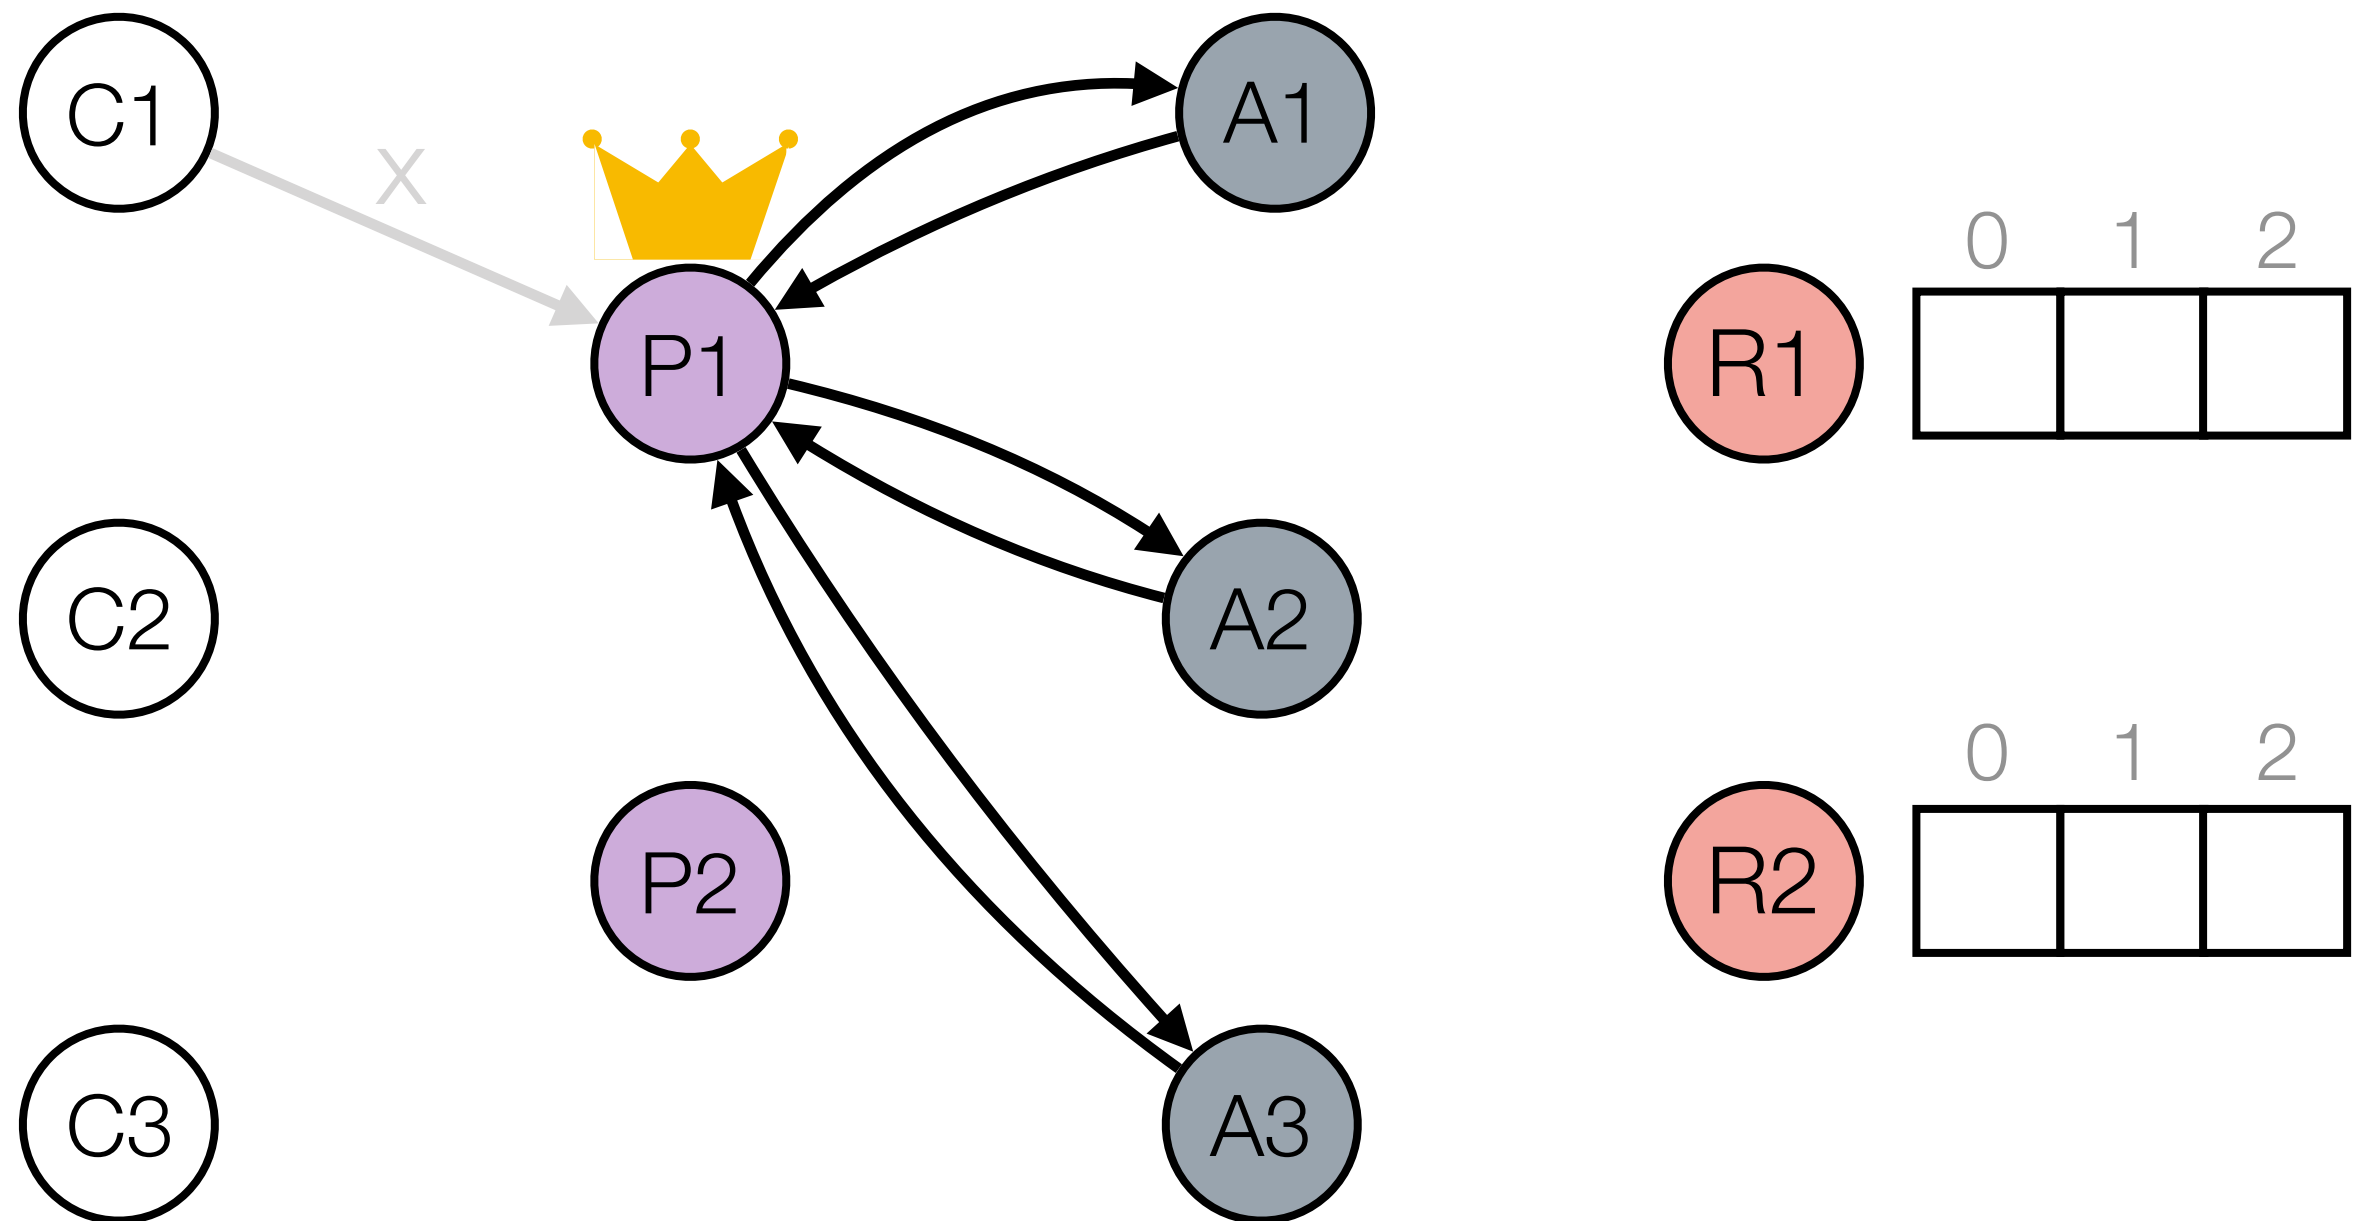


clients

$f+1$
proposers

$2f+1$
acceptors

$f+1$
replicas

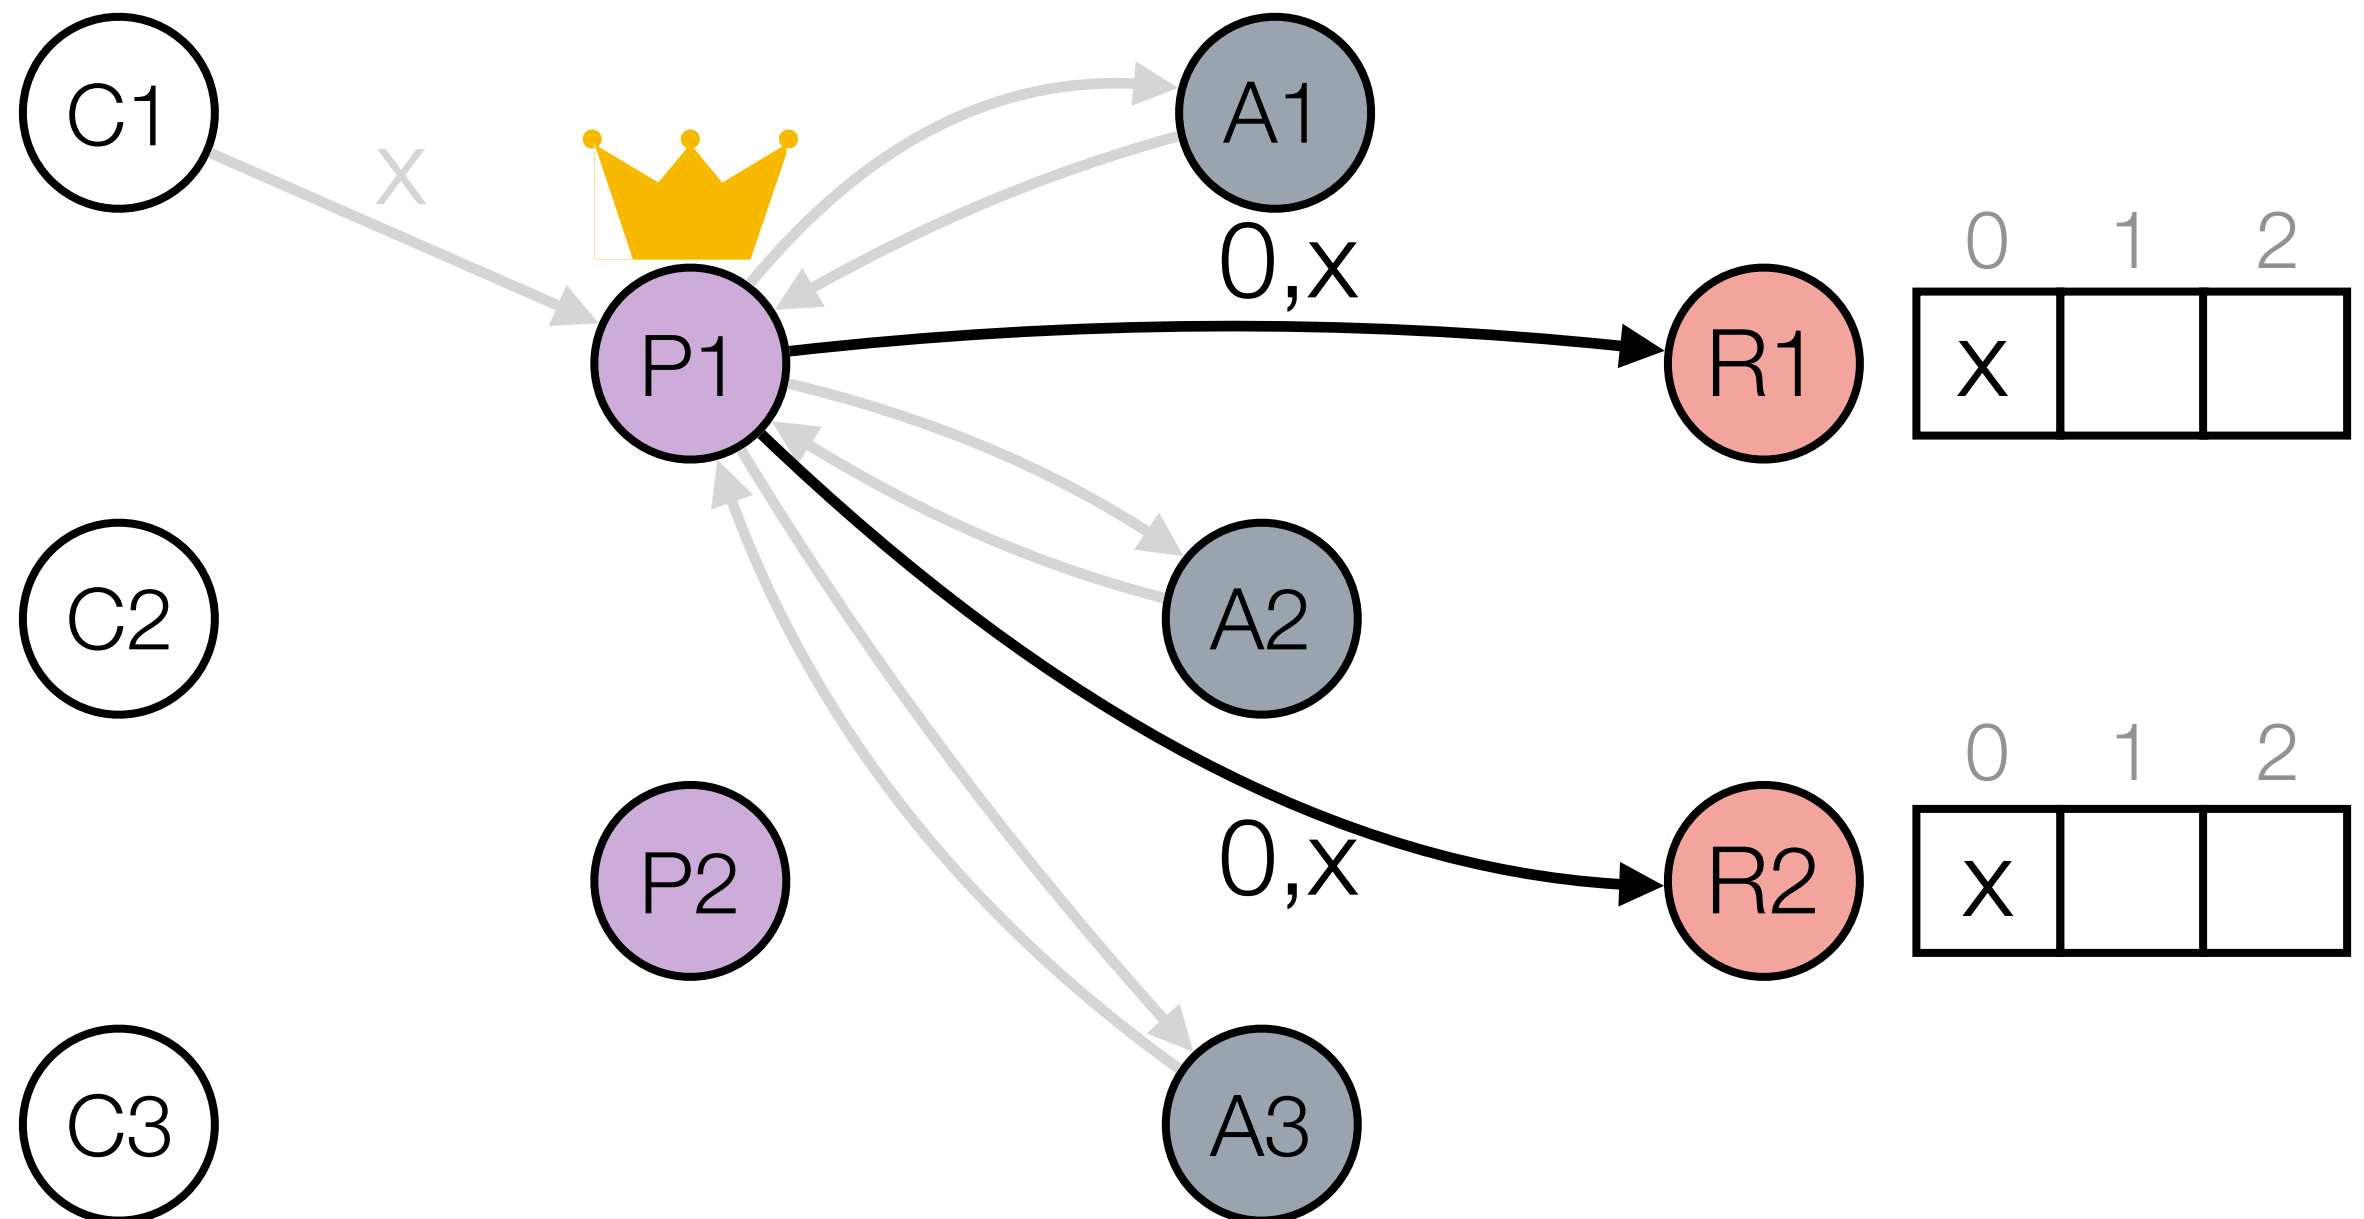


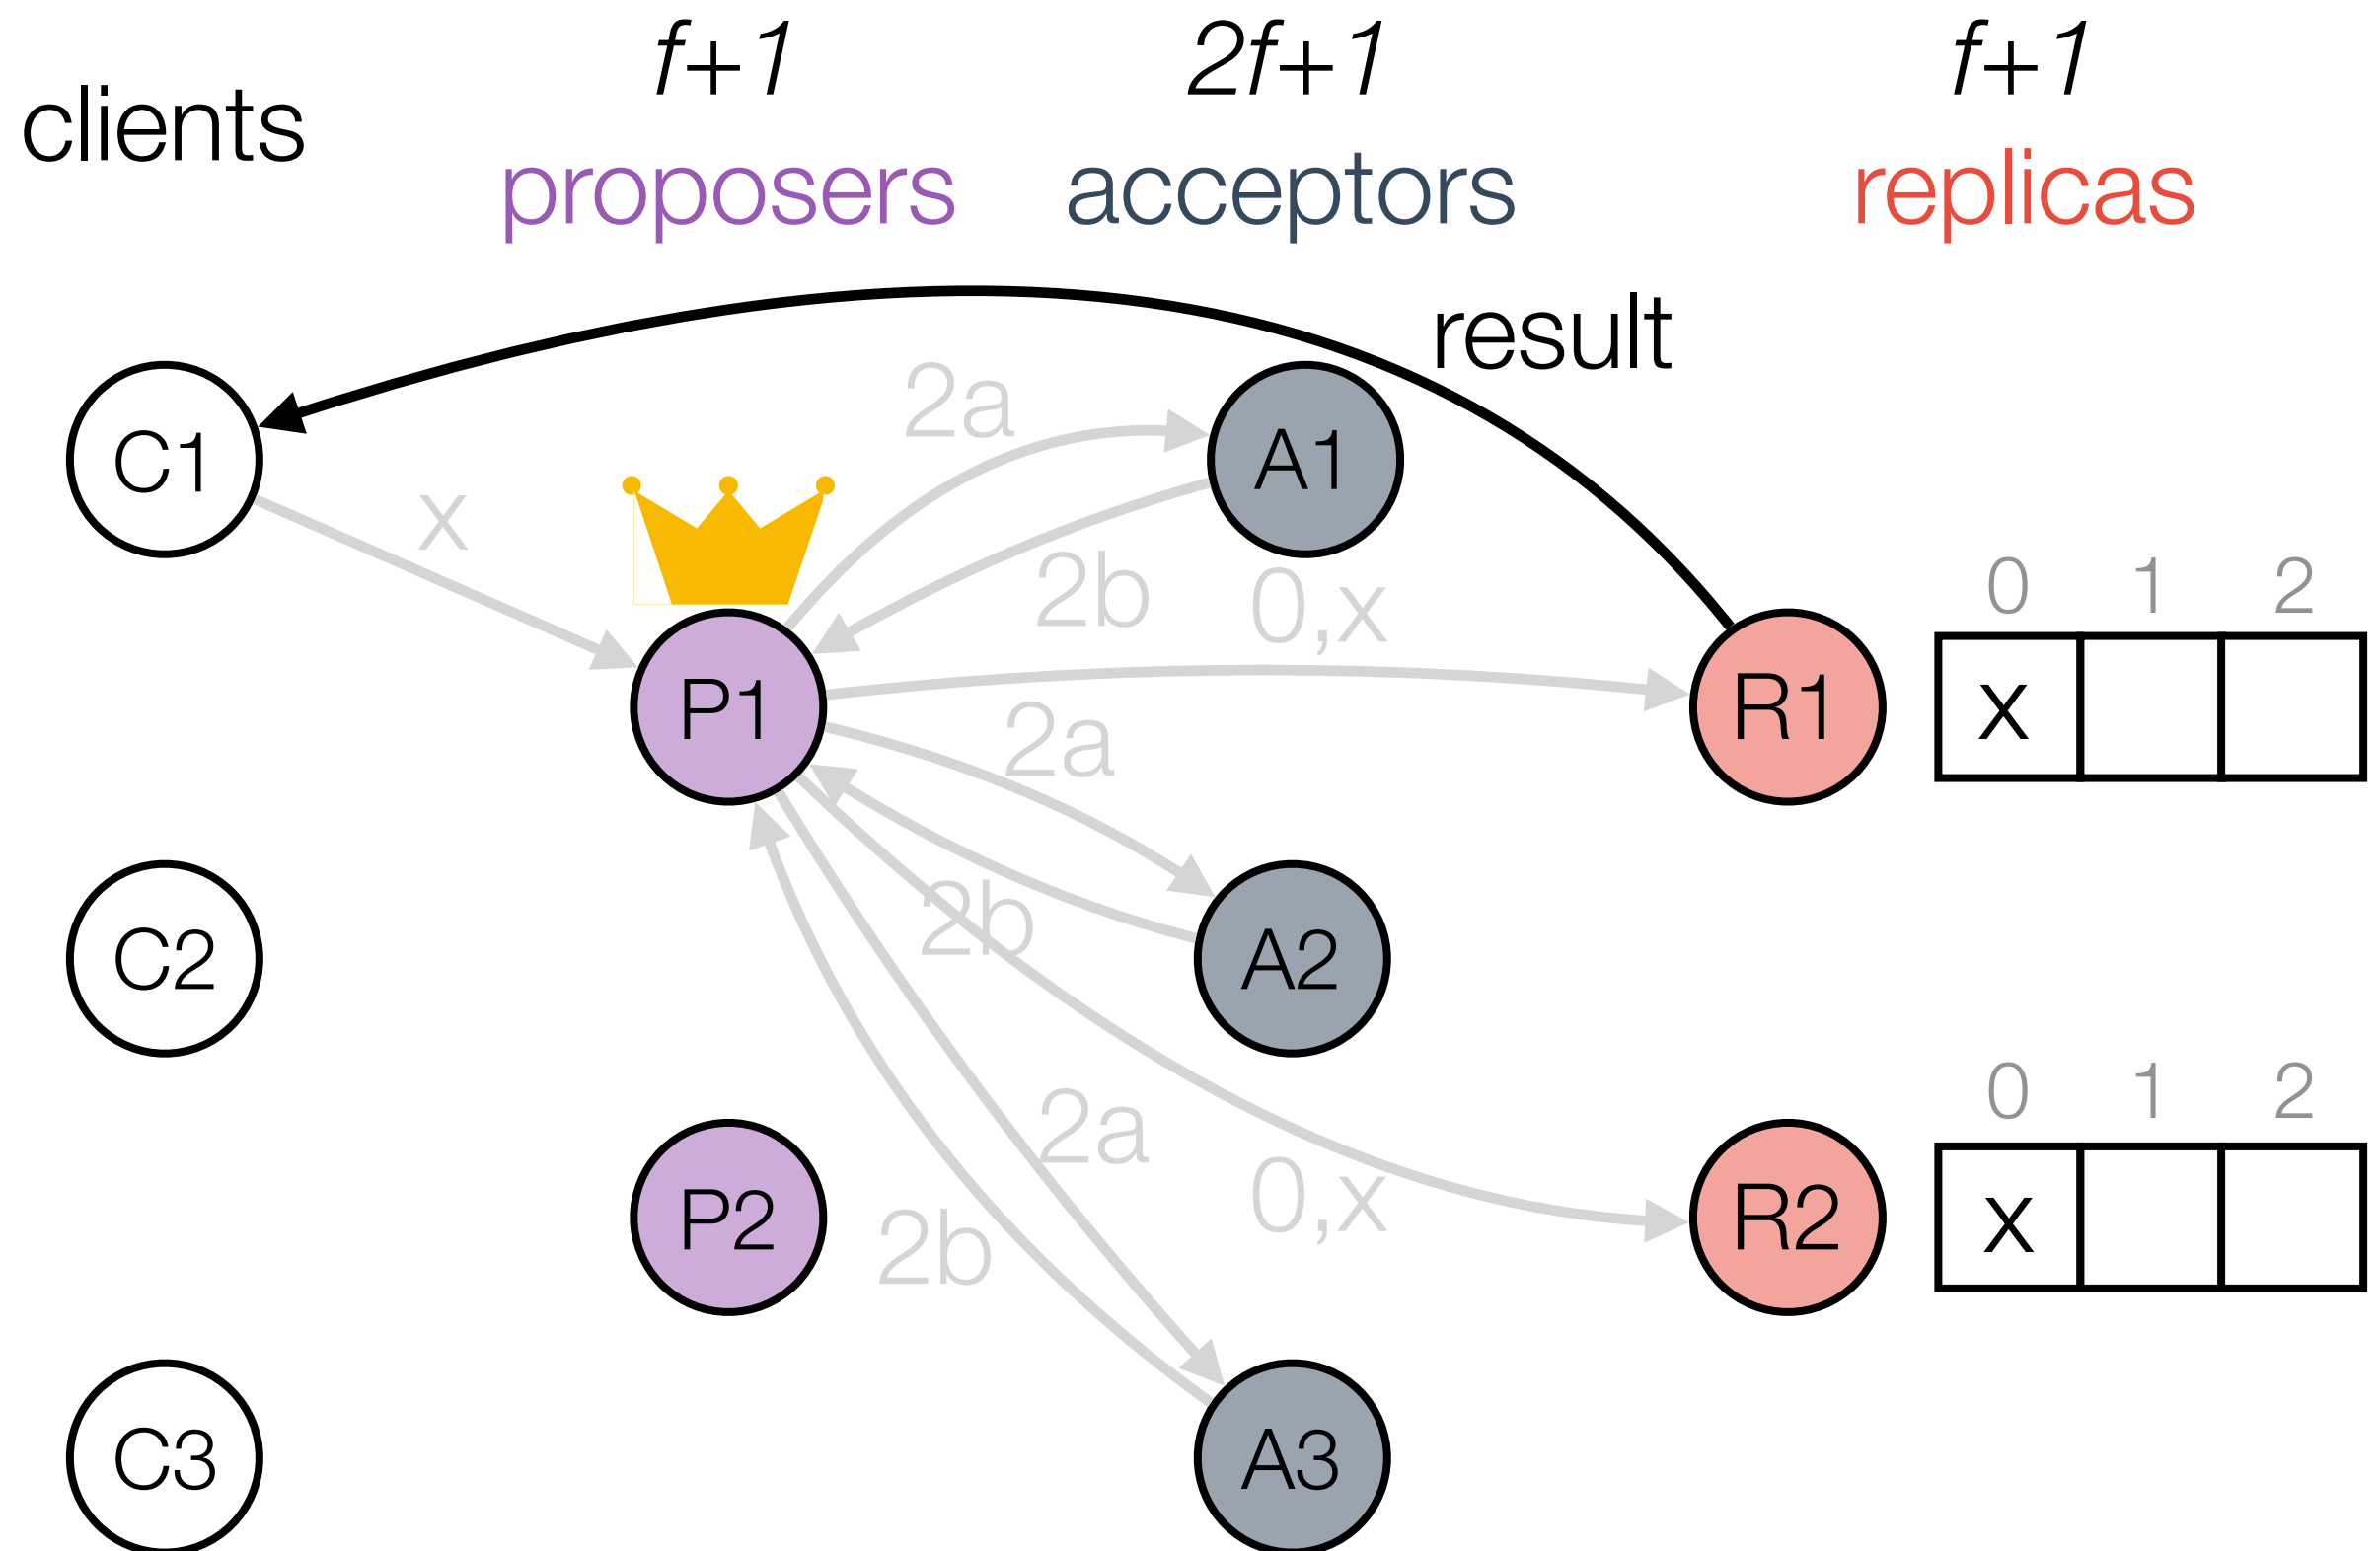
clients

$f+1$
proposers

$2f+1$
acceptors

$f+1$
replicas
















clients

$f+1$
proposers

$2f+1$
acceptors

$f+1$
replicas

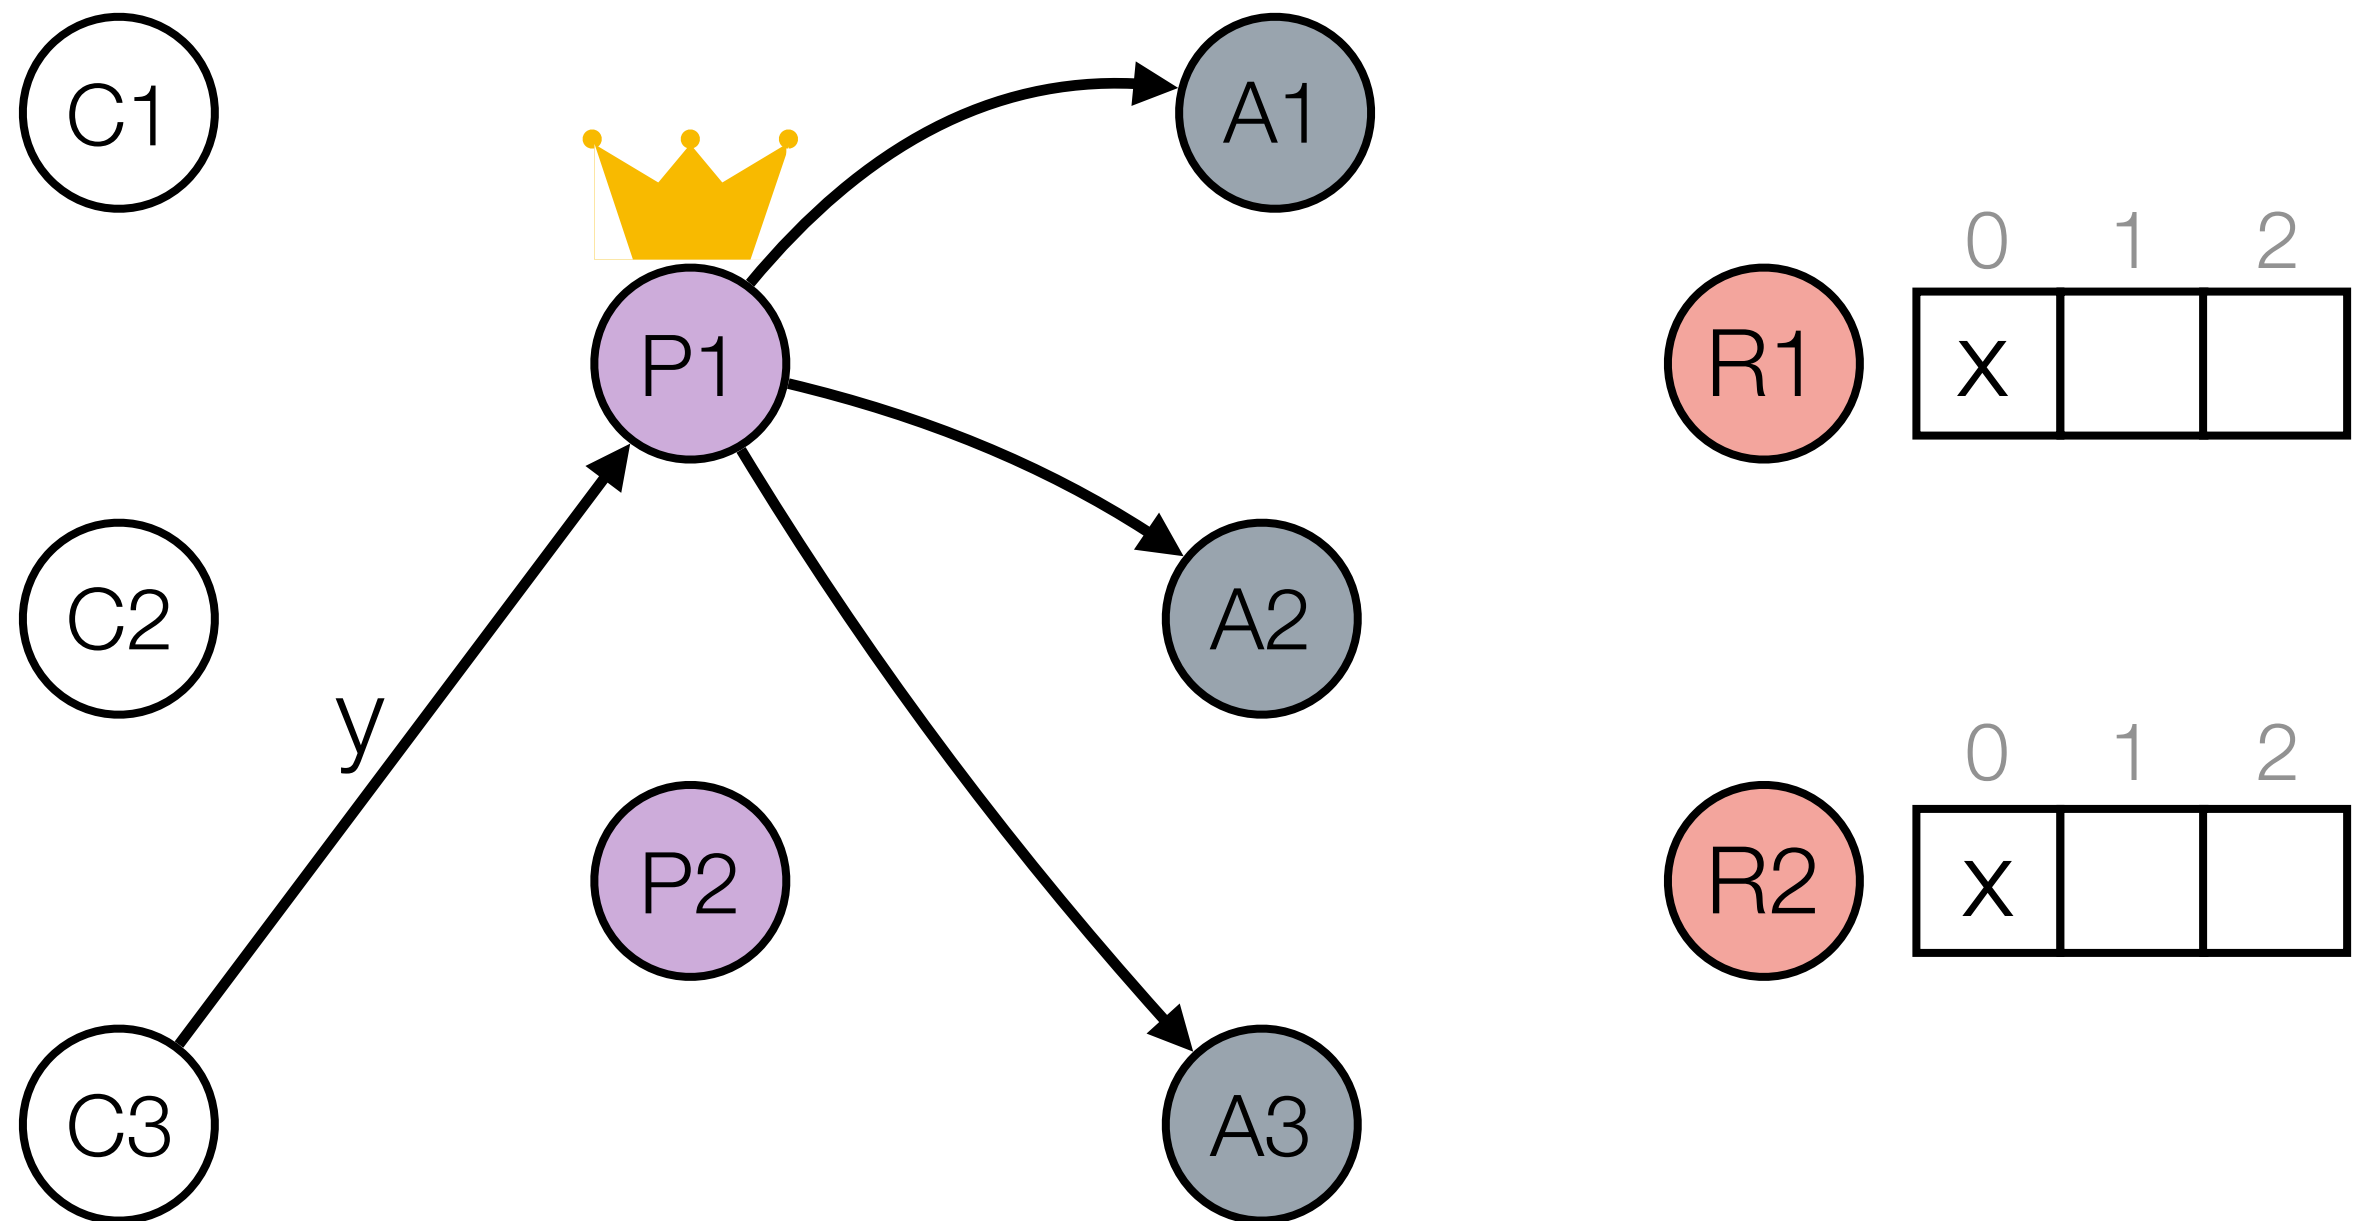


clients

$f+1$
proposers

$2f+1$
acceptors

$f+1$
replicas


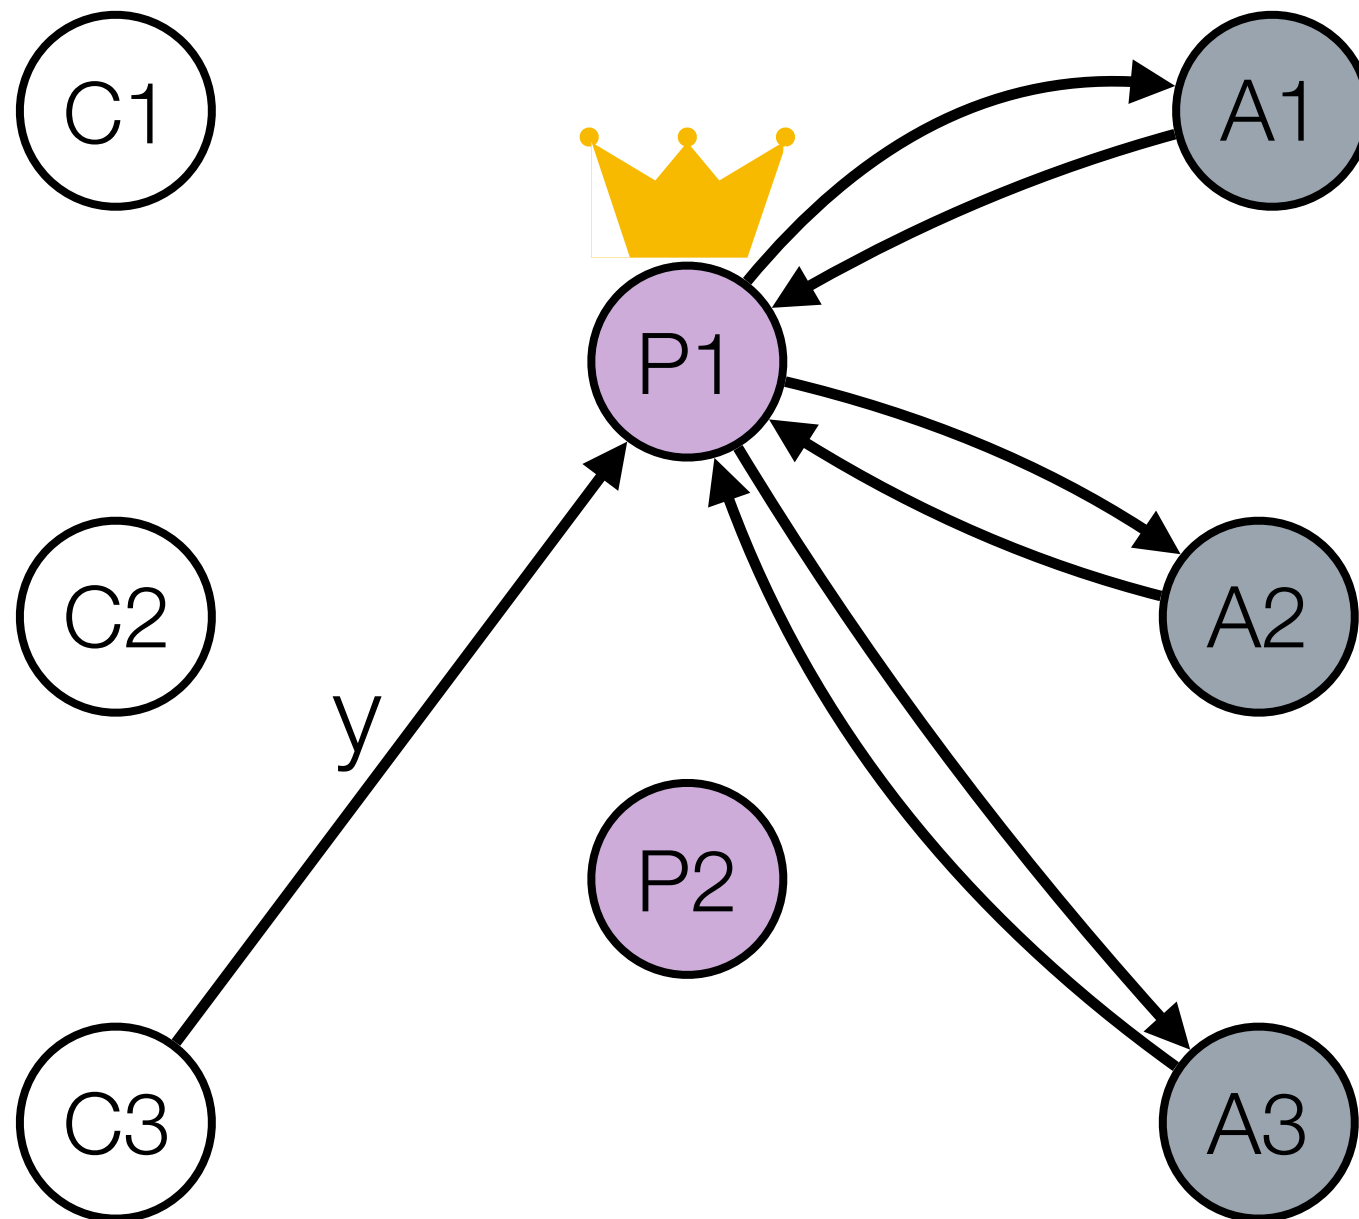


clients

$f+1$
proposers

$2f+1$
acceptors

$f+1$
replicas

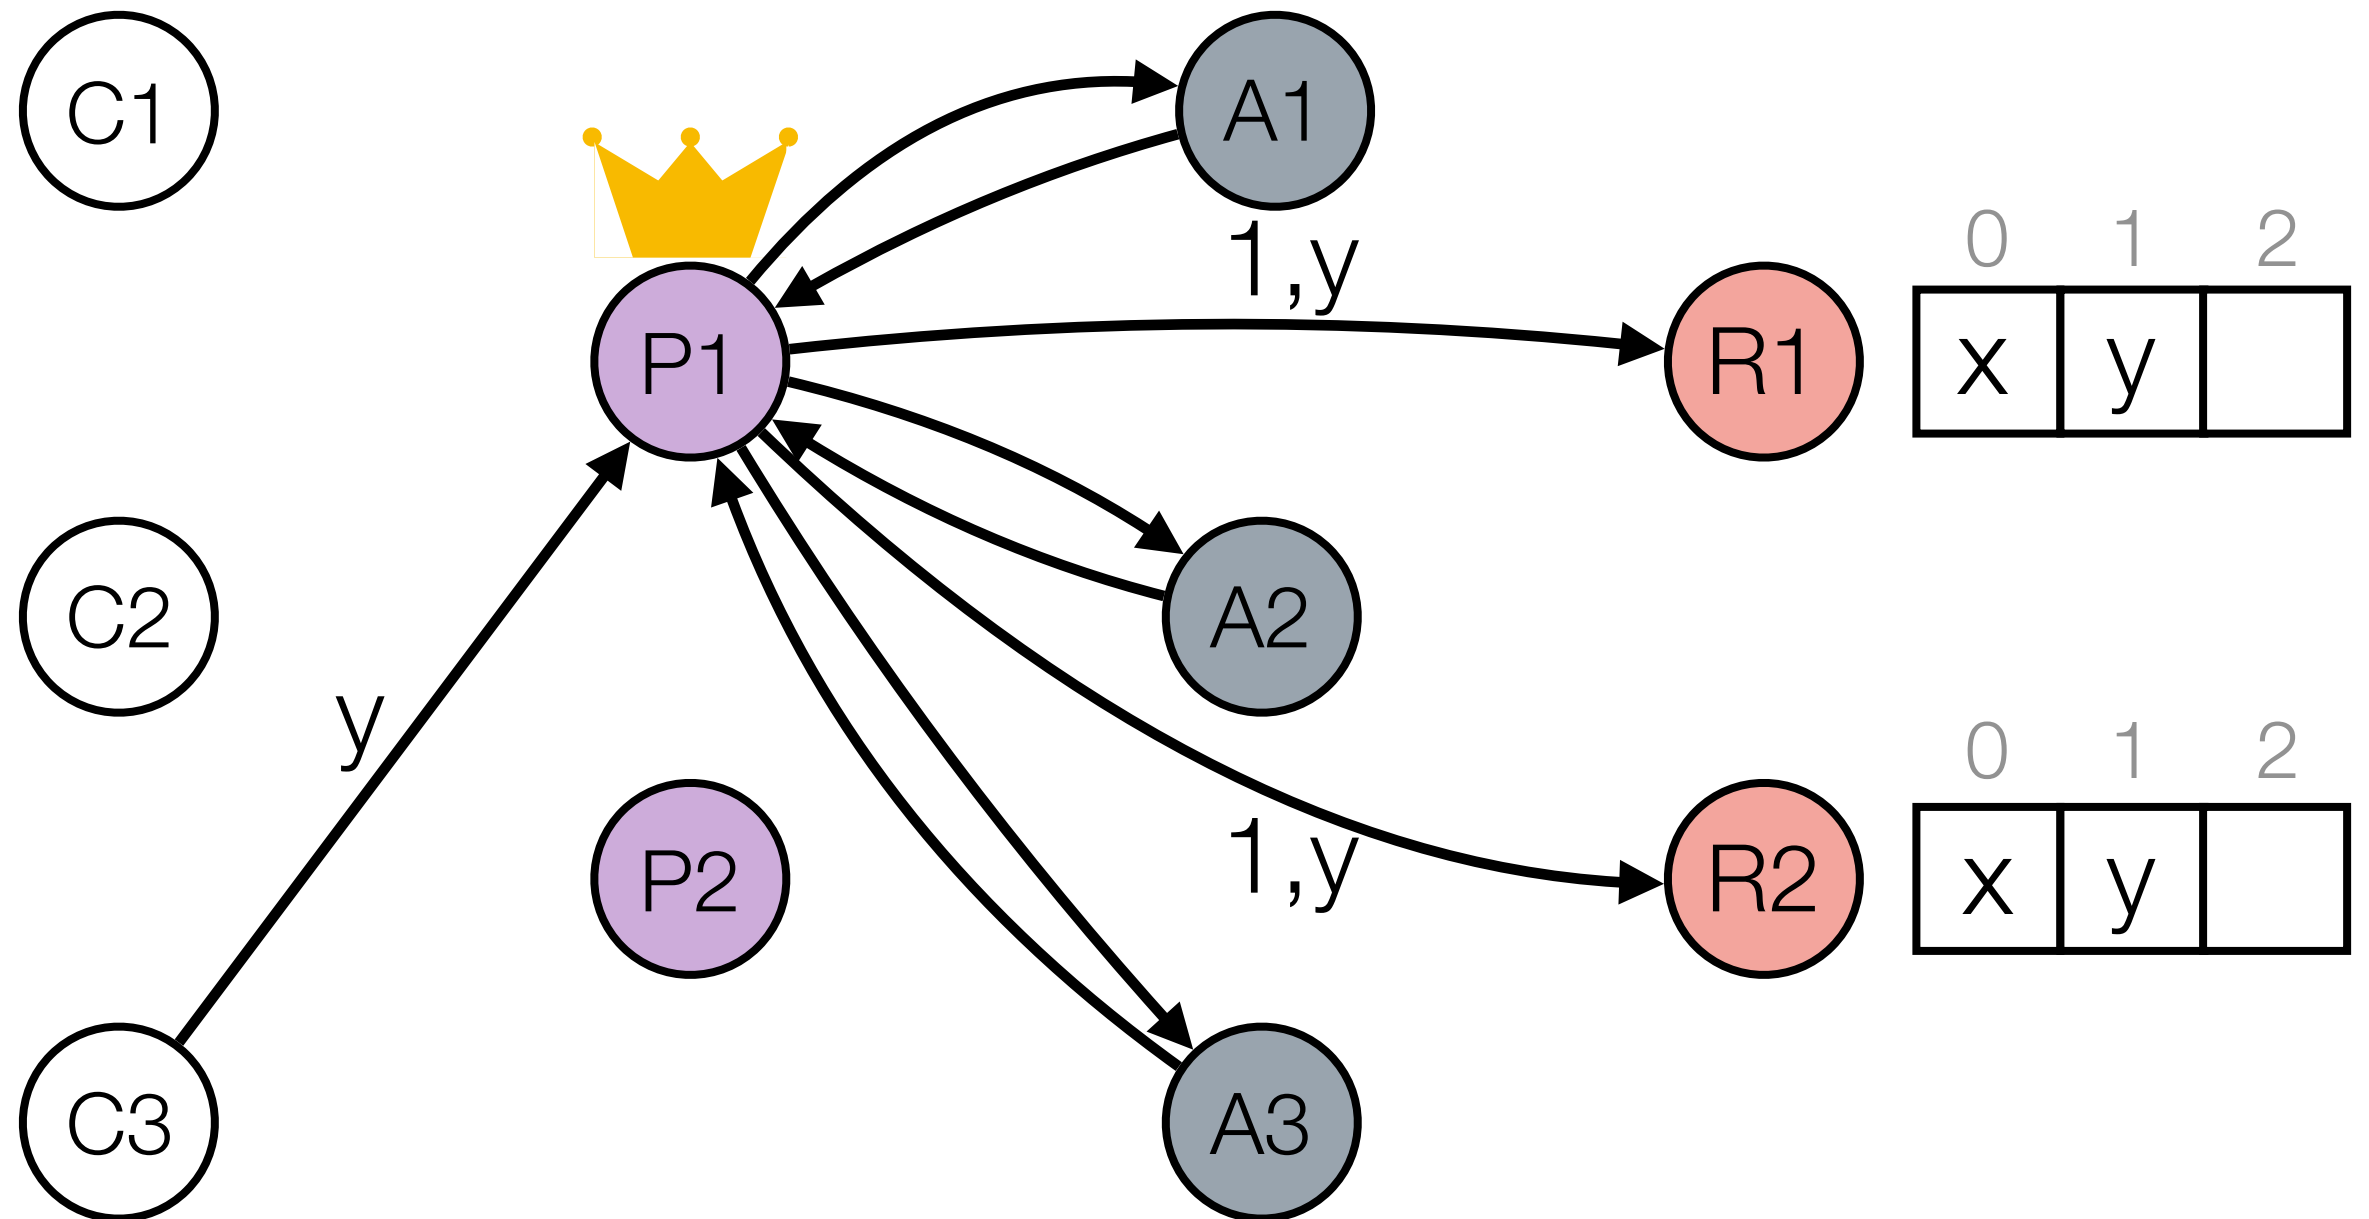


clients

$f+1$
proposers

$2f+1$
acceptors

$f+1$
replicas

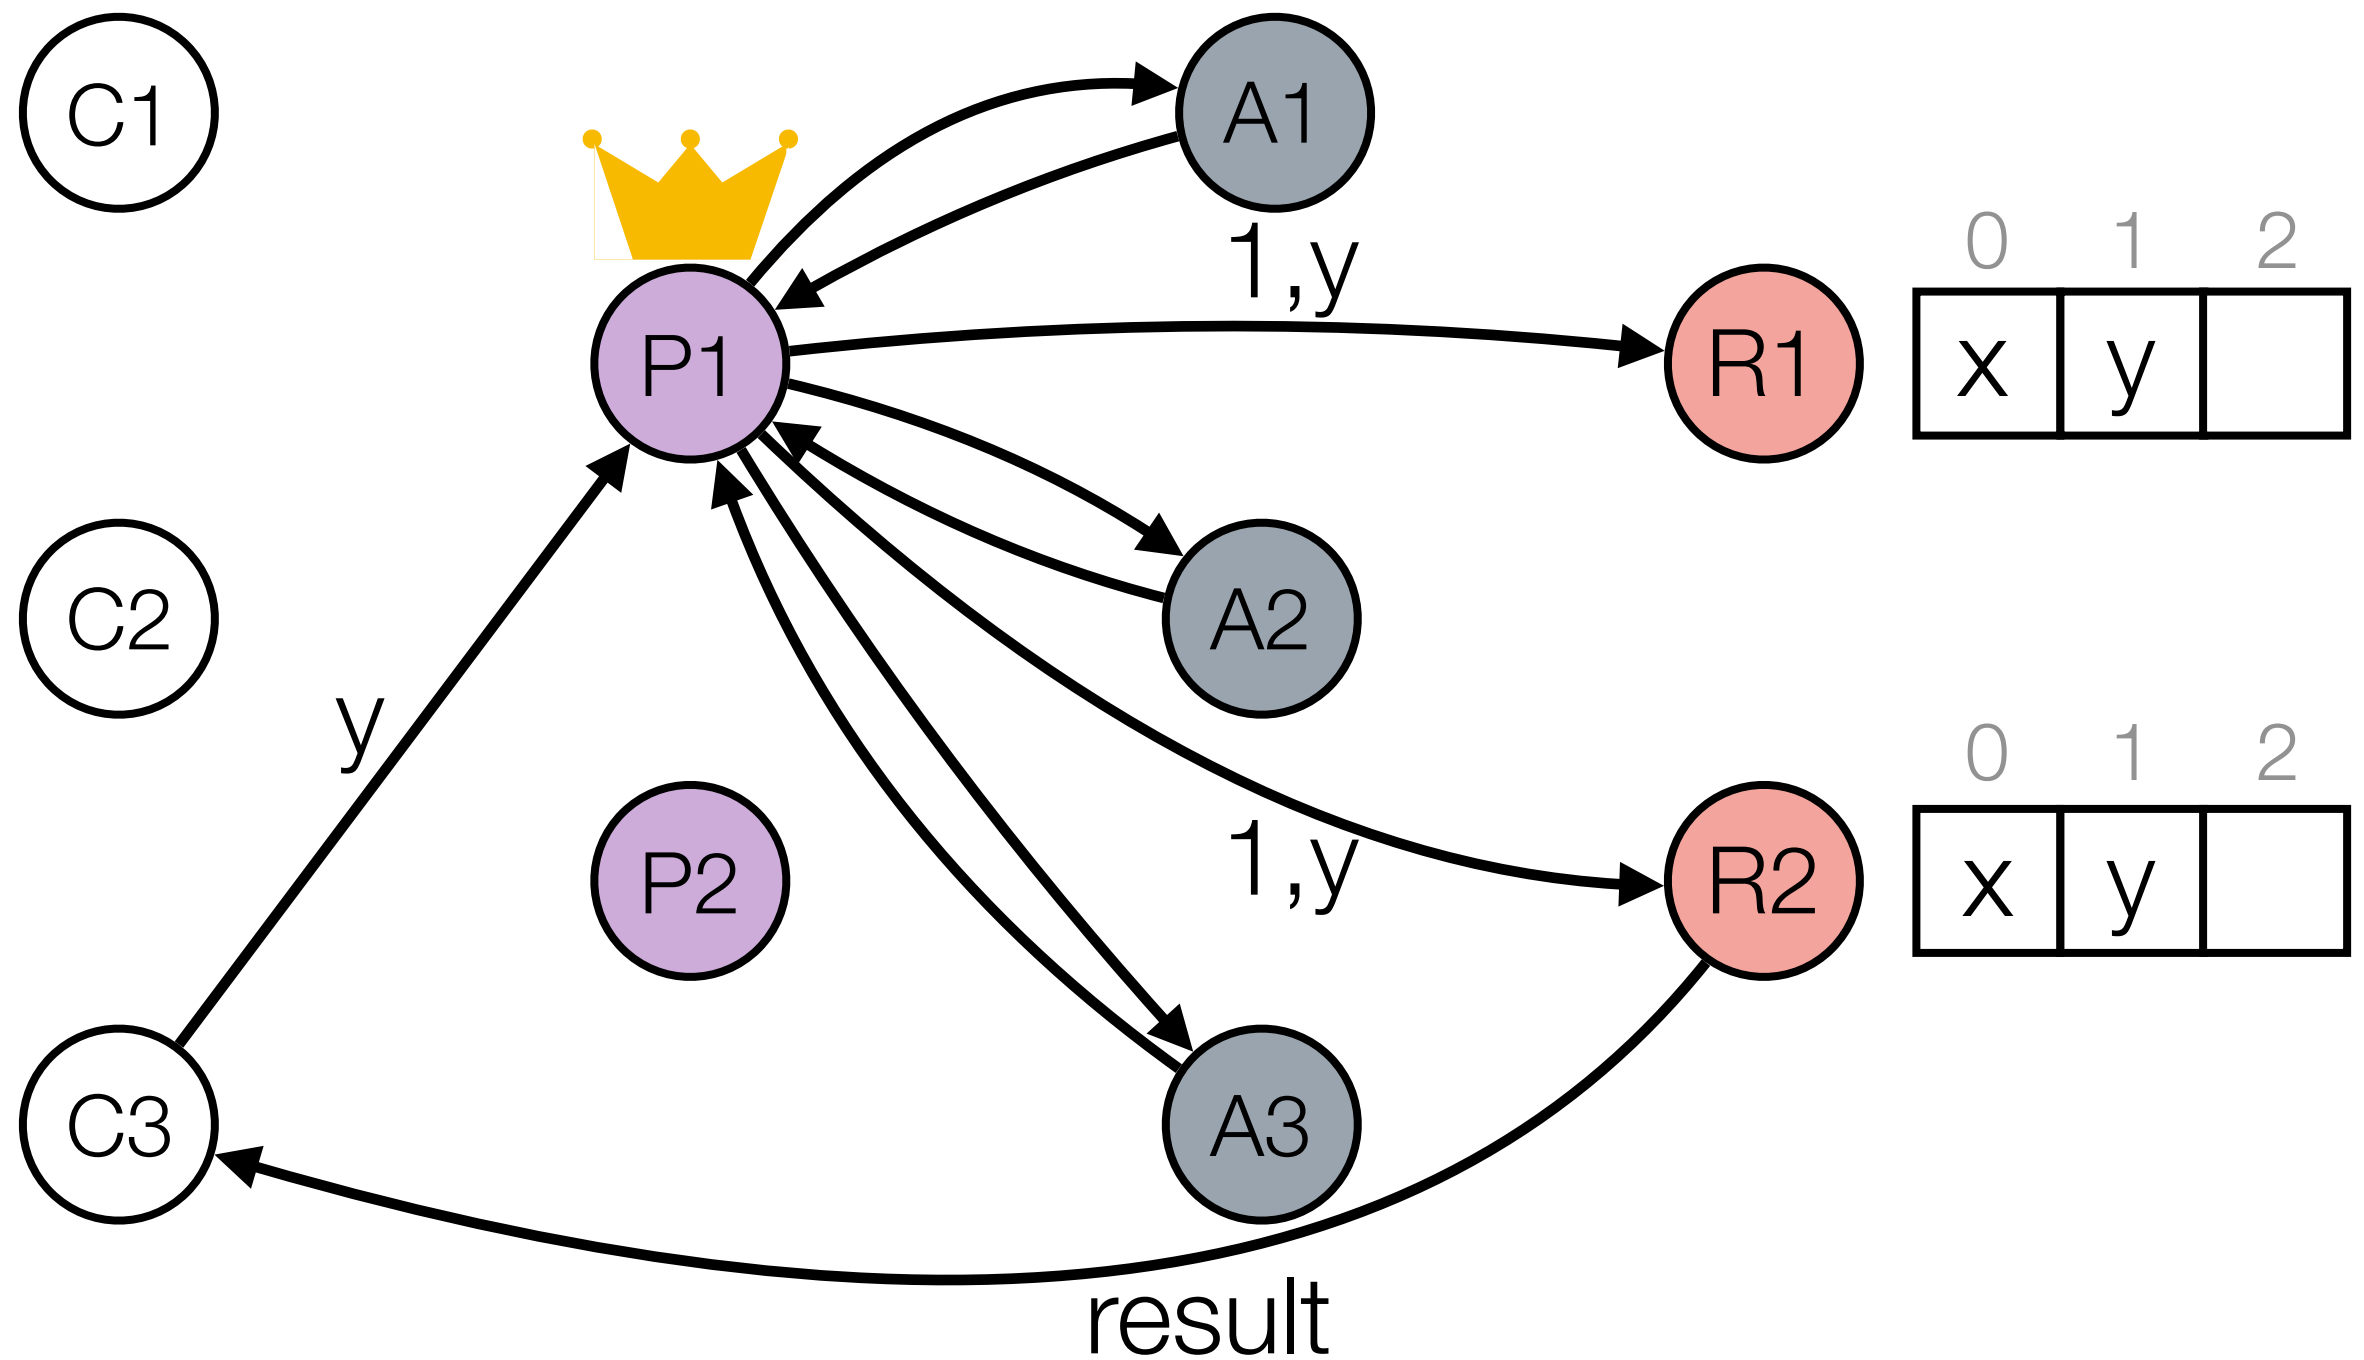


clients

$f+1$
proposers

$2f+1$
acceptors

$f+1$
replicas

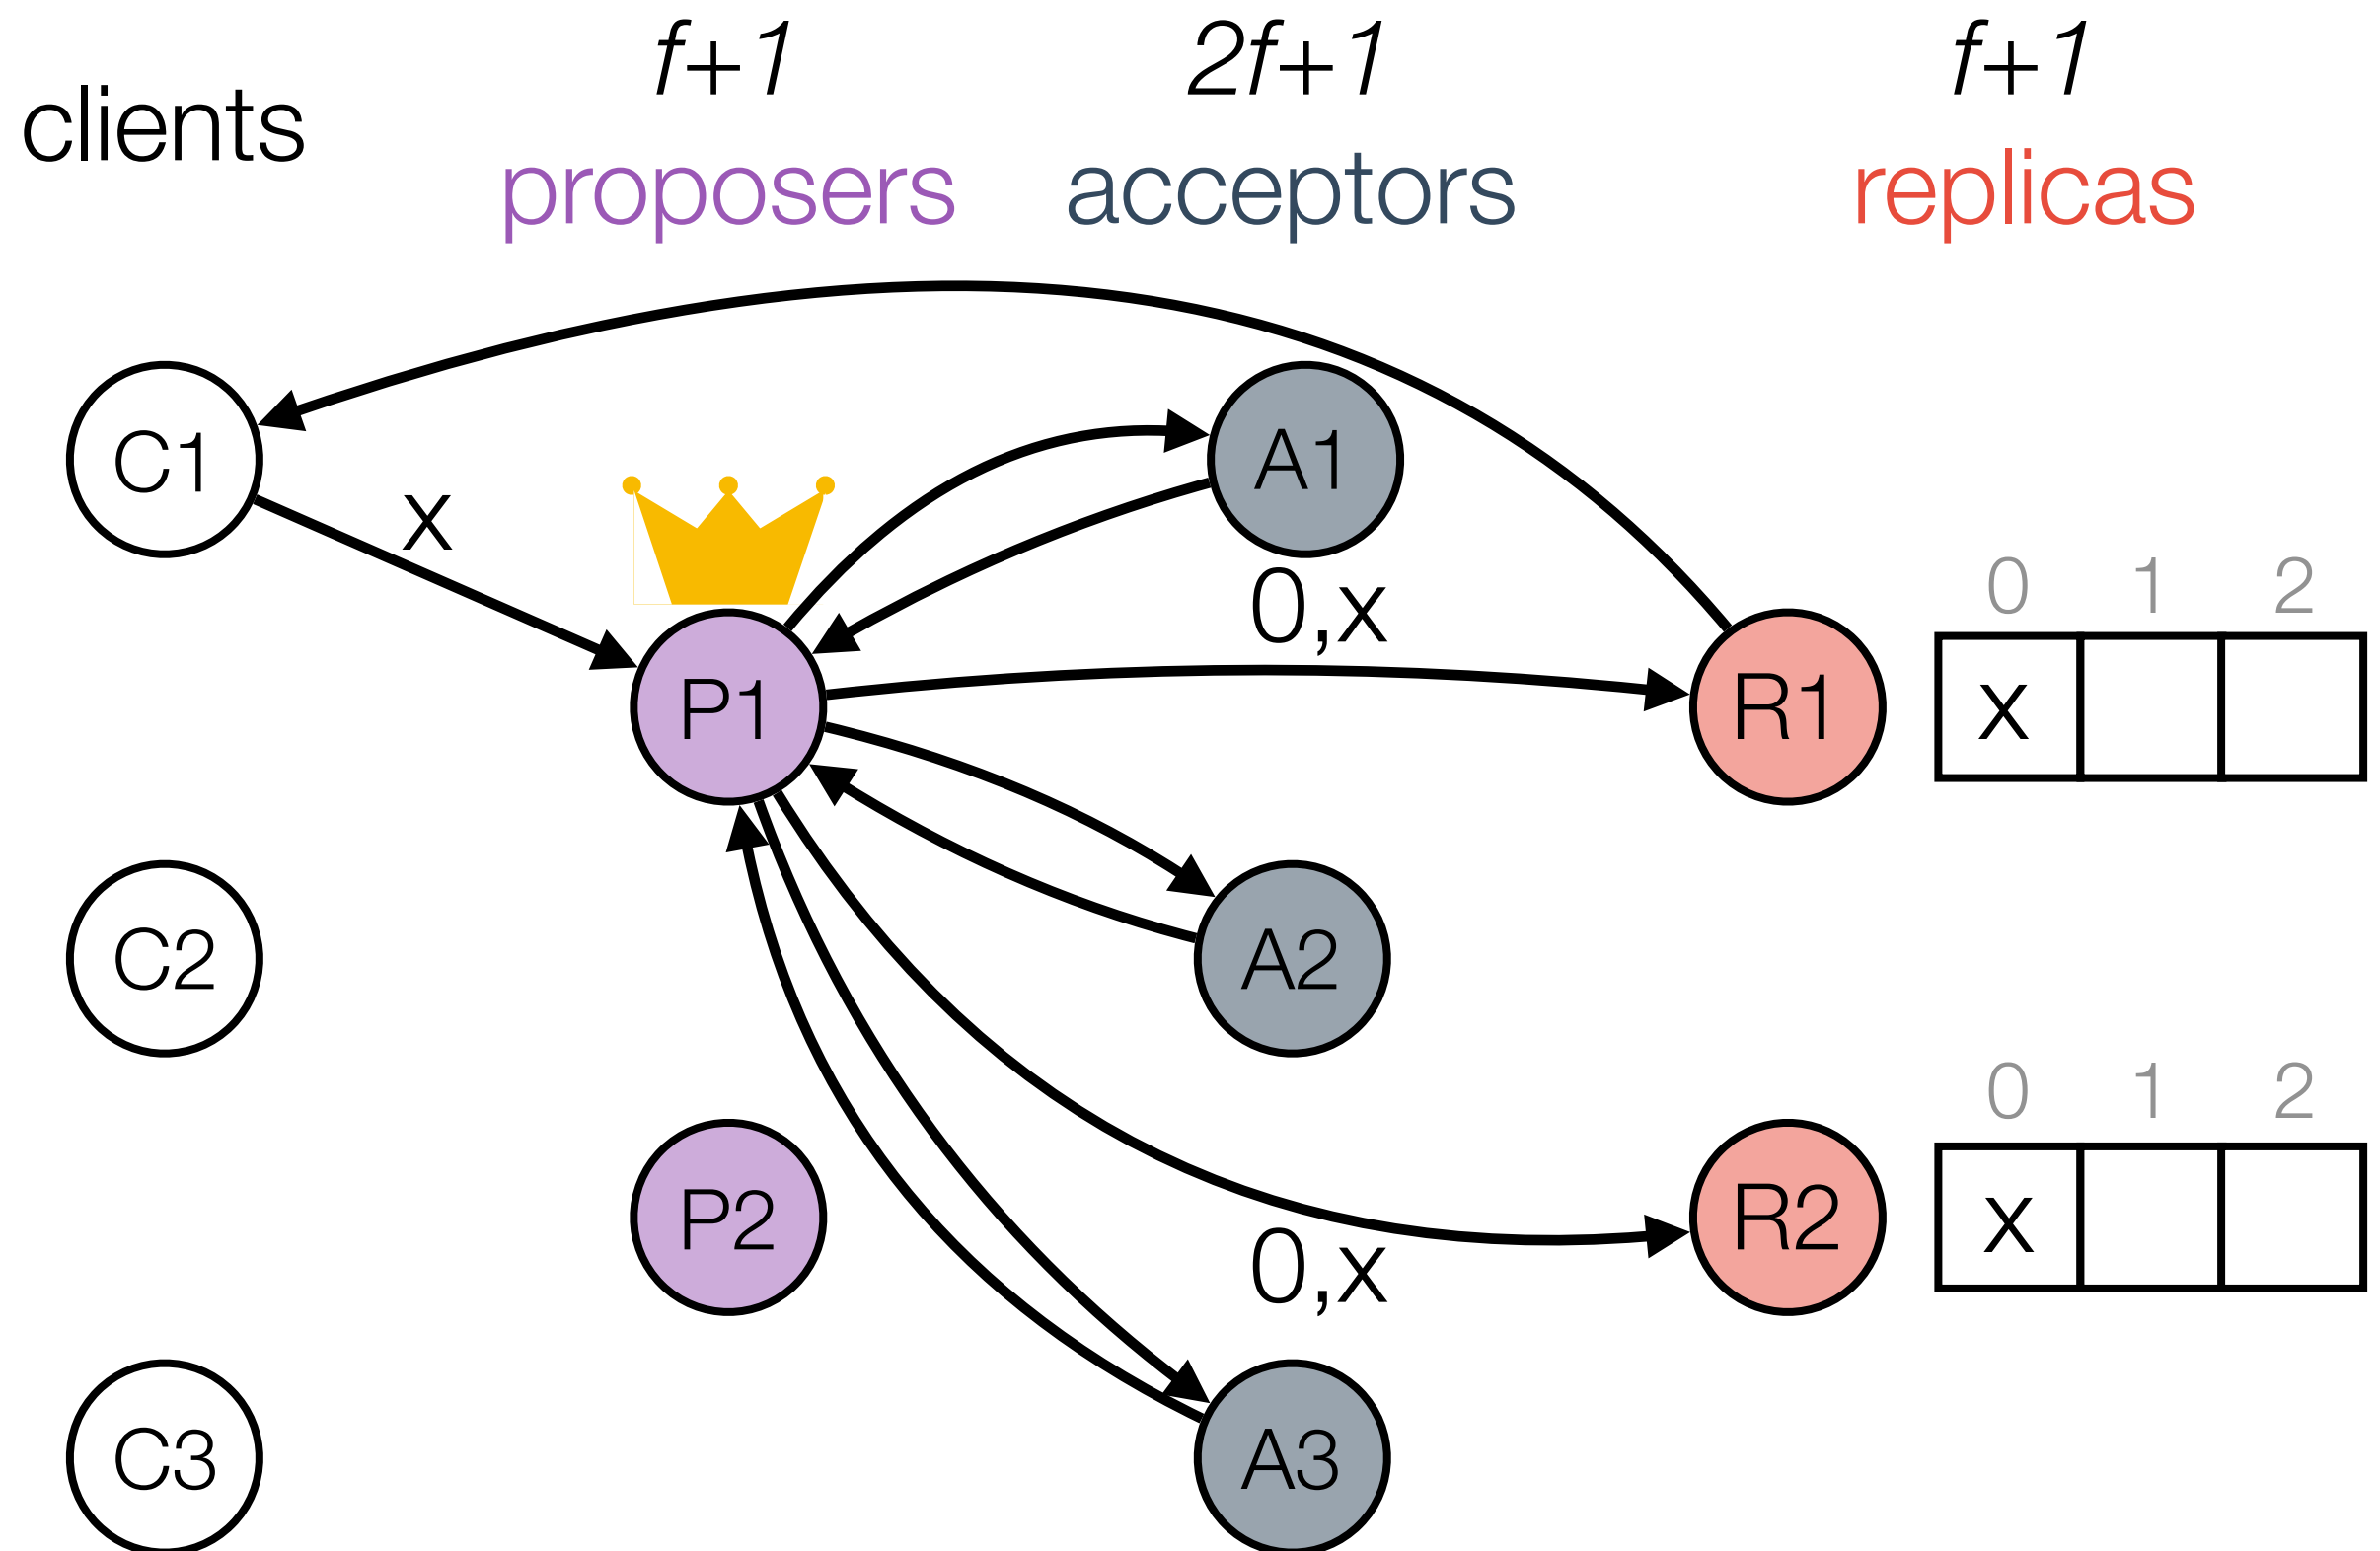


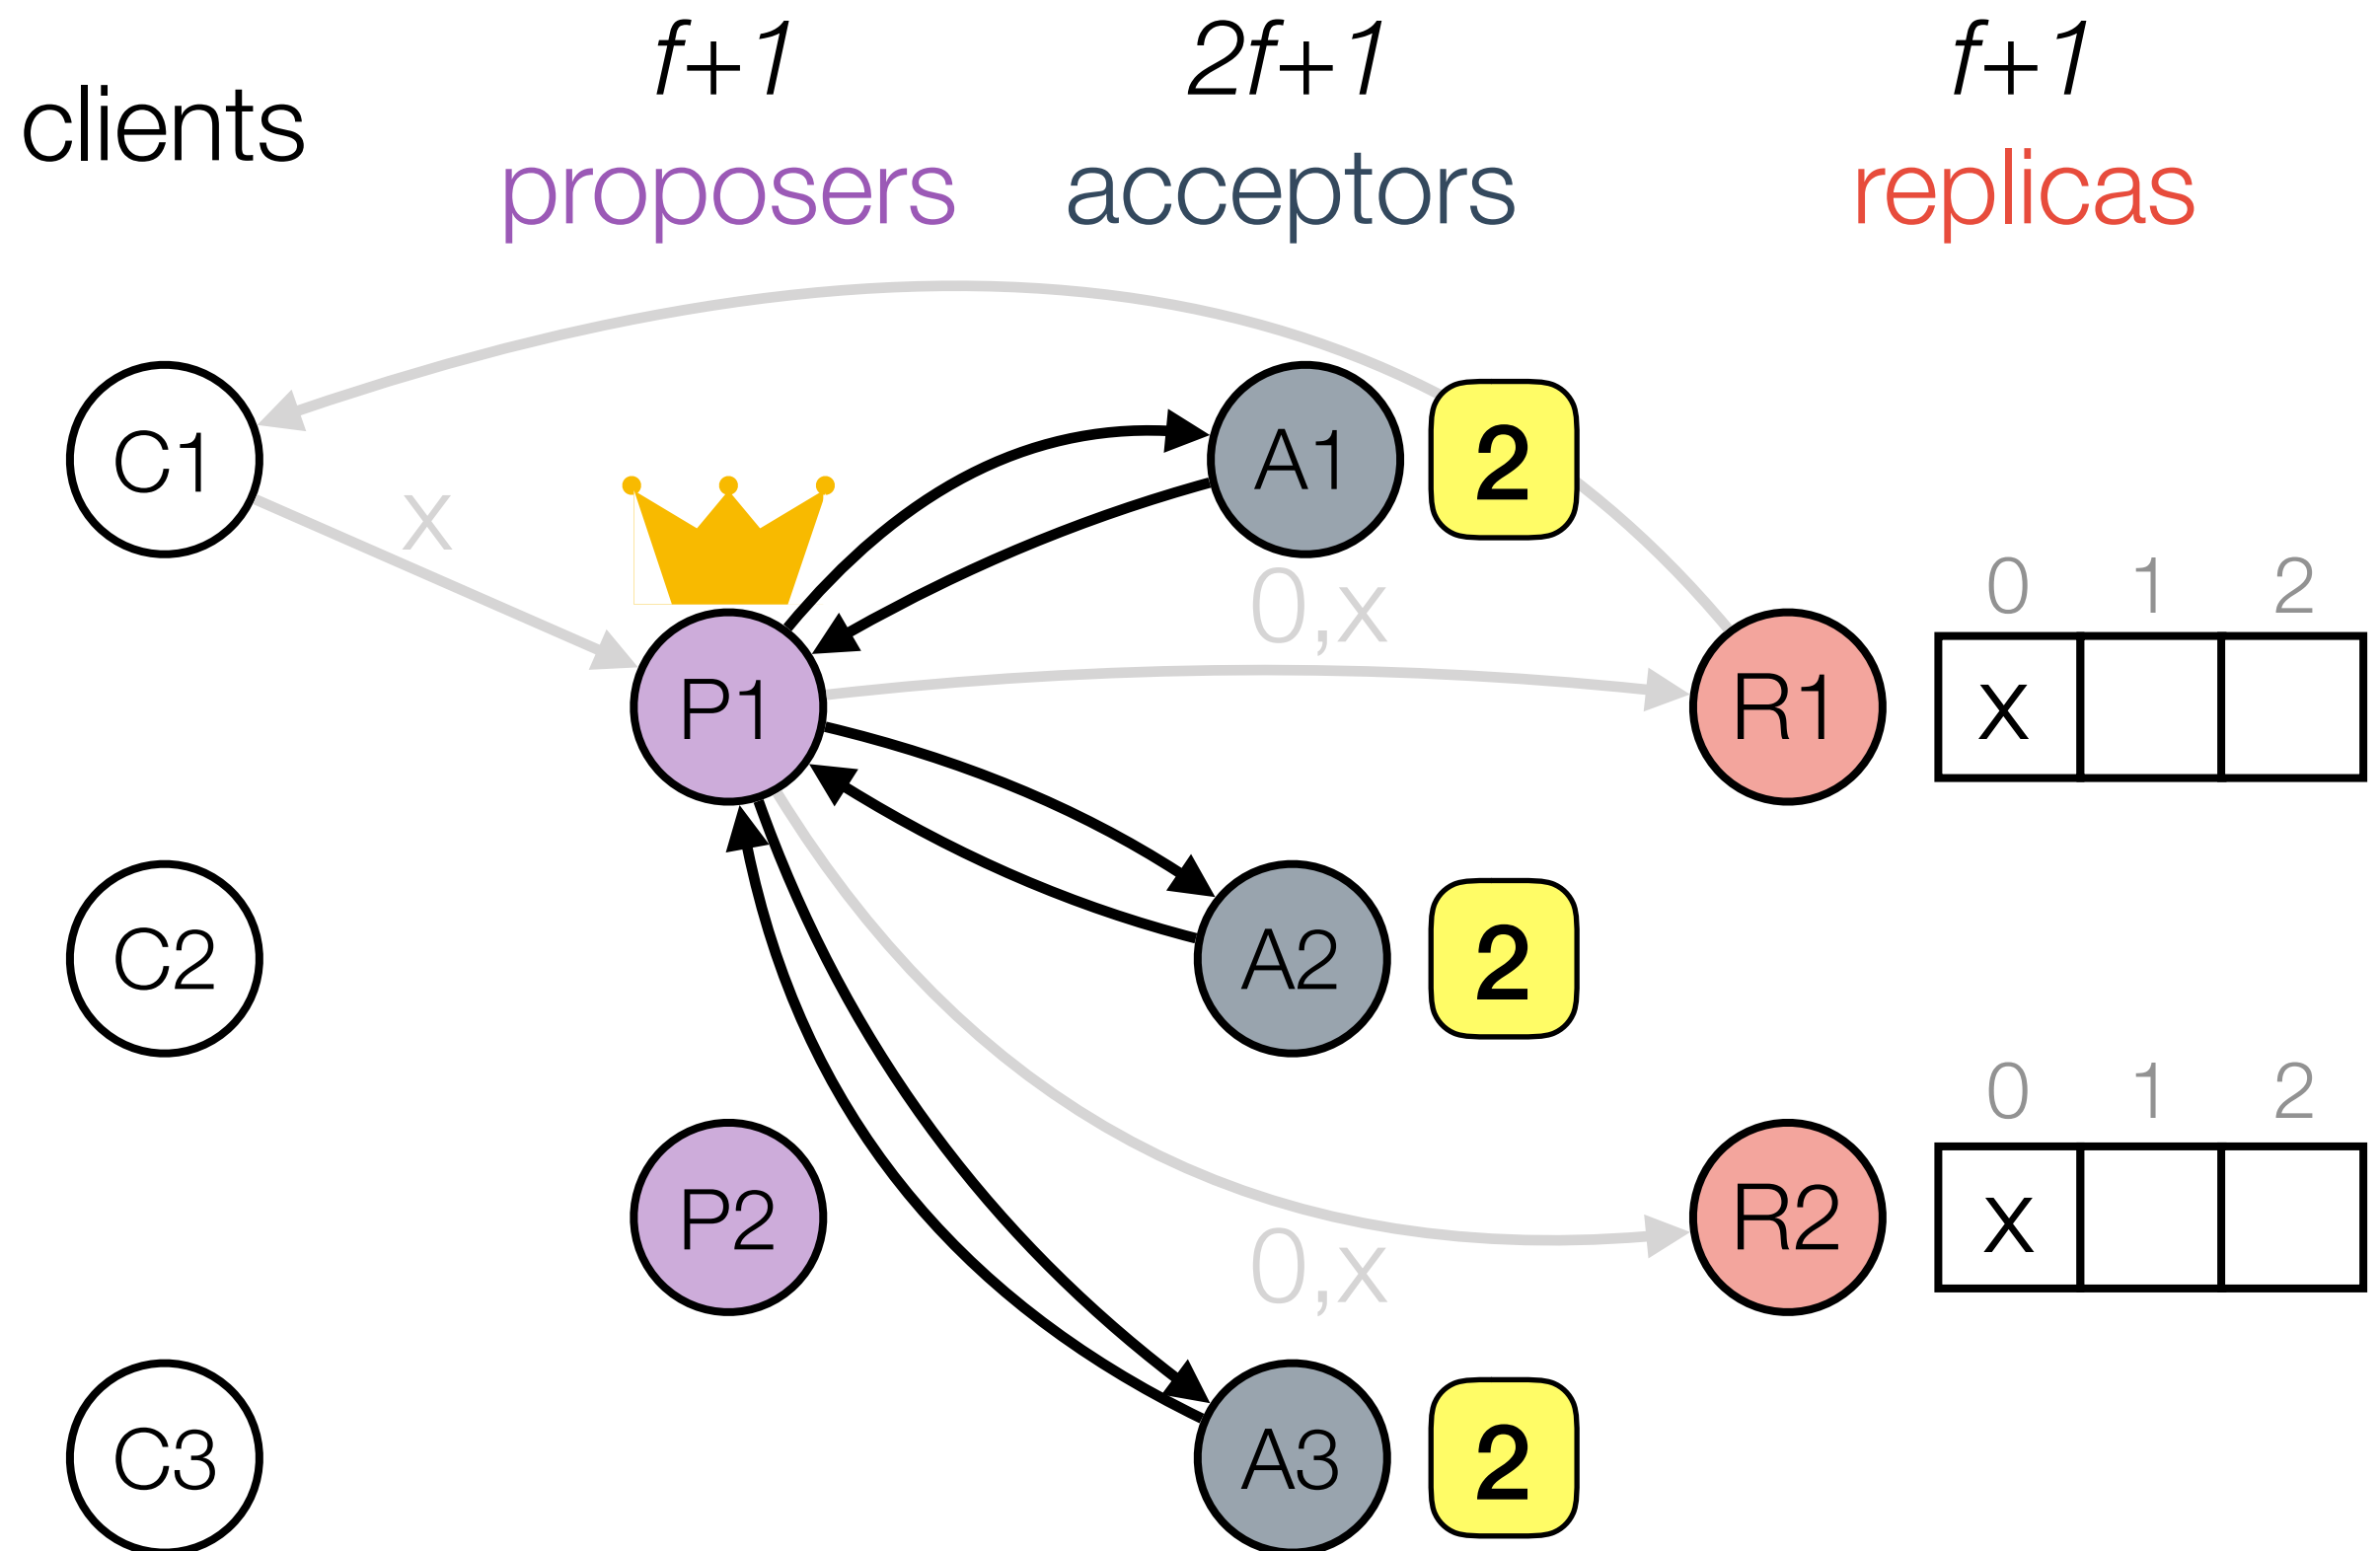
The MultiPaxos
leader is a
throughput
bottleneck

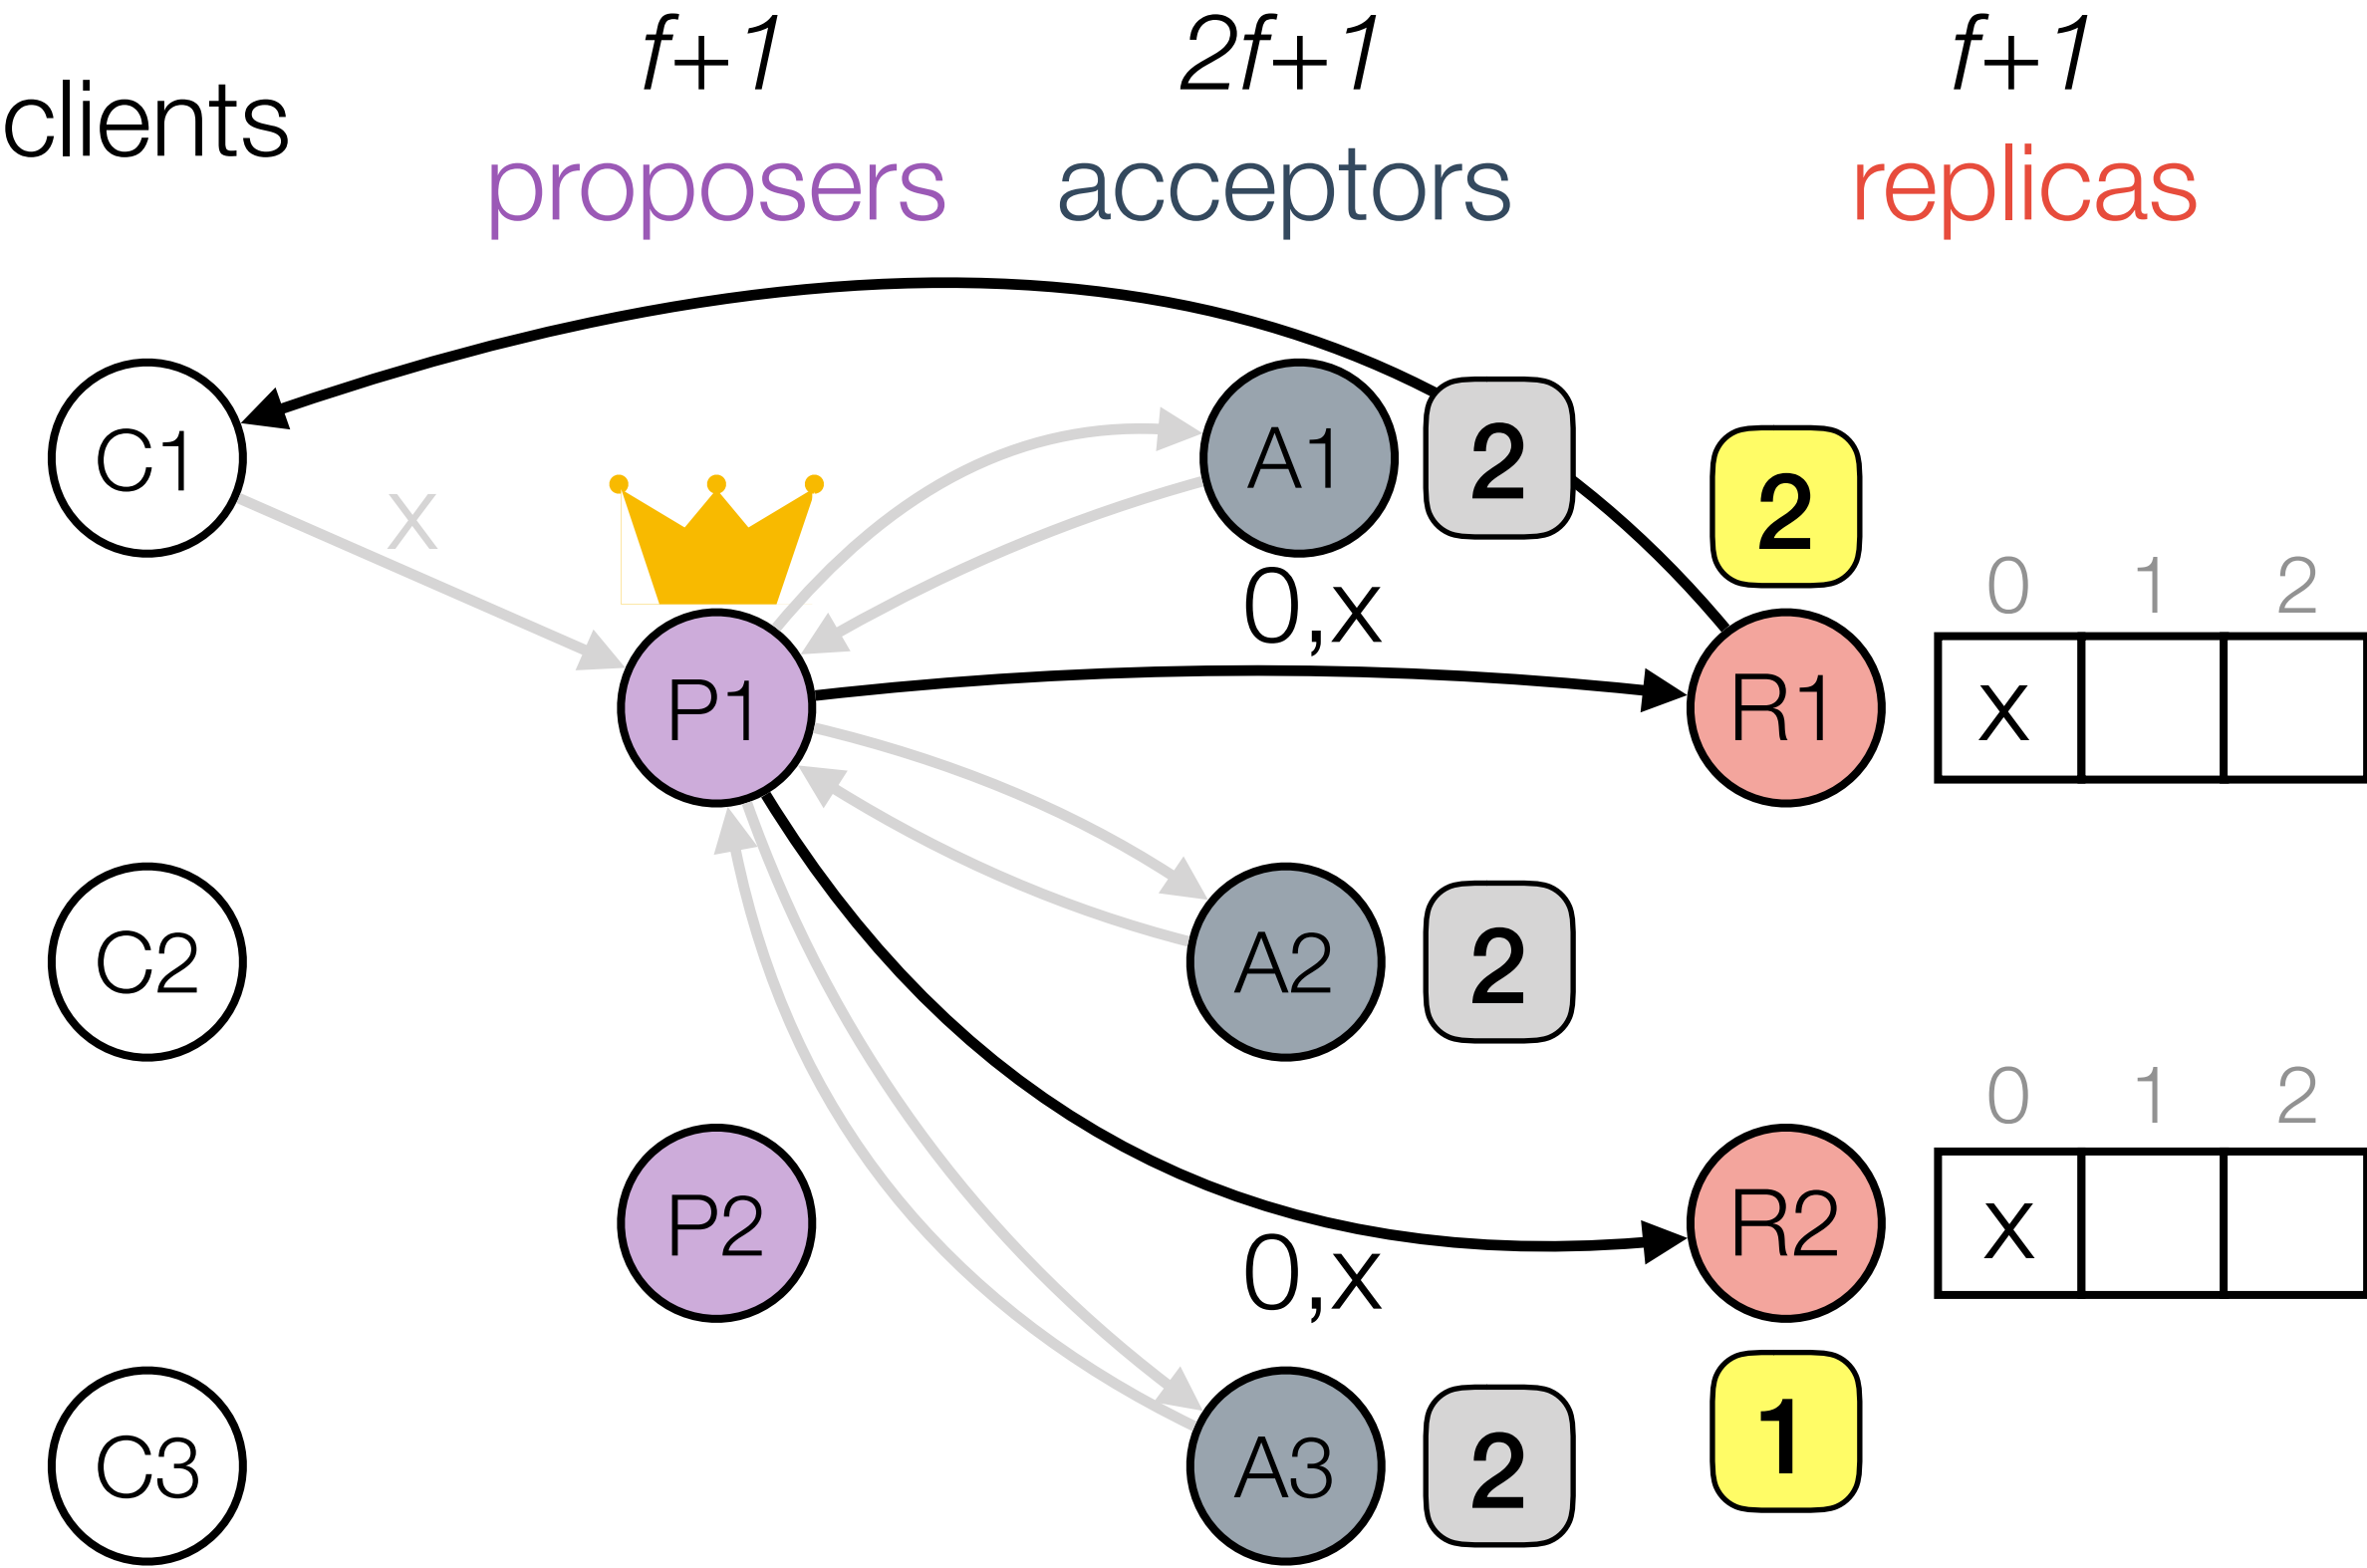
“The leader in Paxos is a bottleneck that limits throughput” -Mencius

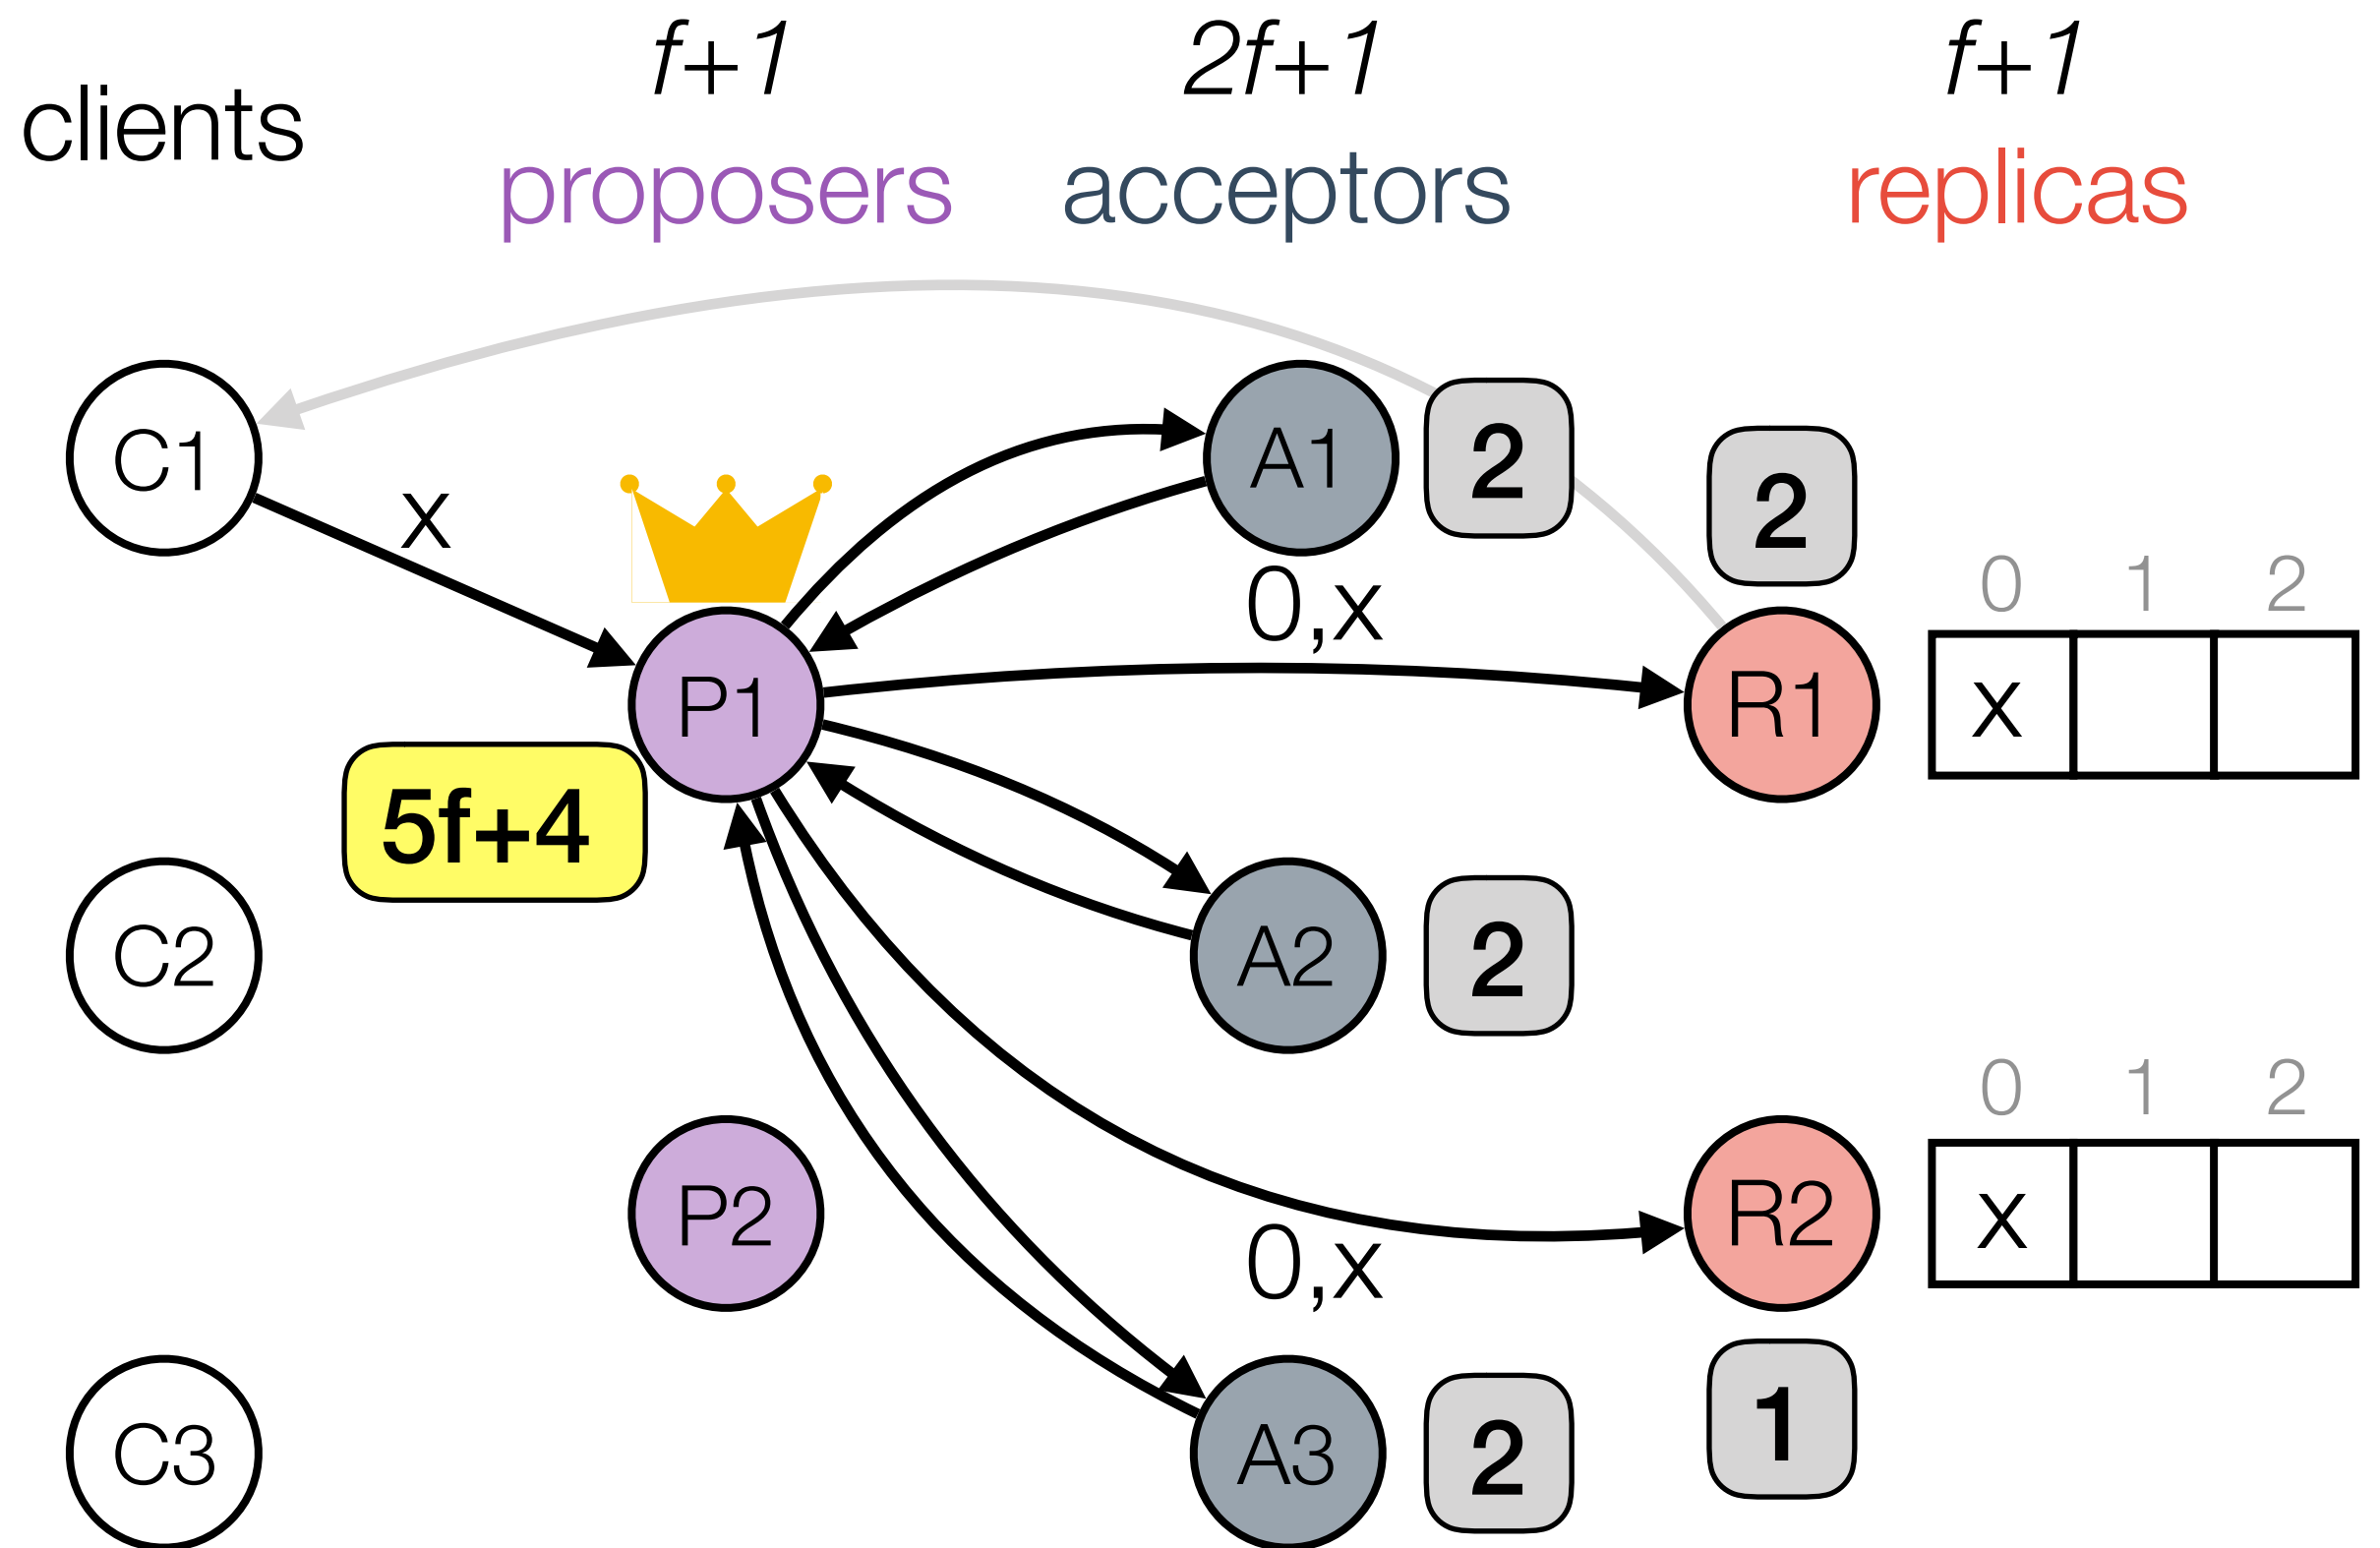
“It impairs scalability by placing a disproportionately high load on the master, which must process more messages than the other replicas” -EPaxos

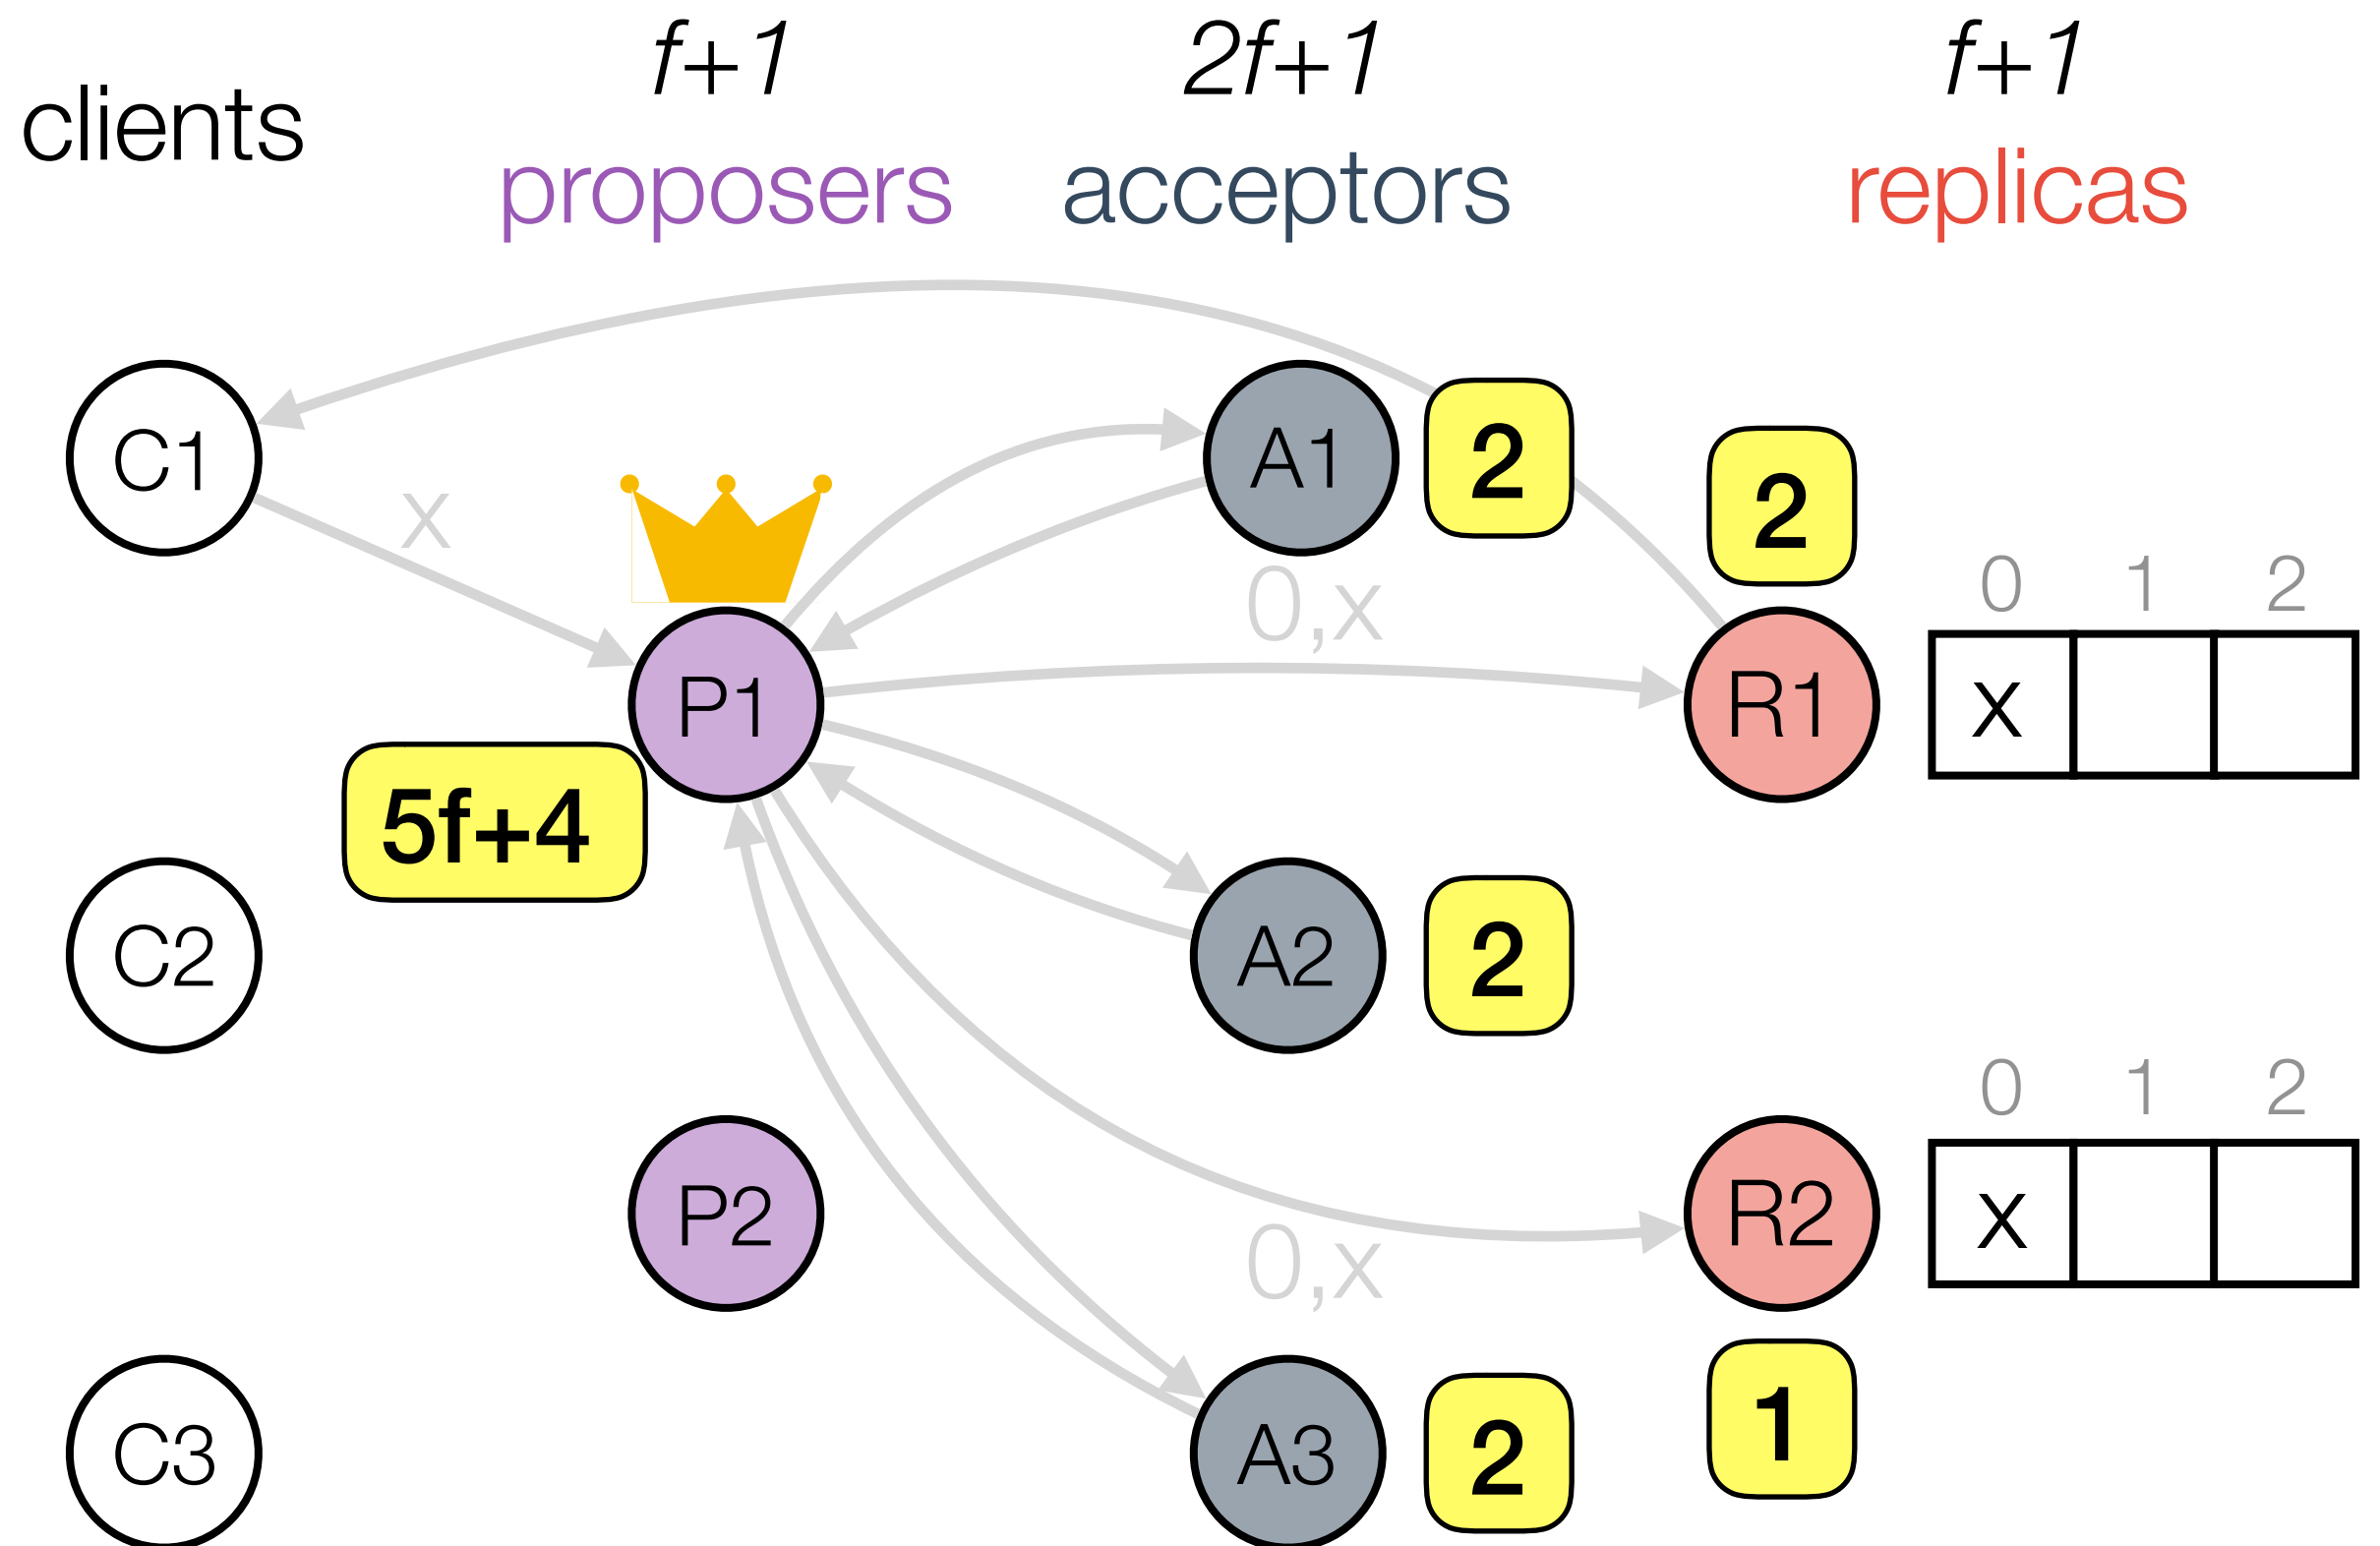
“In practice, [MultiPaxos] performance is tied to the performance of the leader” -
Caesar











Goal: to avoid the leader
bottleneck and increase
MultiPaxos' throughput.

Generalized Paxos

Generalized Consensus and Paxos

Leslie Lamport

3 March 2004
revised 15 March 2005
corrected 28 April 2005

Microsoft Research Technical Report MSR-TR-2005-33

EPaxos

There Is More Consensus in Egalitarian Parliaments

Julian Marra, David G. Andersen, Michael Kaminsky
Carnegie Mellon University and Intel Labs

Abstract

This paper describes the design and implementation of Egalitarian Paxos (EPaxos), a new distributed consensus algorithm based on Paxos. EPaxos achieves three goals: (1) optimal commit latency in the wide-area when tolerating one and two failures, under realistic conditions; (2) uniform load balancing across all replicas (thus achieving high throughput); and (3) graceful performance degradation when replicas are slow or crash.

Egalitarian Paxos is to our knowledge the first protocol to achieve the previously stated goals efficiently—that is, requiring only a simple majority of replicas to be non-faulty, using a number of messages linear in the number of replicas to choose a consensus, and committing consensus after just one communication round (one round trip) in the common case, or after at most two rounds in any case. We prove Egalitarian Paxos's properties theoretically and demonstrate its advantages empirically through an implementation running on Amazon EC2.

1. Introduction

Distributed computing places two main demands on replication protocols: (1) high throughput for replication inside a computing cluster; and (2) low latency for replication across data centers. Today's clusters are fault-tolerant distributed systems such as Chubby [4], Hadoop [21], or ZooKeeper [12] for activities including operation spanning, coordination, leader election, and resource discovery. Many databases are replicated redundantly from different continents, requiring geo-replication [2, 8].

An important limitation on these systems is that during efficient, failure-free operation, all clients communicate with a single master (or leader) server at all times.

Permission is granted to digital or hard copies of part or all of this work for personal or classroom use, provided that the copies are not made or distributed for profit or commercial advantage, and that copies bear this notice and the full citation on the first page. Copyright for third-party components of this work must be honored. For all other uses, contact the Owner/Author.

Copyright is held by the Owner/Author(s).
SIGOPS, Nov. 3–6, 2014, Farnham, Surrey, UK.
ACM 978-1-4555-2588-8/14/0001.
http://dx.doi.org/10.1145/2513291.2513328

This optimization—termed “Multi-Paxos” for systems based on the Paxos protocol [18]—is important to achieving high throughput in practical systems [7]. Changing the leader negates involving additional consensus machines that substantially reduce performance.

This algorithmic limitation has several important consequences. First, it impairs scalability by placing a disproportionate high load on the master, which must process more messages than the other replicas [23]. Second, when performing geo-replication, clients incur additional latency for communicating with a remote master. Third, as we show in this paper, traditional Paxos variants are sensitive to both long-term and transient load spikes and network delays that increase latency at the master. Finally, this single-master optimization can harm availability: if the master fails, the system cannot service requests until a new master is elected. Previously proposed solutions—such as partitioning or using proxy servers—are undesirable because they restrict the type of operations the cluster can perform. For example, a partitioned cluster cannot perform atomic operations across partitions without using additional techniques.

Egalitarian Paxos (EPaxos) has no designated leader process. Instead, clients can choose at every step which replica to submit a command to, and in most cases the command is committed without interfering with other consensus attempts. This allows the system to evenly distribute the load to all replicas, eliminating the first bottleneck identified above (i.e., having one server that must be on the critical path for all communication). EPaxos's flexible load distribution better handles permanently or transiently slow nodes, as well as the latency heterogeneity caused by geographical distribution of replicas; this substantially reduces both the median and tail commit latency. Finally, the system can provide higher availability and higher performance under failures because there is no transient interruption caused by leader election; there is no leader, and hence, no need for leader election, as long as more than half of the replicas are available.

We begin by reviewing the core Paxos algorithm, and the intuition behind Egalitarian Paxos in Section 2. We then describe several Paxos variants that reduce overhead or commit latency in Section 3. Throughout the paper we compare extensively against Multi-Paxos and two im-

Generalized Paxos

Generalized Consensus and Paxos

Leslie Lamport

3 March 2004
revised 15 March 2005
corrected 28 April 2005

Microsoft Research Technical Report MSR-TR-2003-33


Generalized Paxos

Generalized Consensus and Paxos

Leslie Lamport

3 March 2004
revised 15 March 2005
corrected 28 April 2005

Microsoft Research Technical Report MSR-TR-2003-33



To increase MultiPaxos'
throughput, we must
trade off complexity?

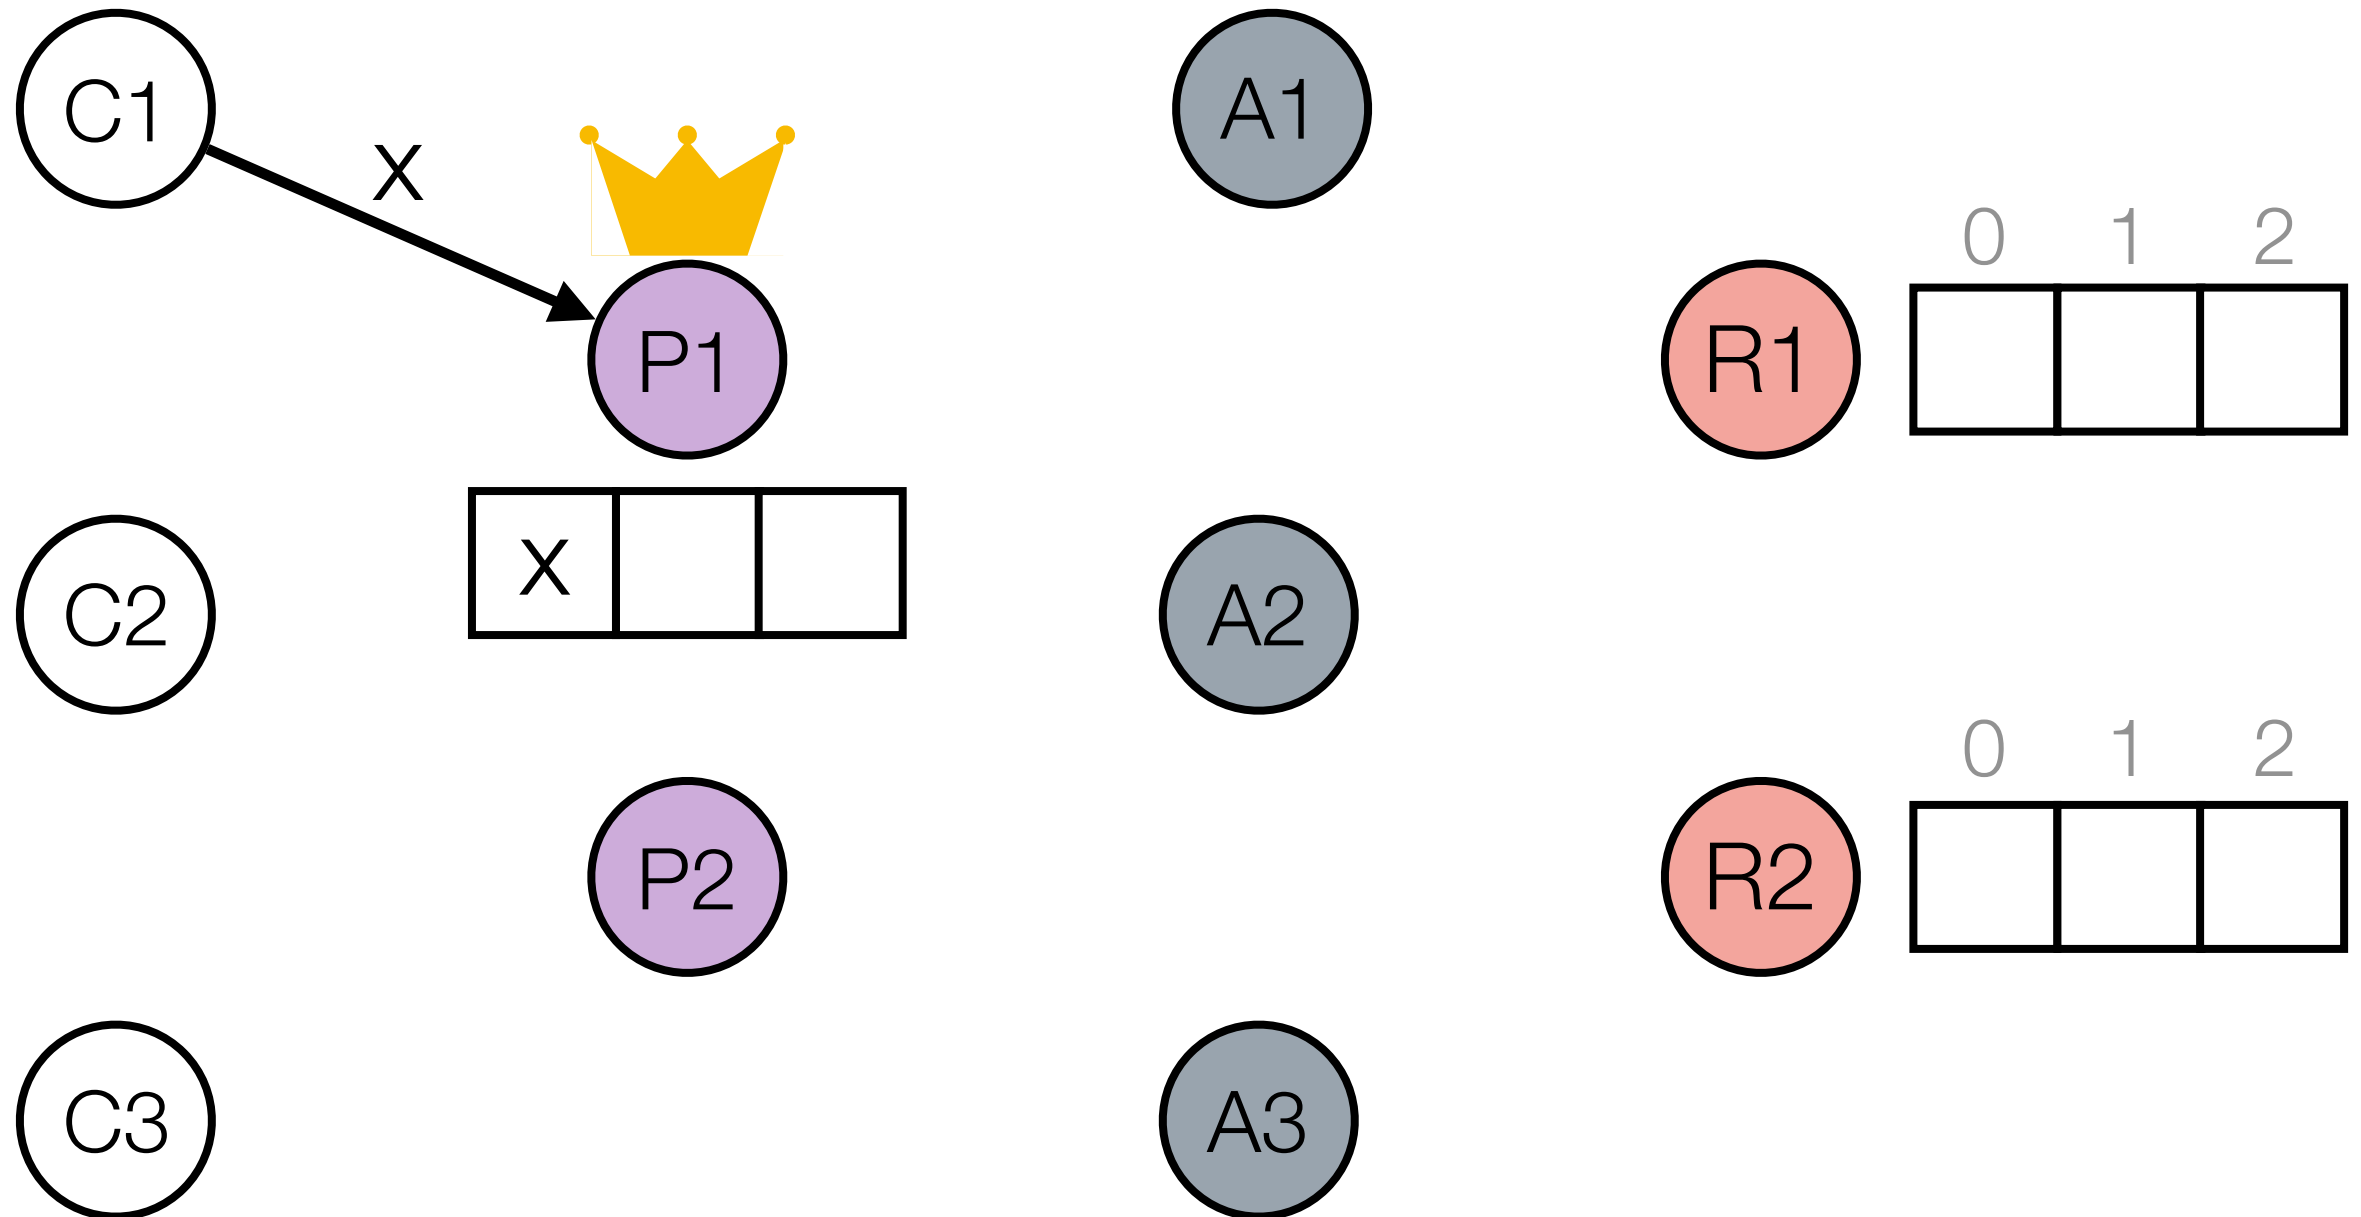
Key Insight: The MultiPaxos leader has **two independent** responsibilities.

clients

$f+1$
proposers

$2f+1$
acceptors

$f+1$
replicas

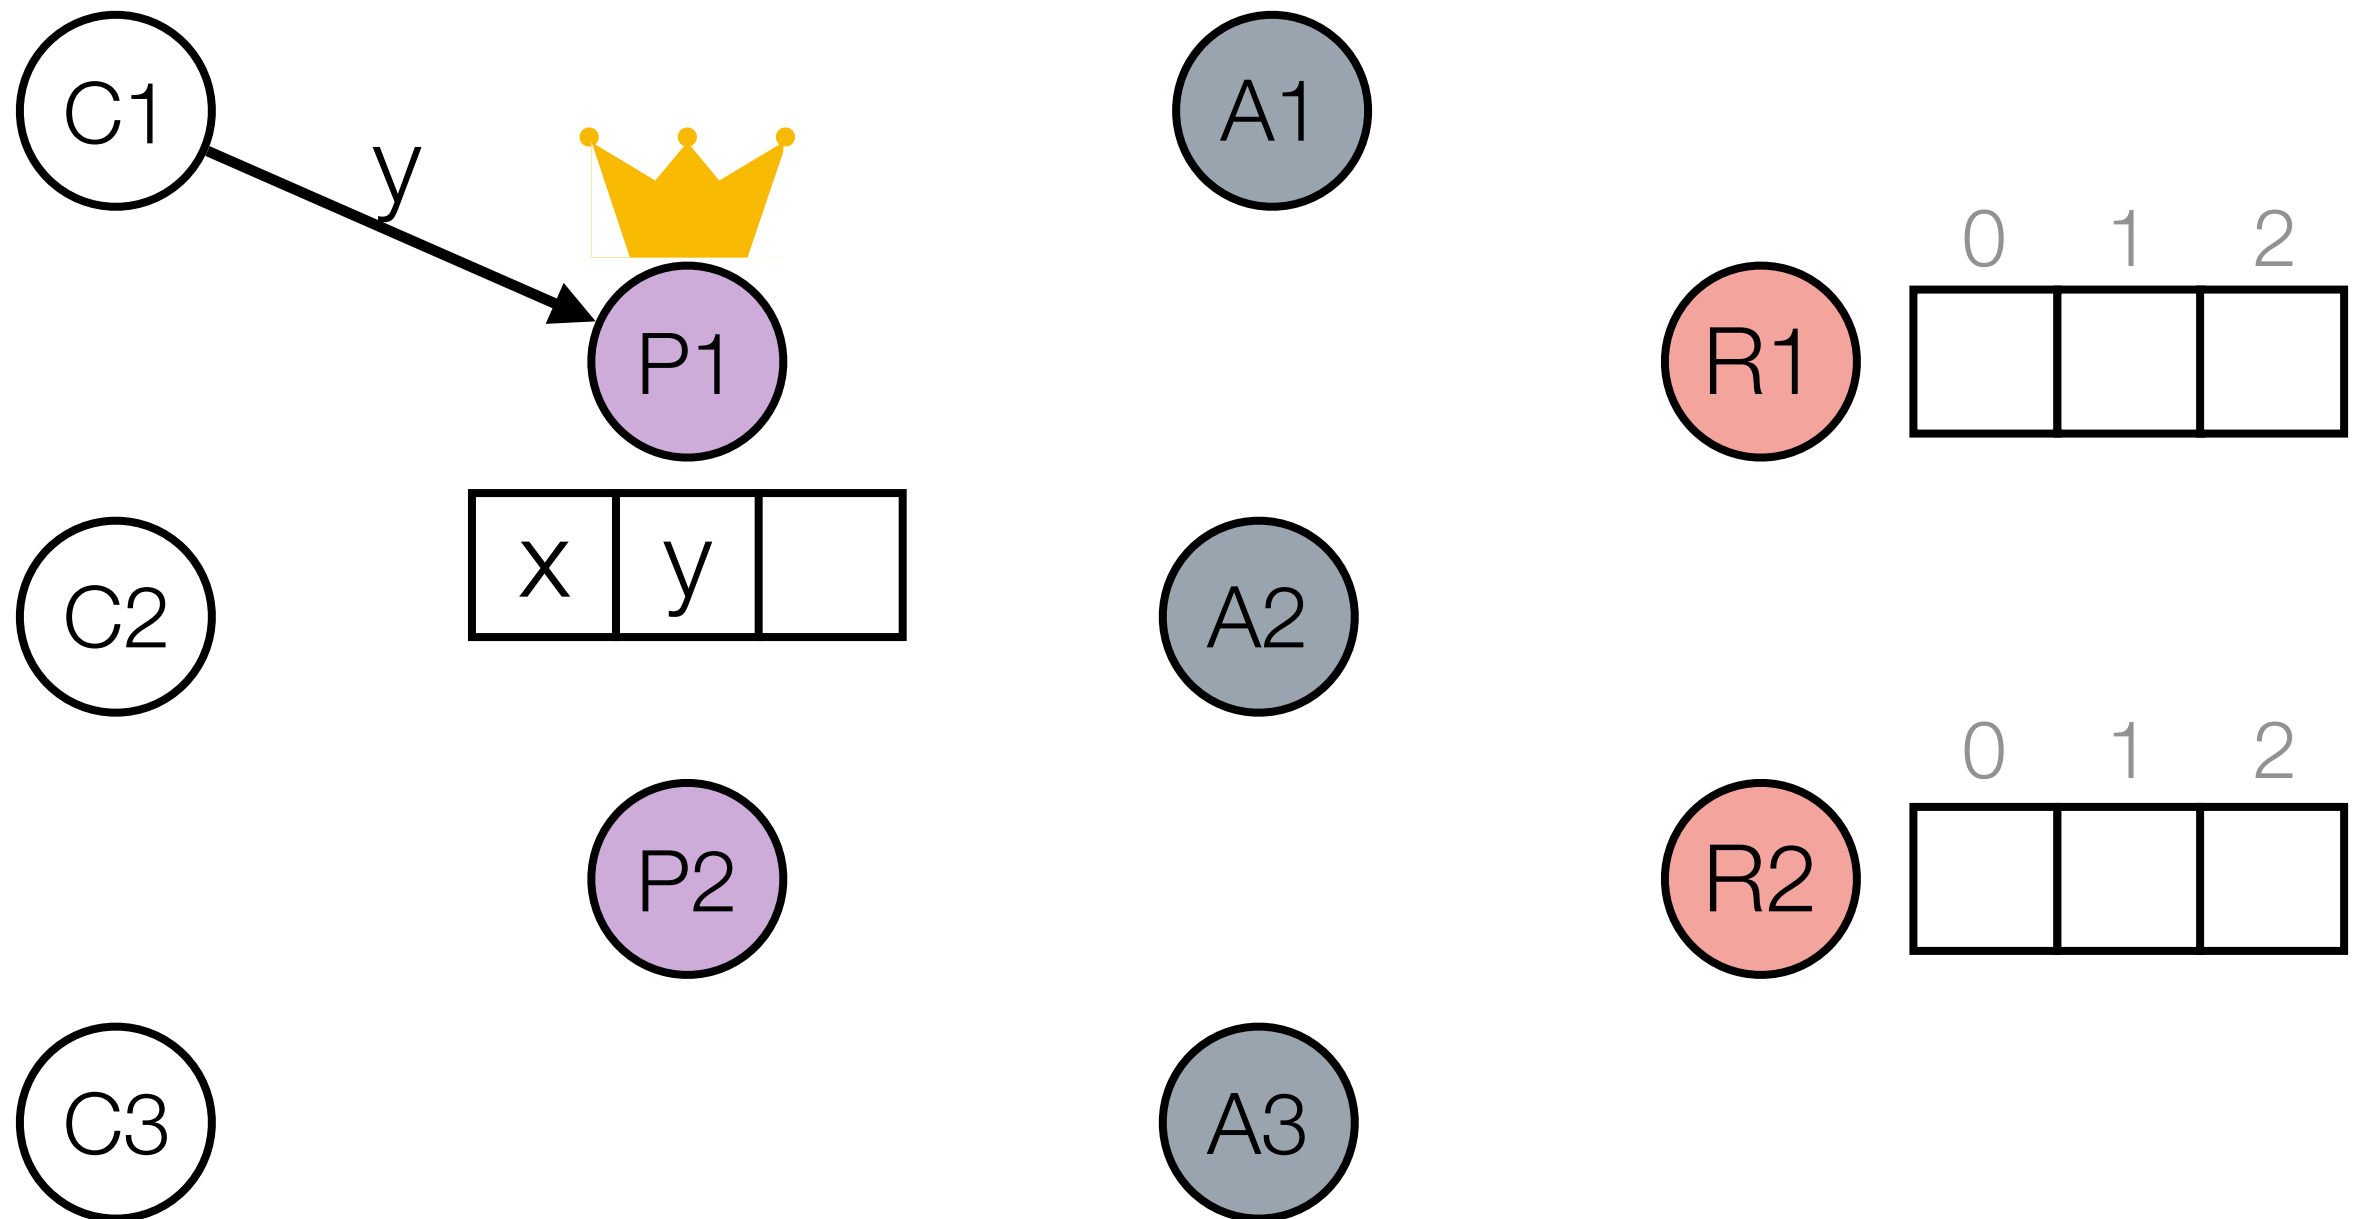


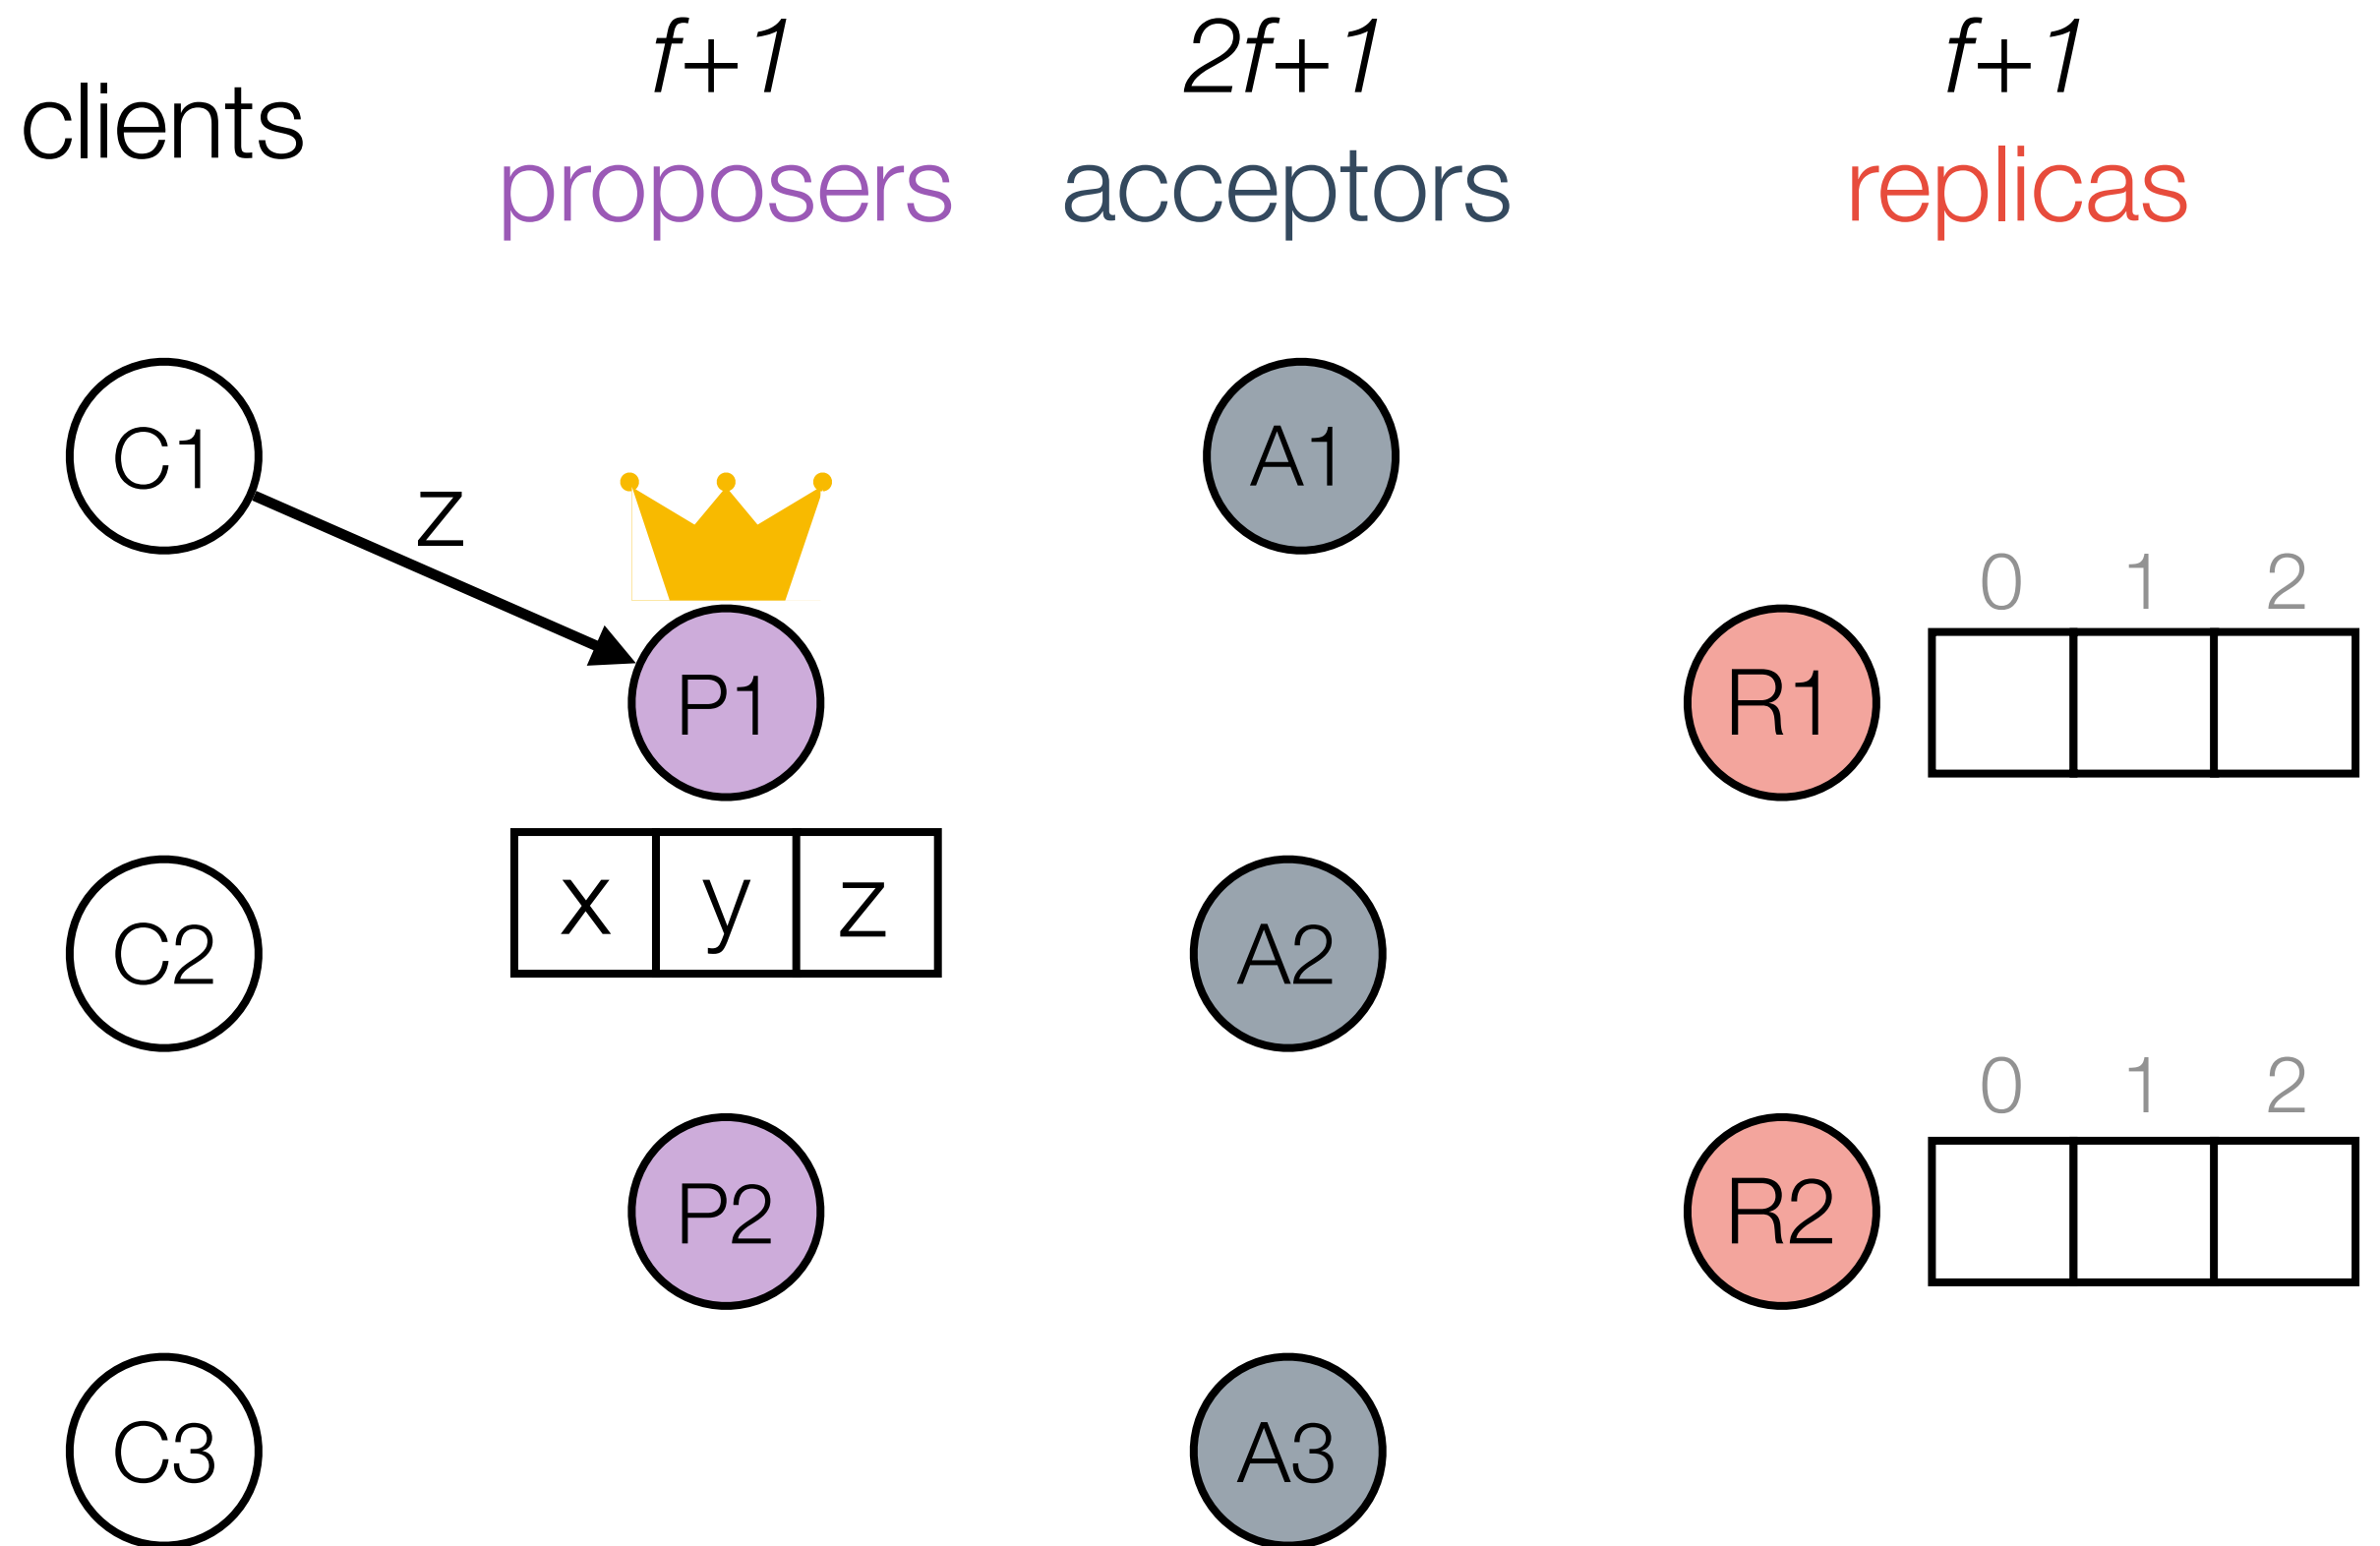
clients

$f+1$
proposers

$2f+1$
acceptors

$f+1$
replicas





clients

$f+1$
proposers


$2f+1$
acceptors

$f+1$
replicas

C1

C2

C3



R1

0	1	2

R2

0	1	2













Key Insight: **Decouple** the
MultiPaxos leader's two
responsibilities

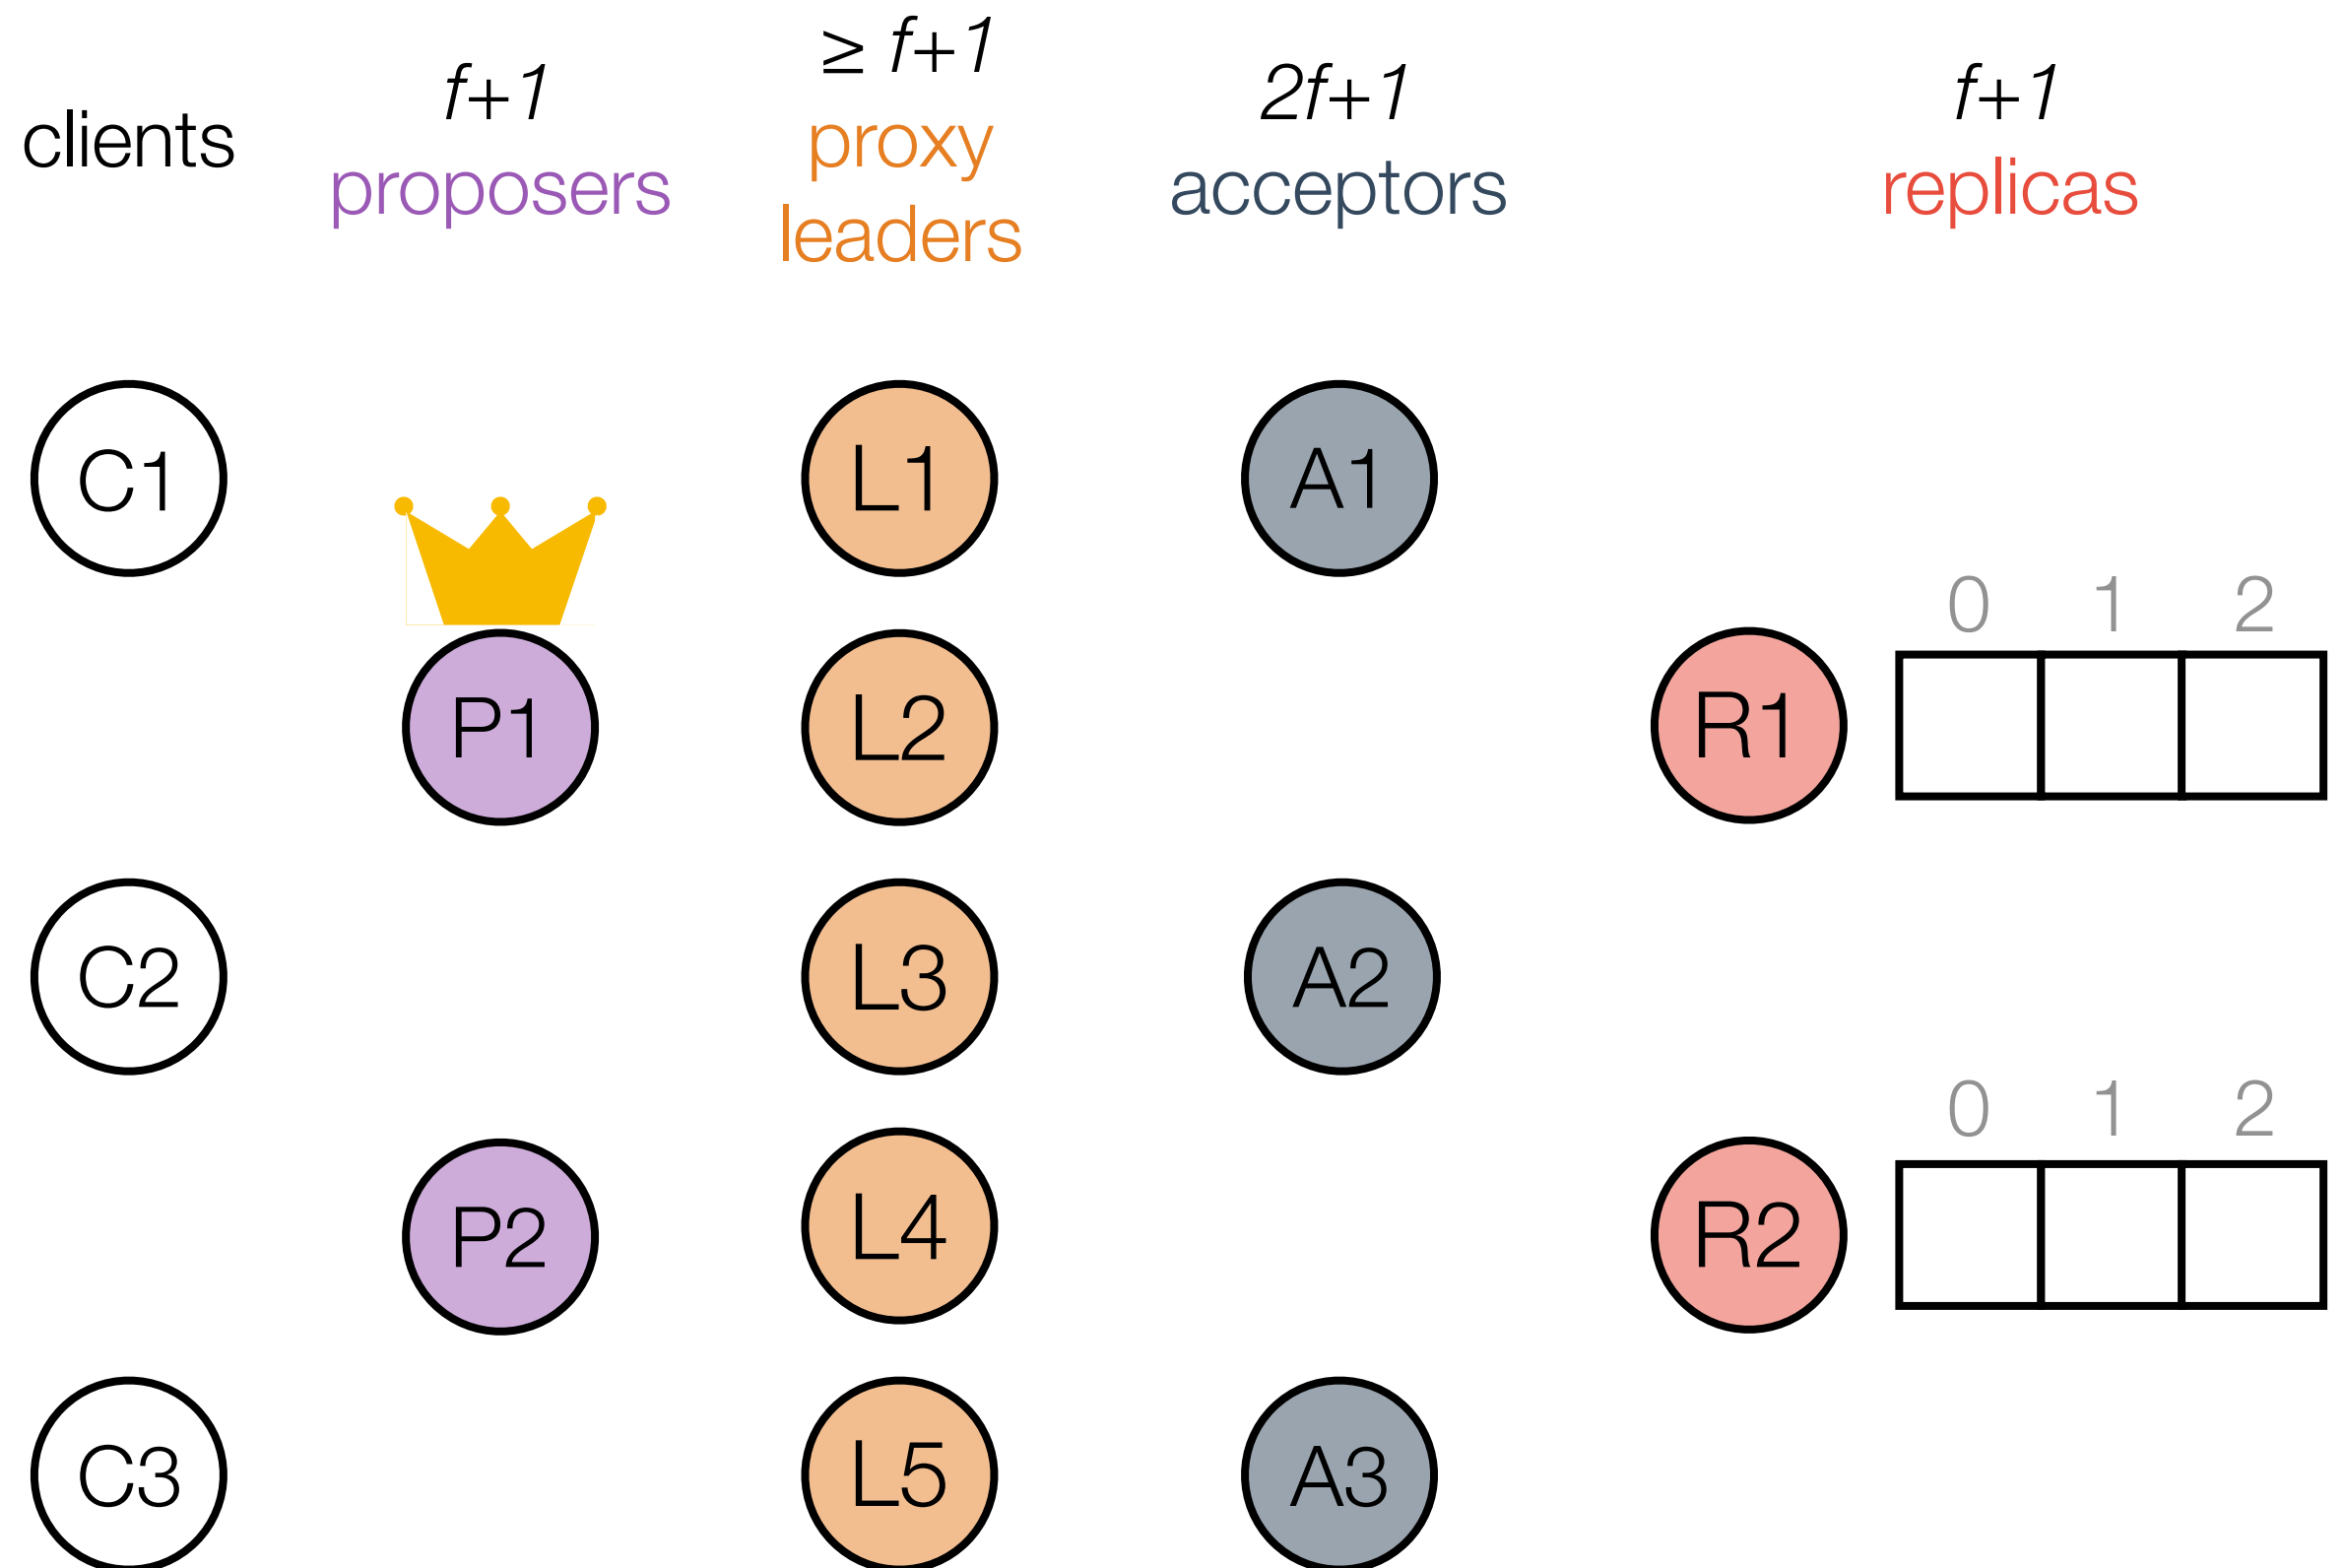
clients

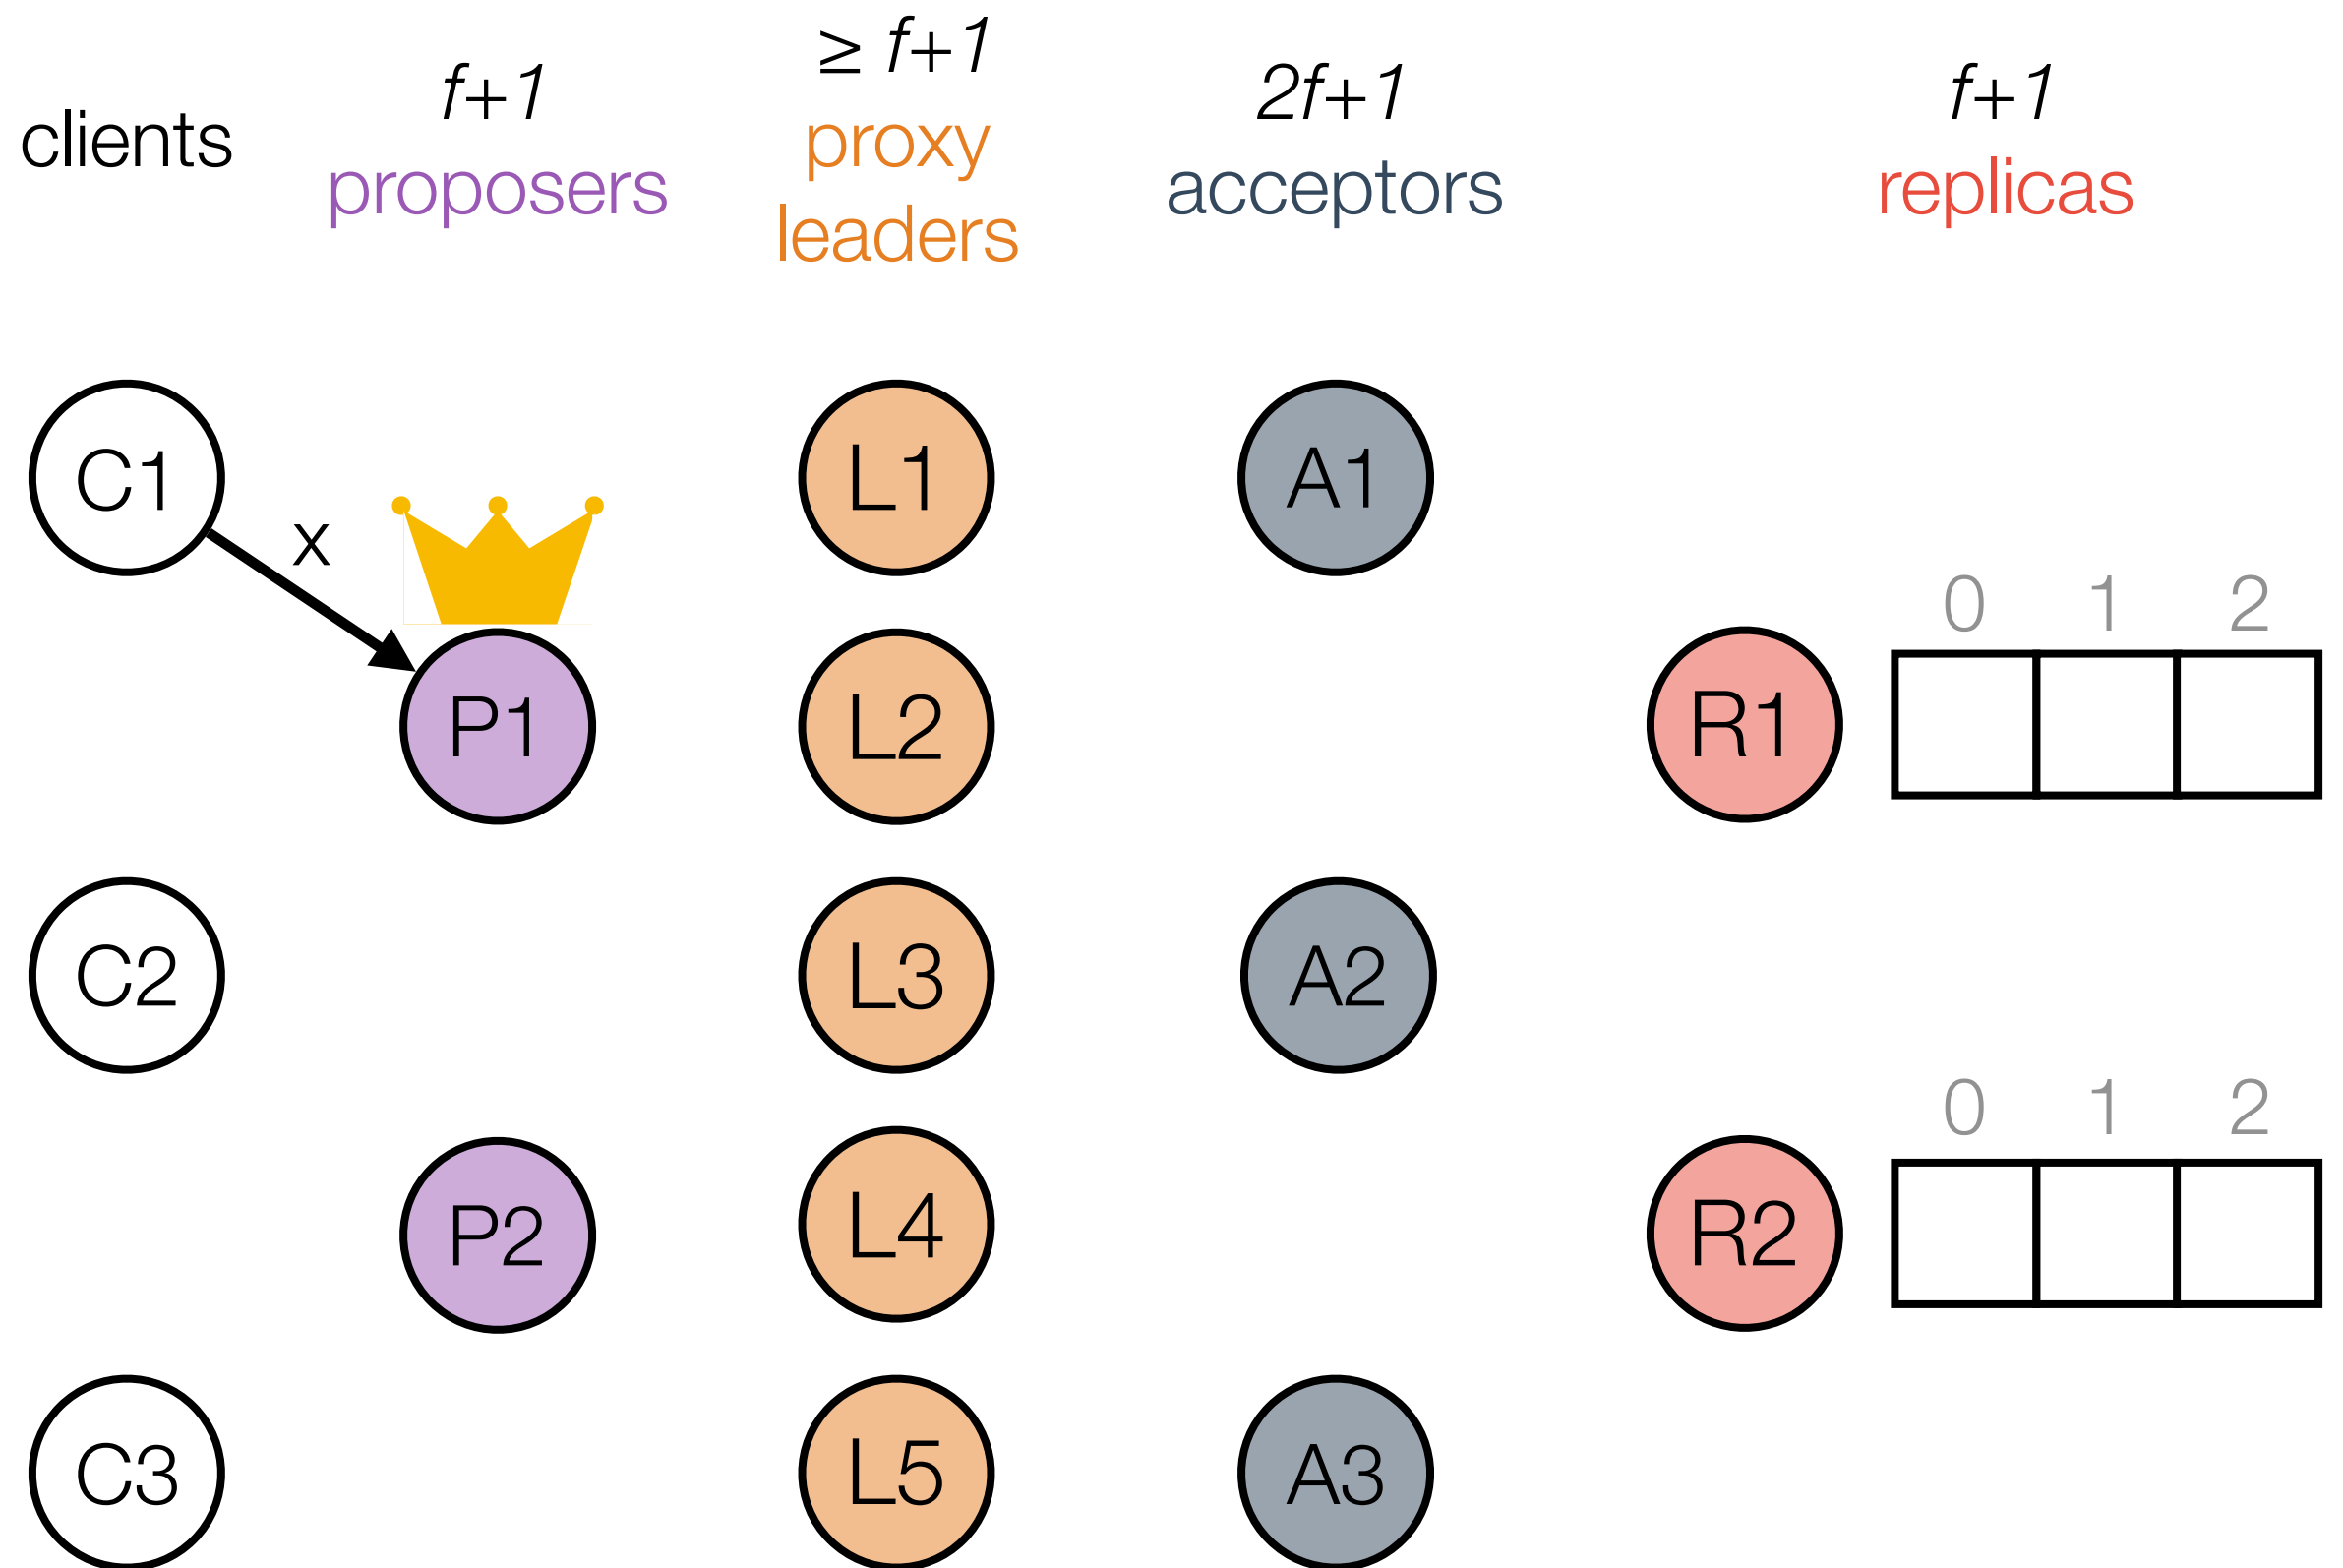
$f+1$
proposers

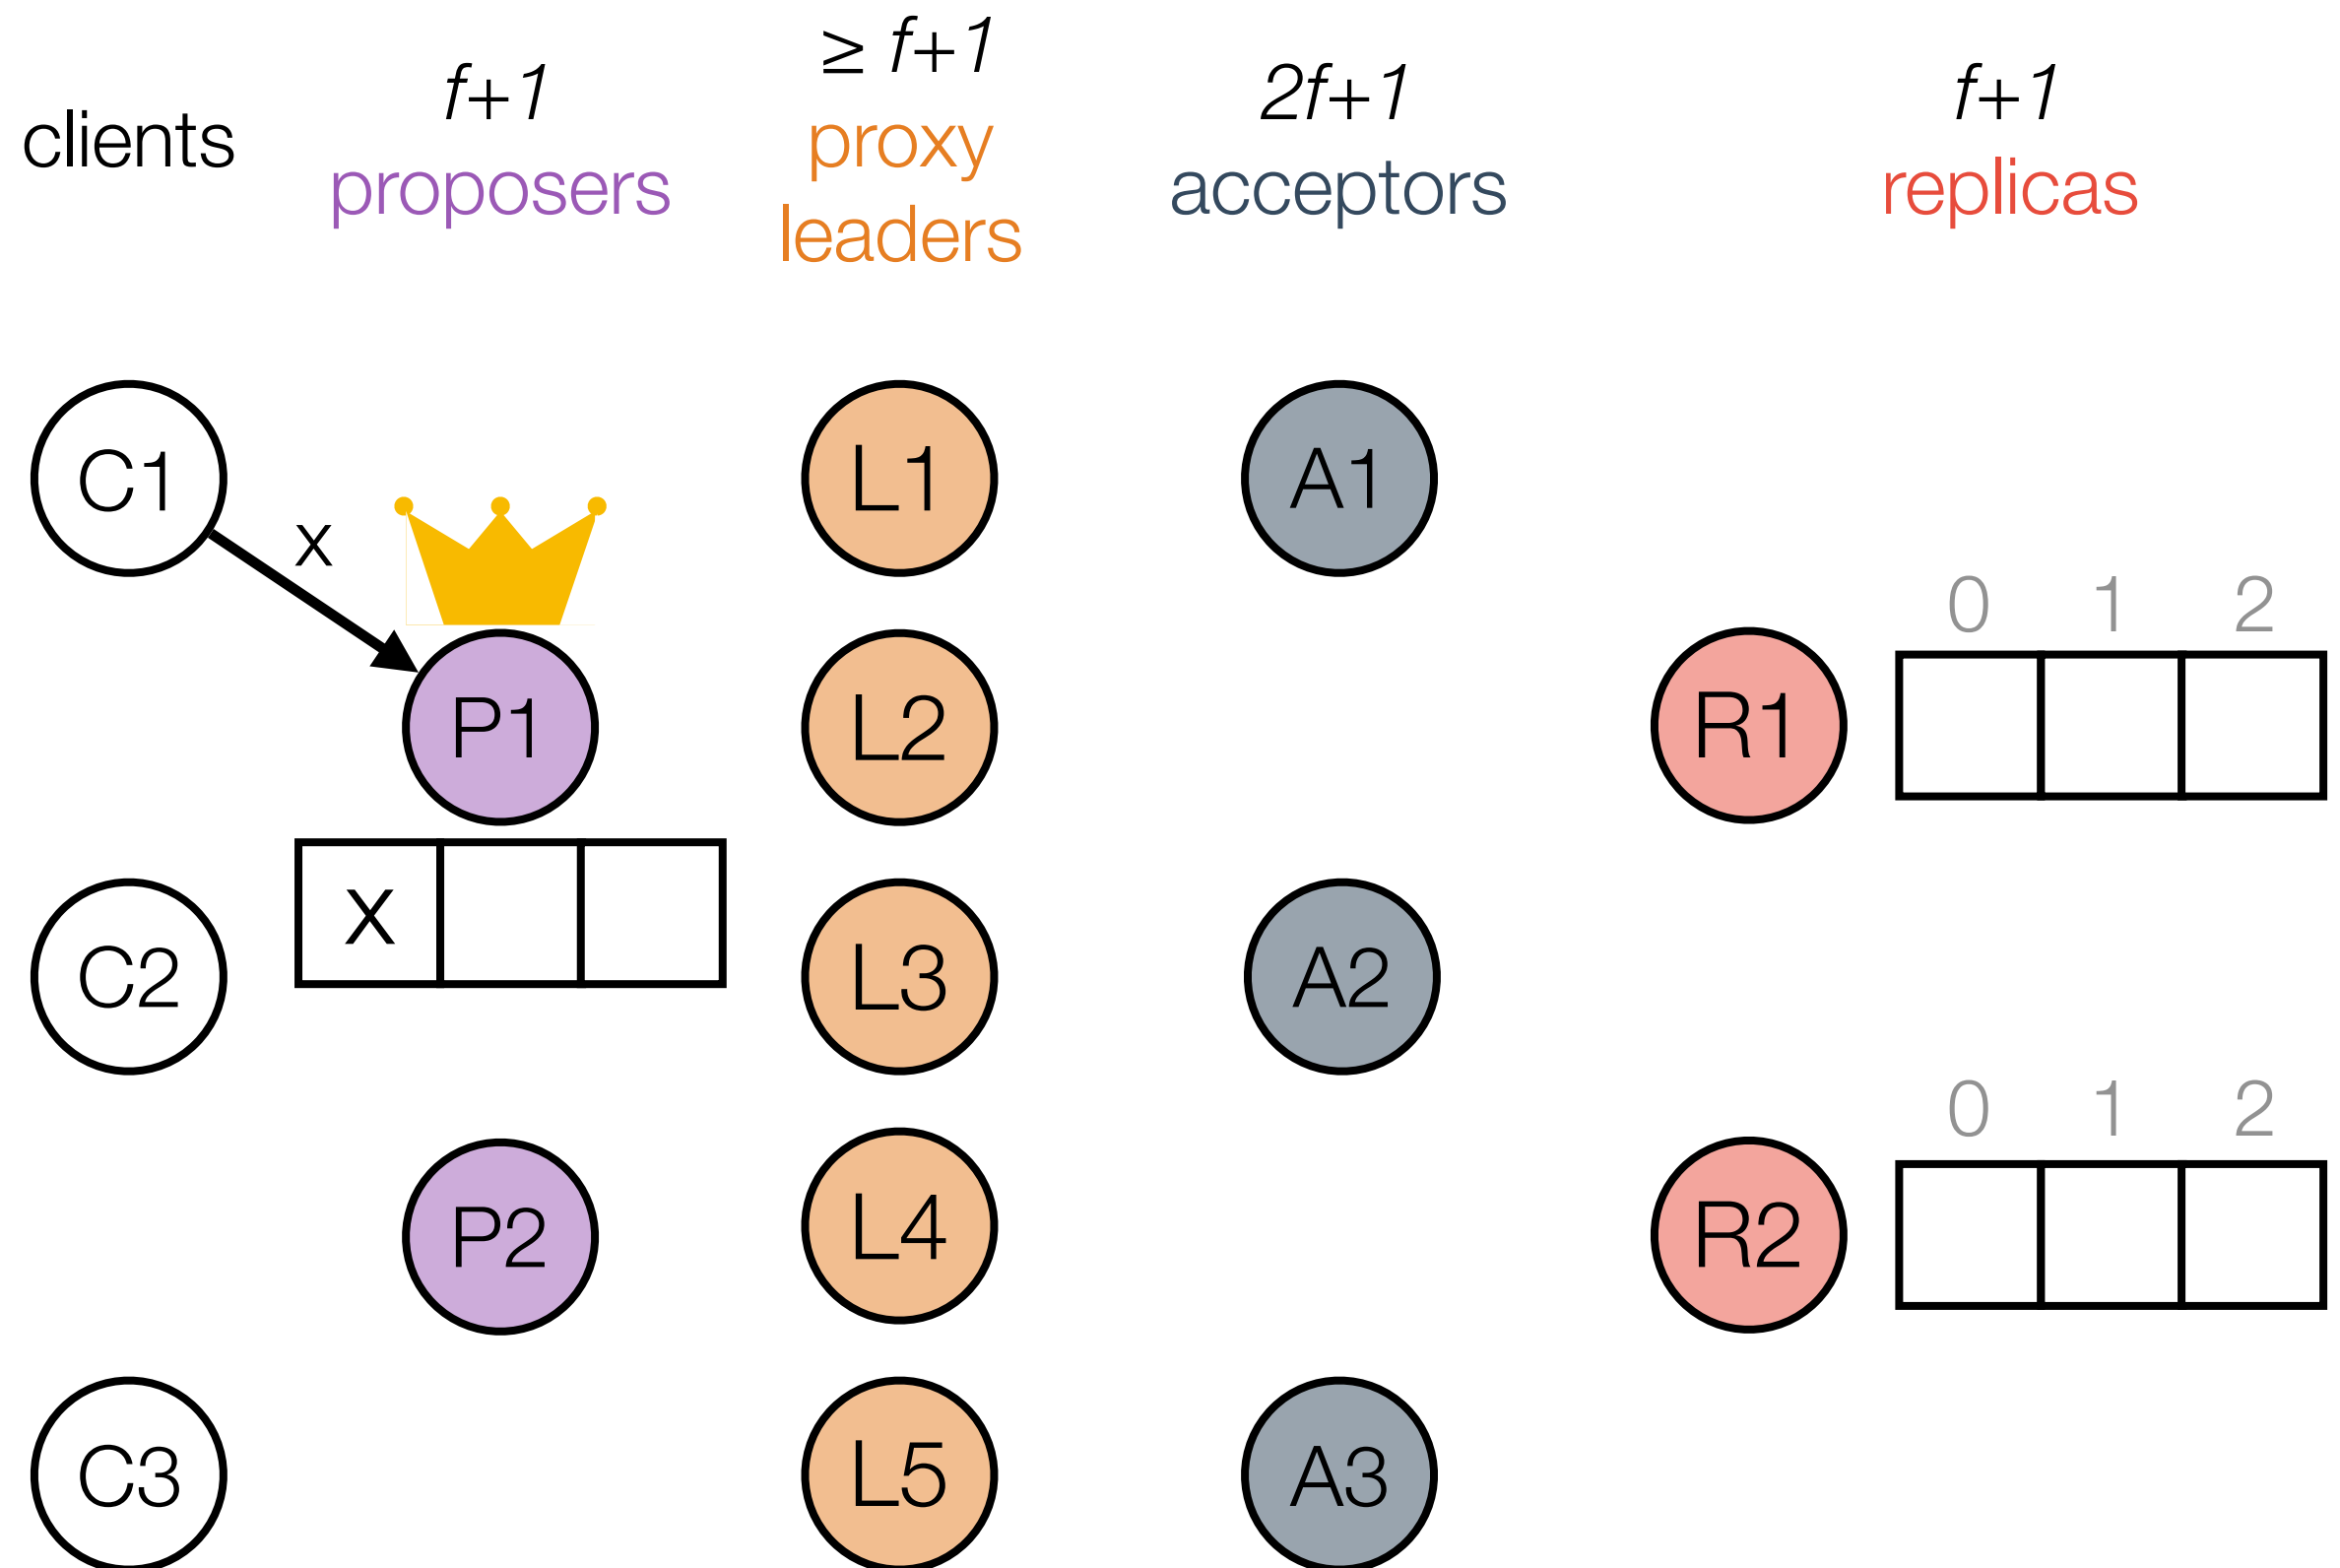
$2f+1$
acceptors

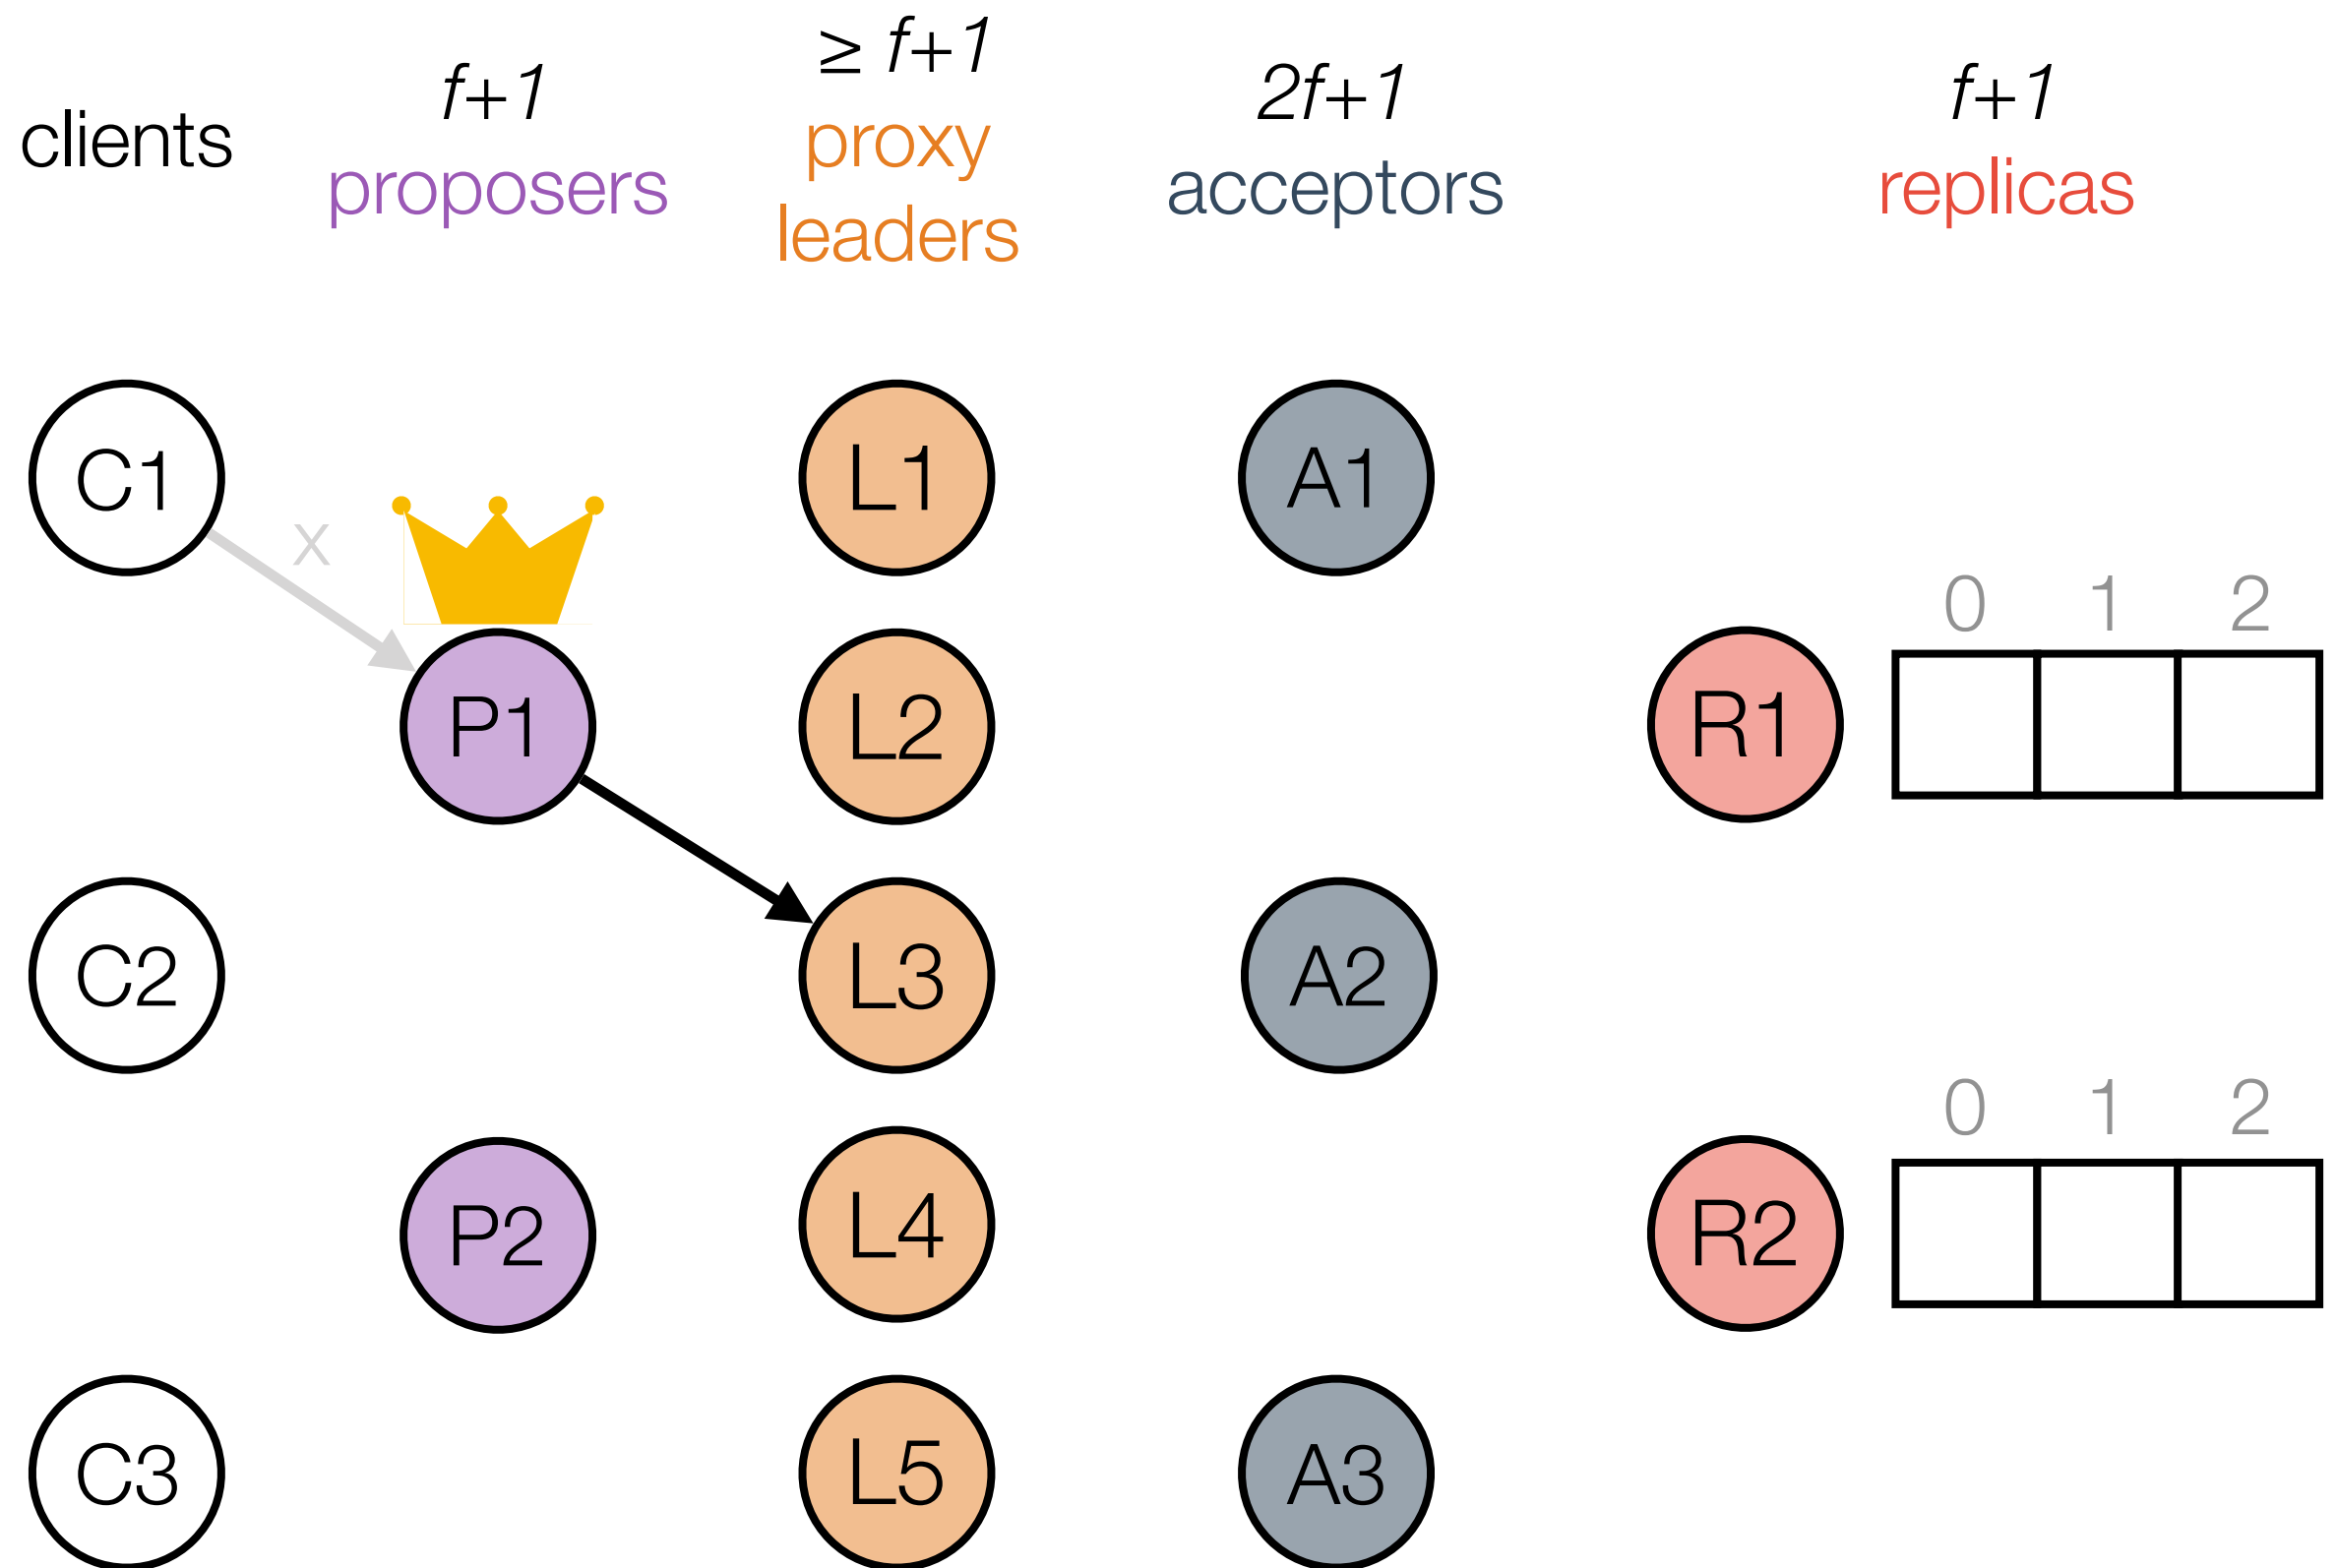
$f+1$
replicas

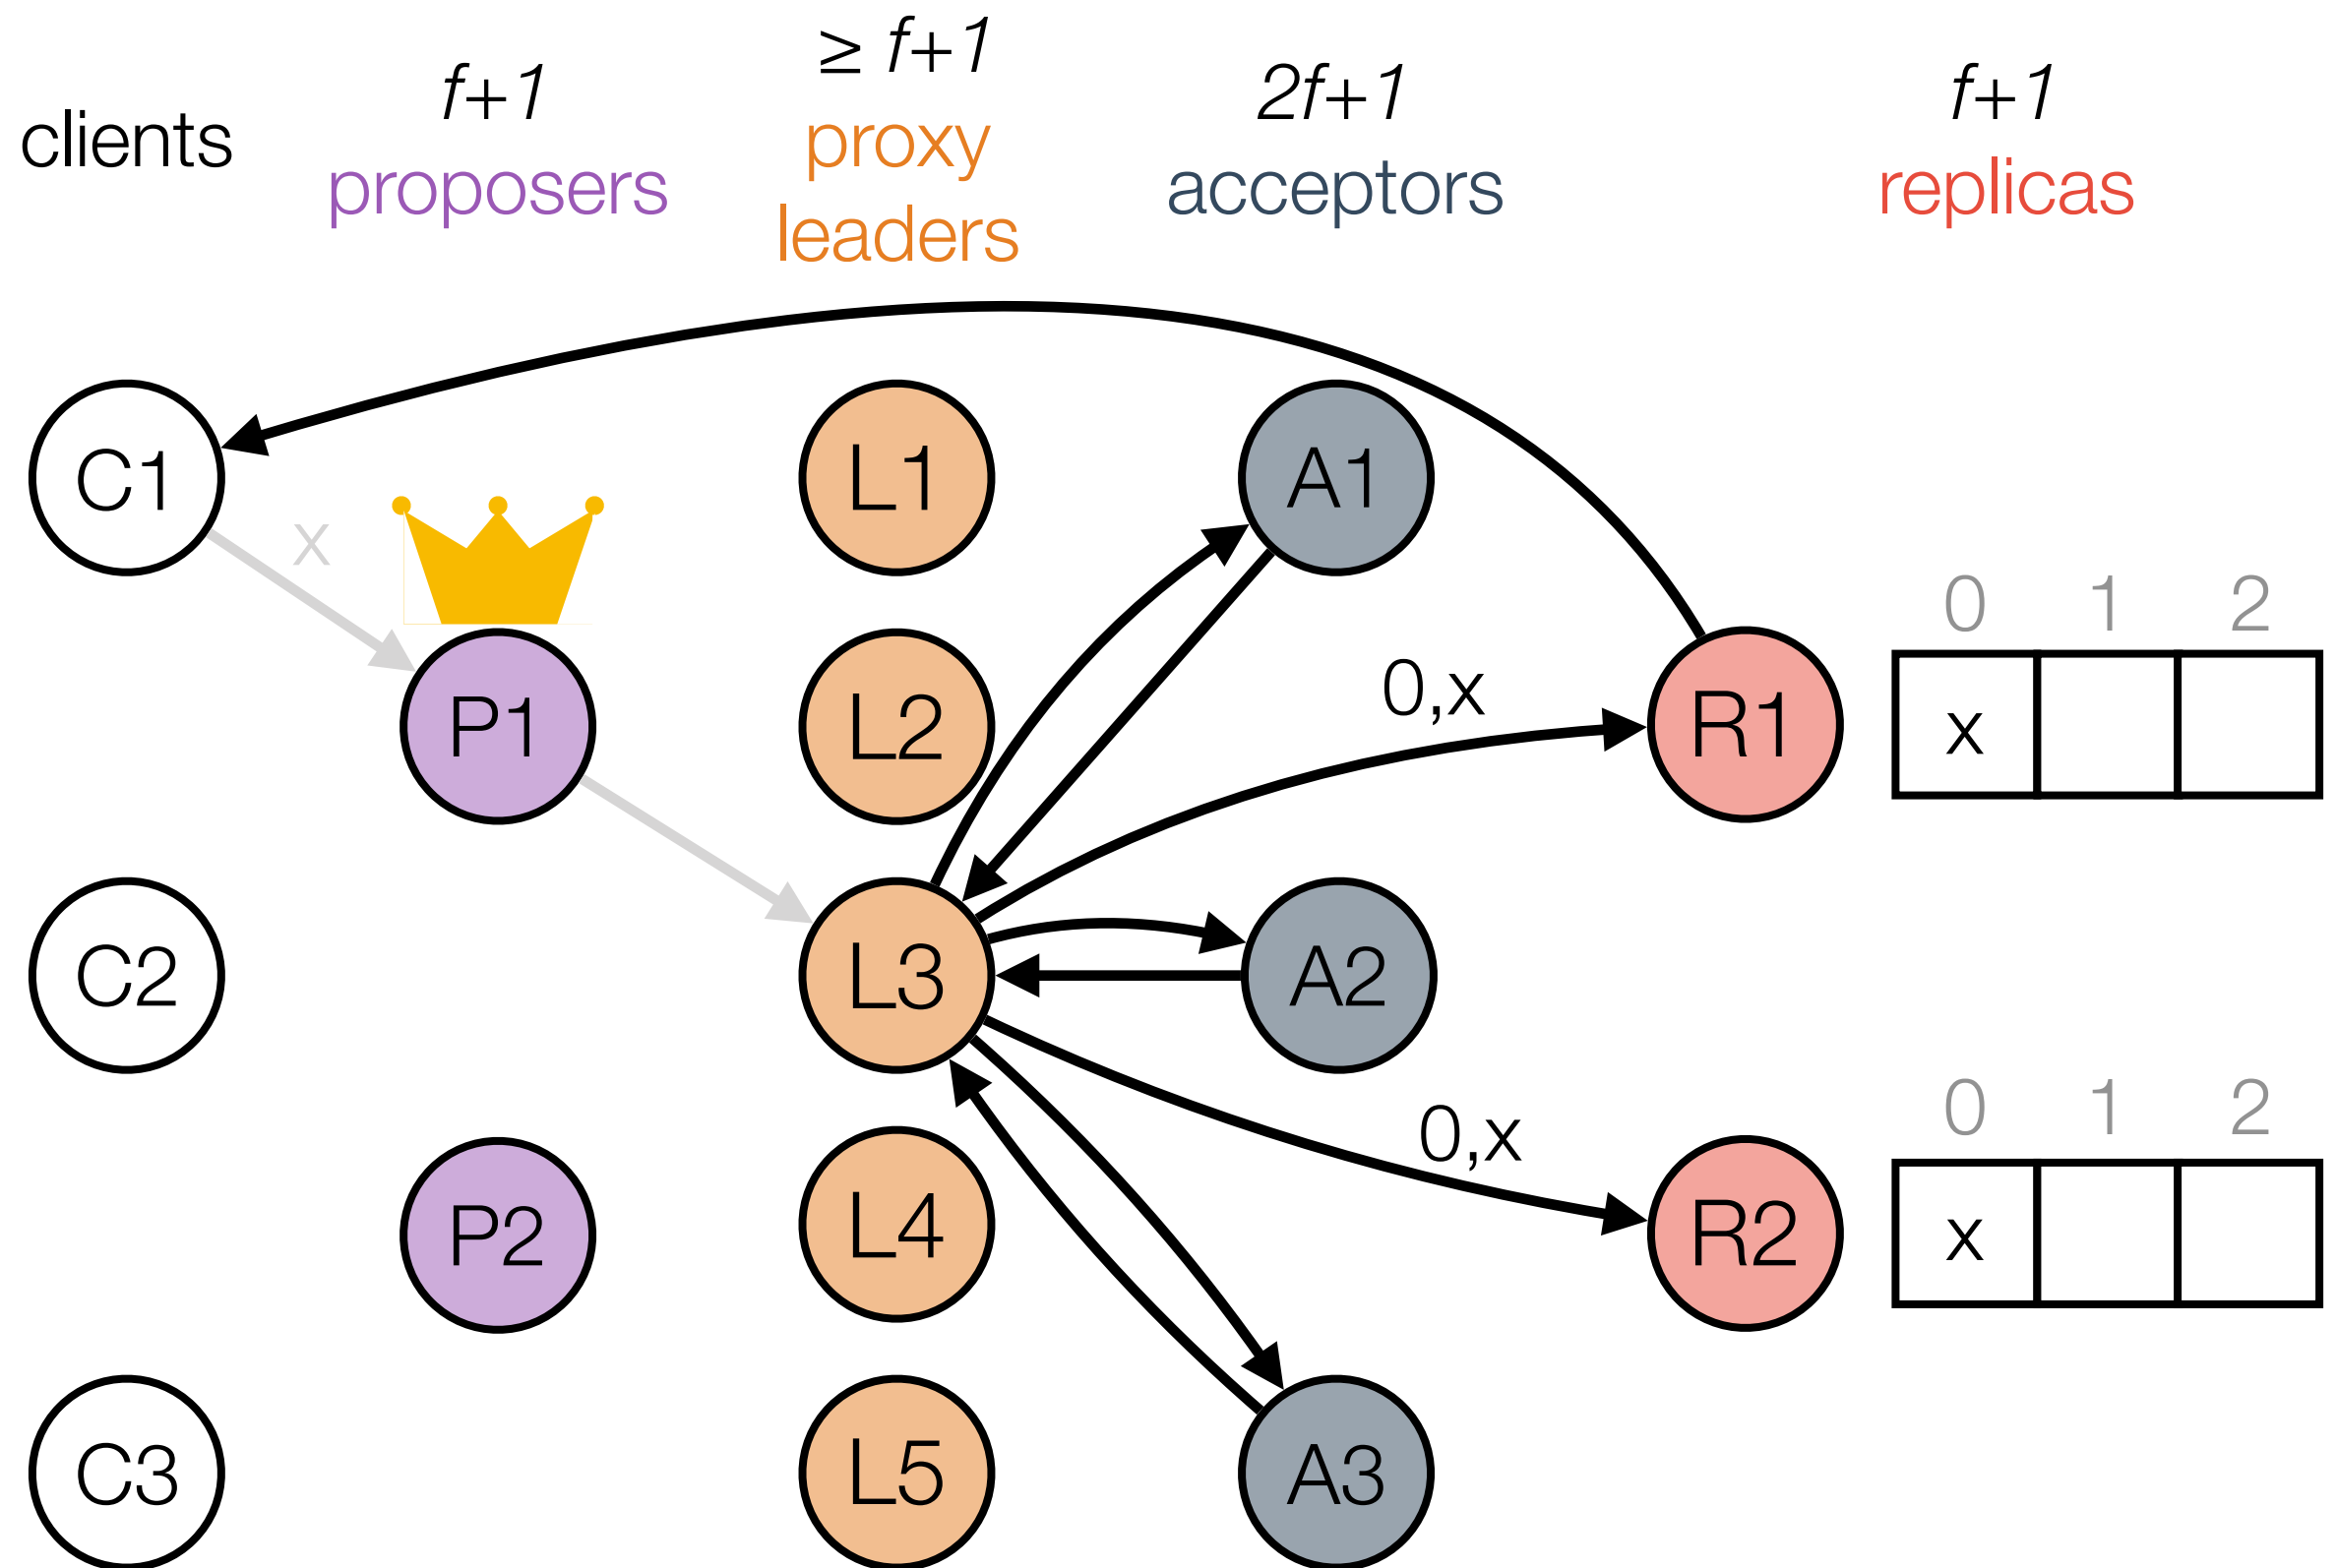




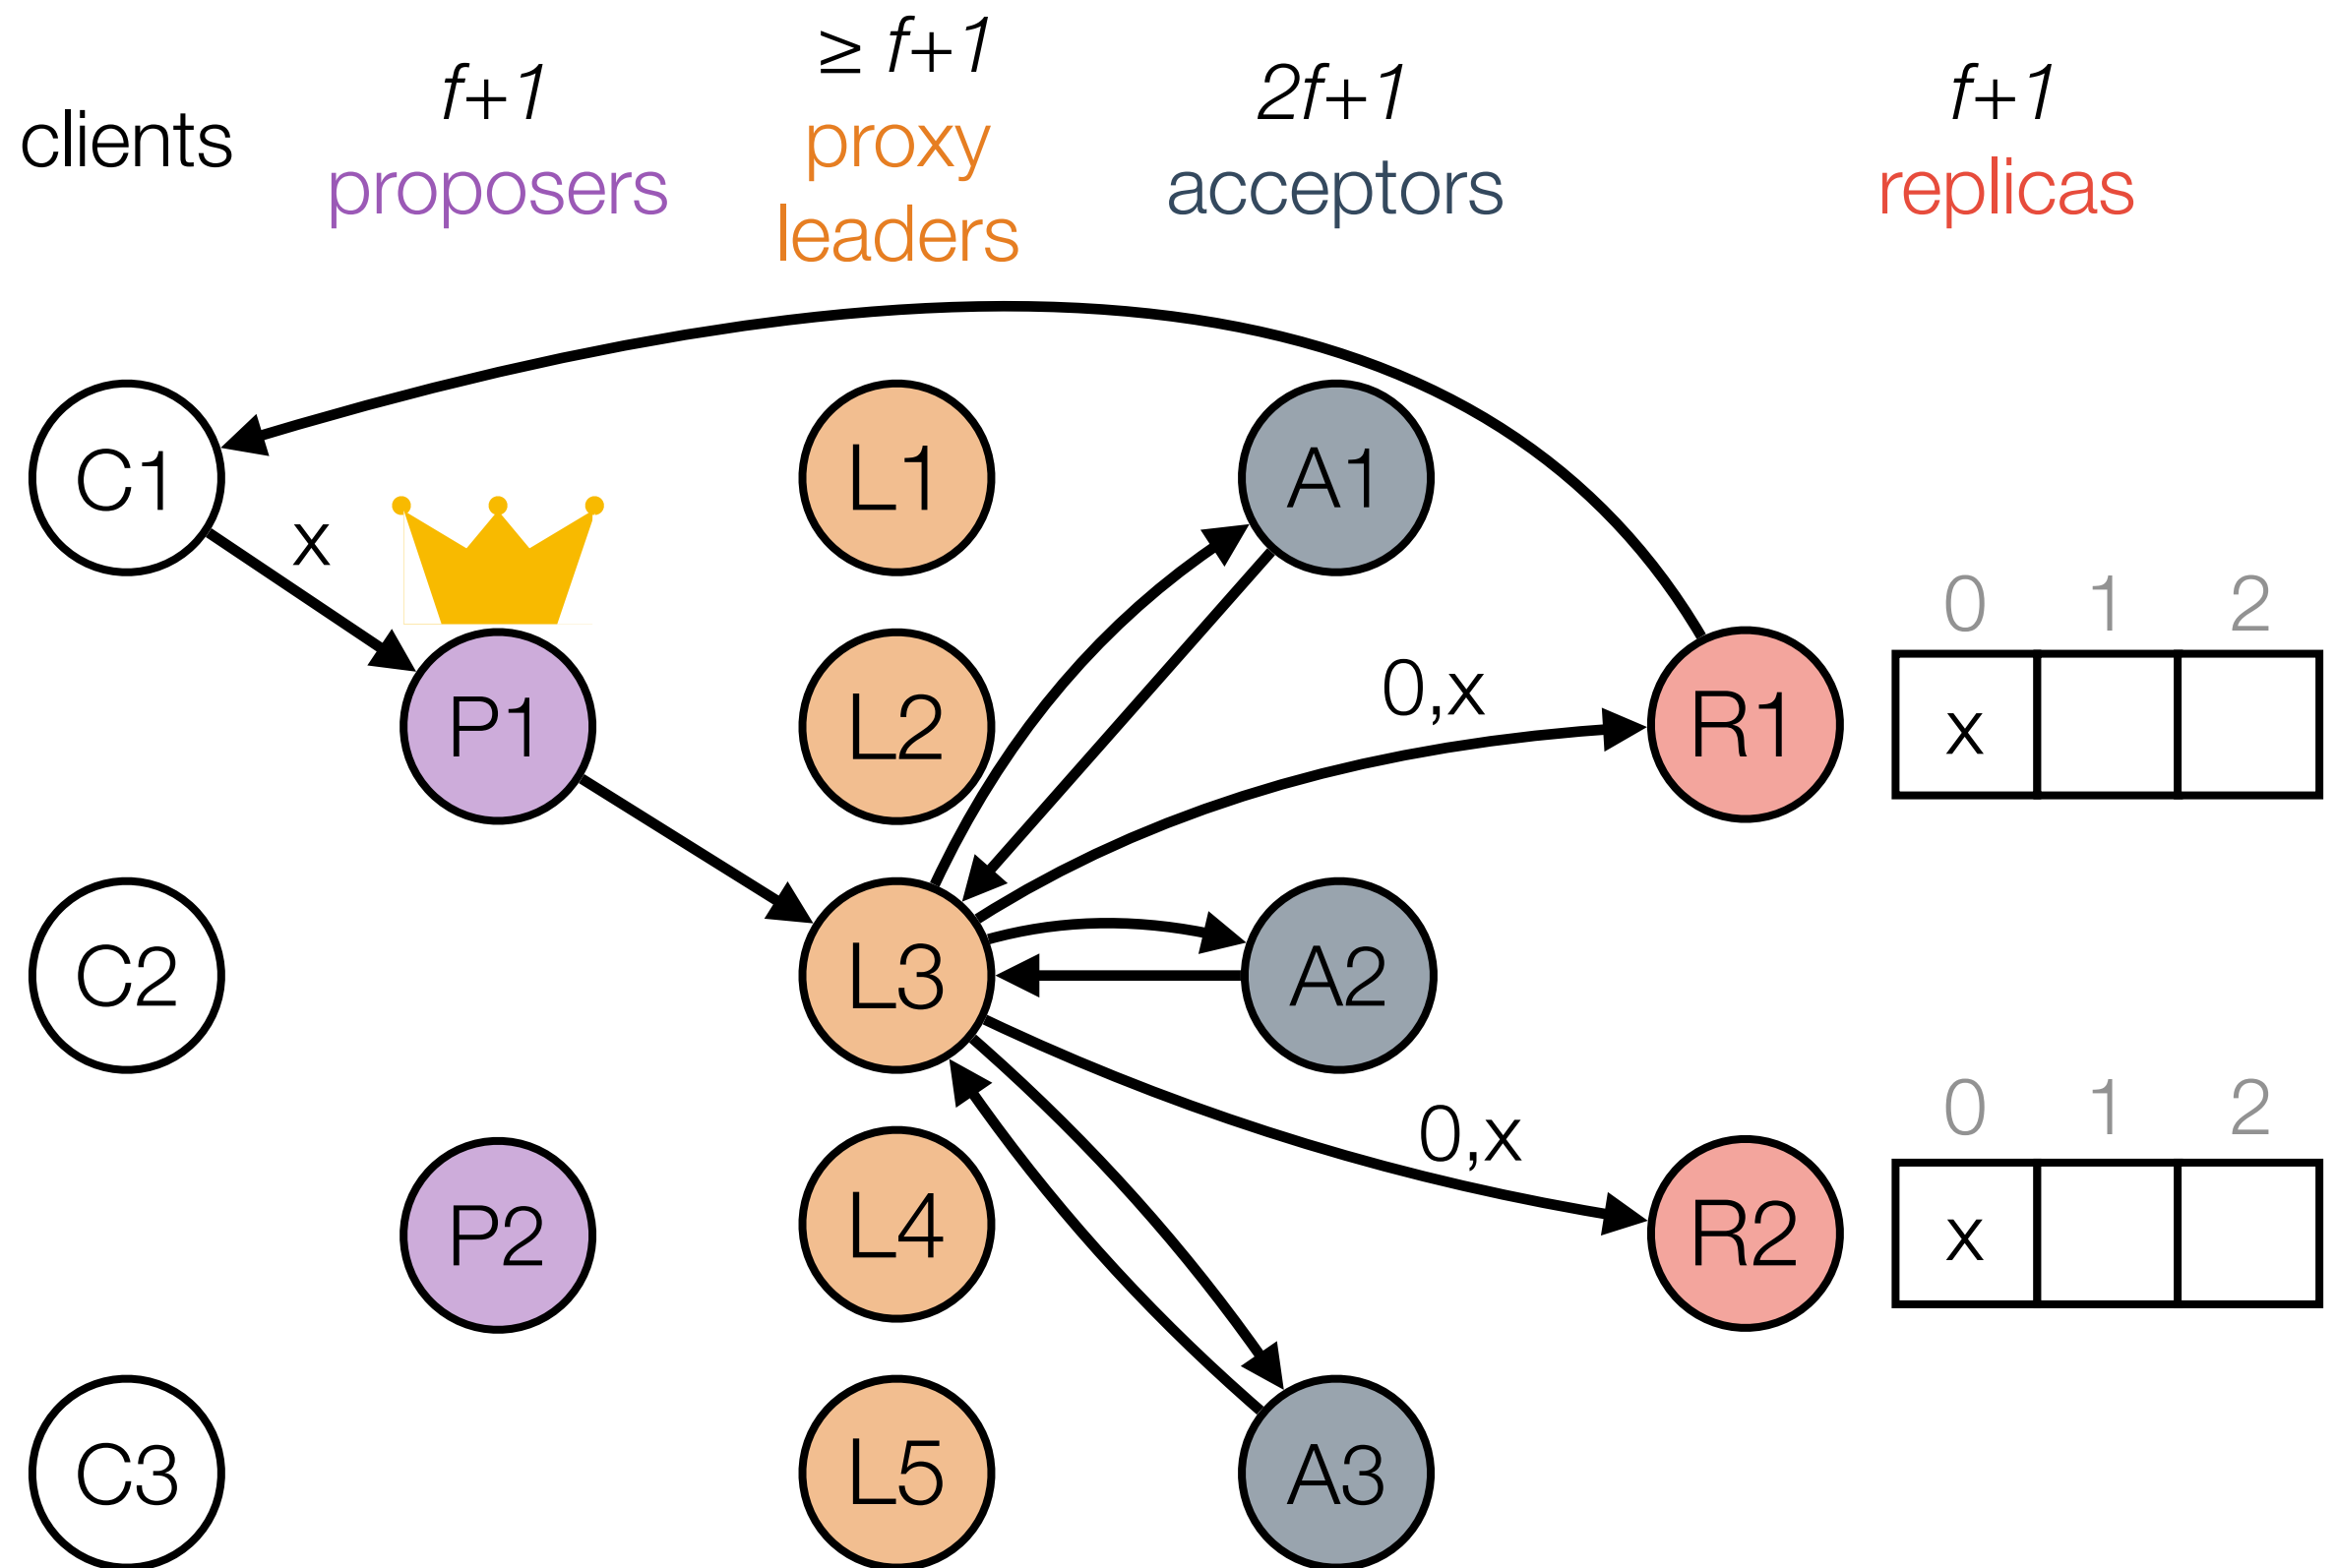


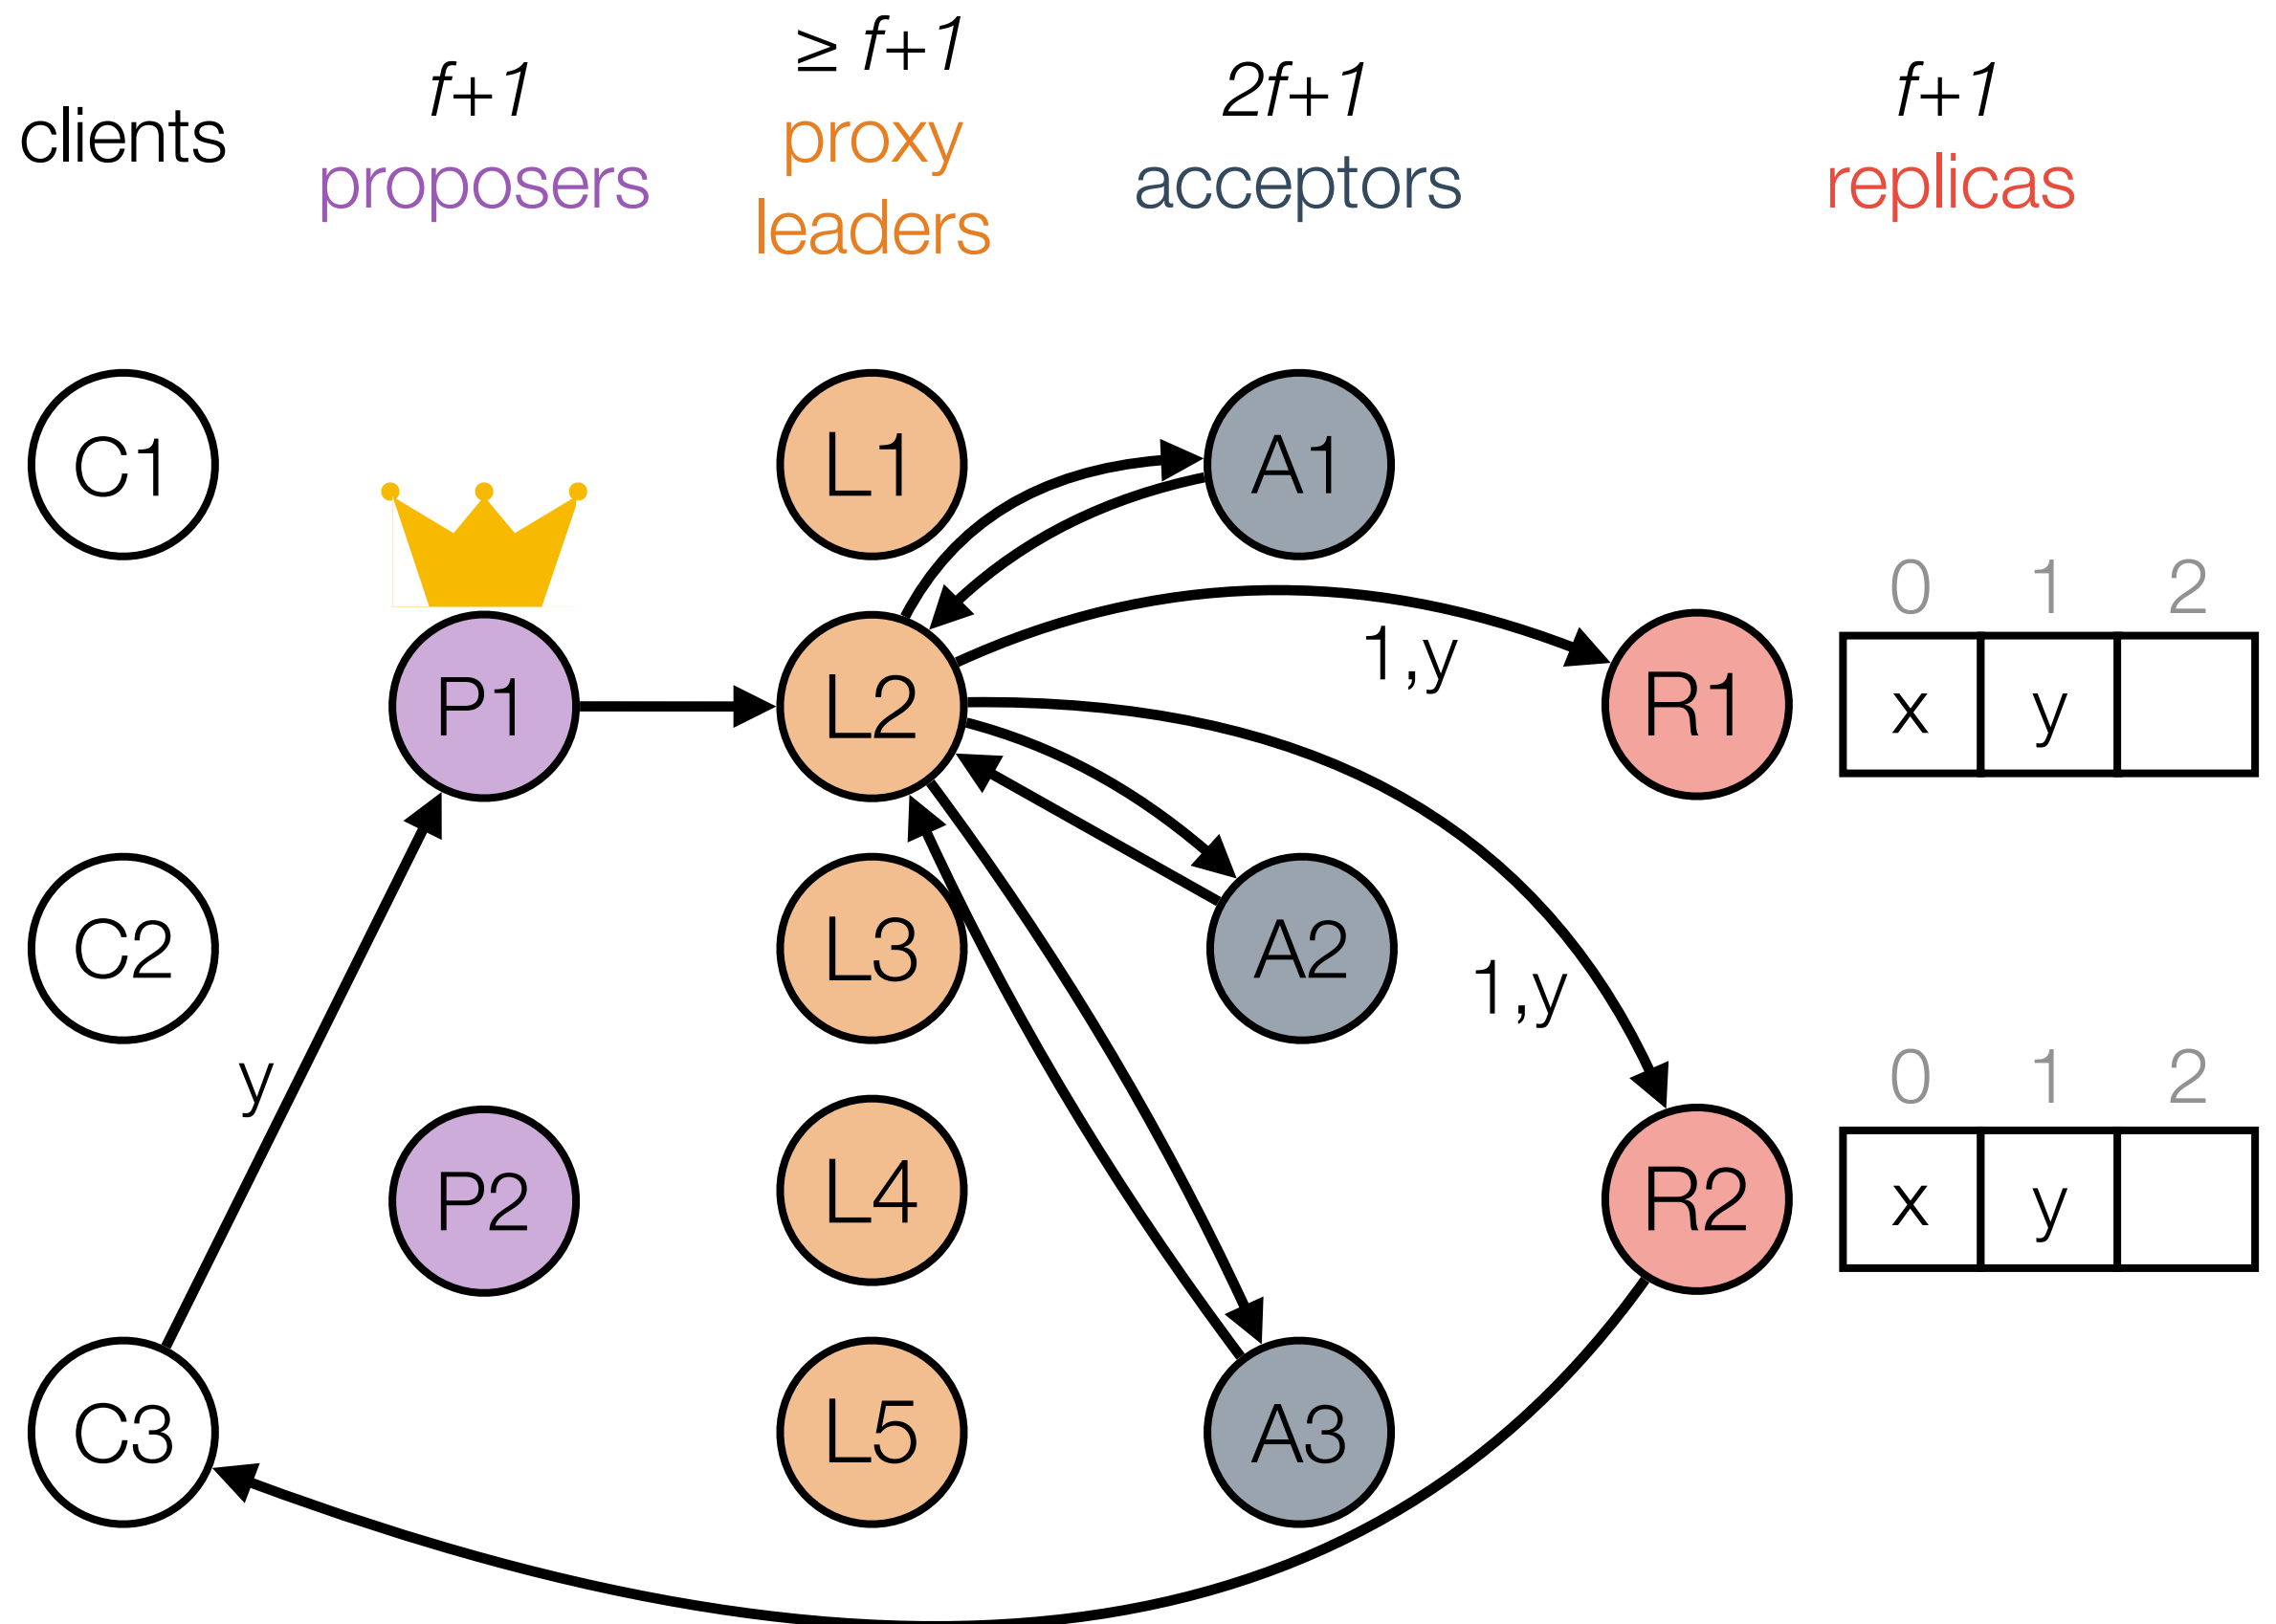


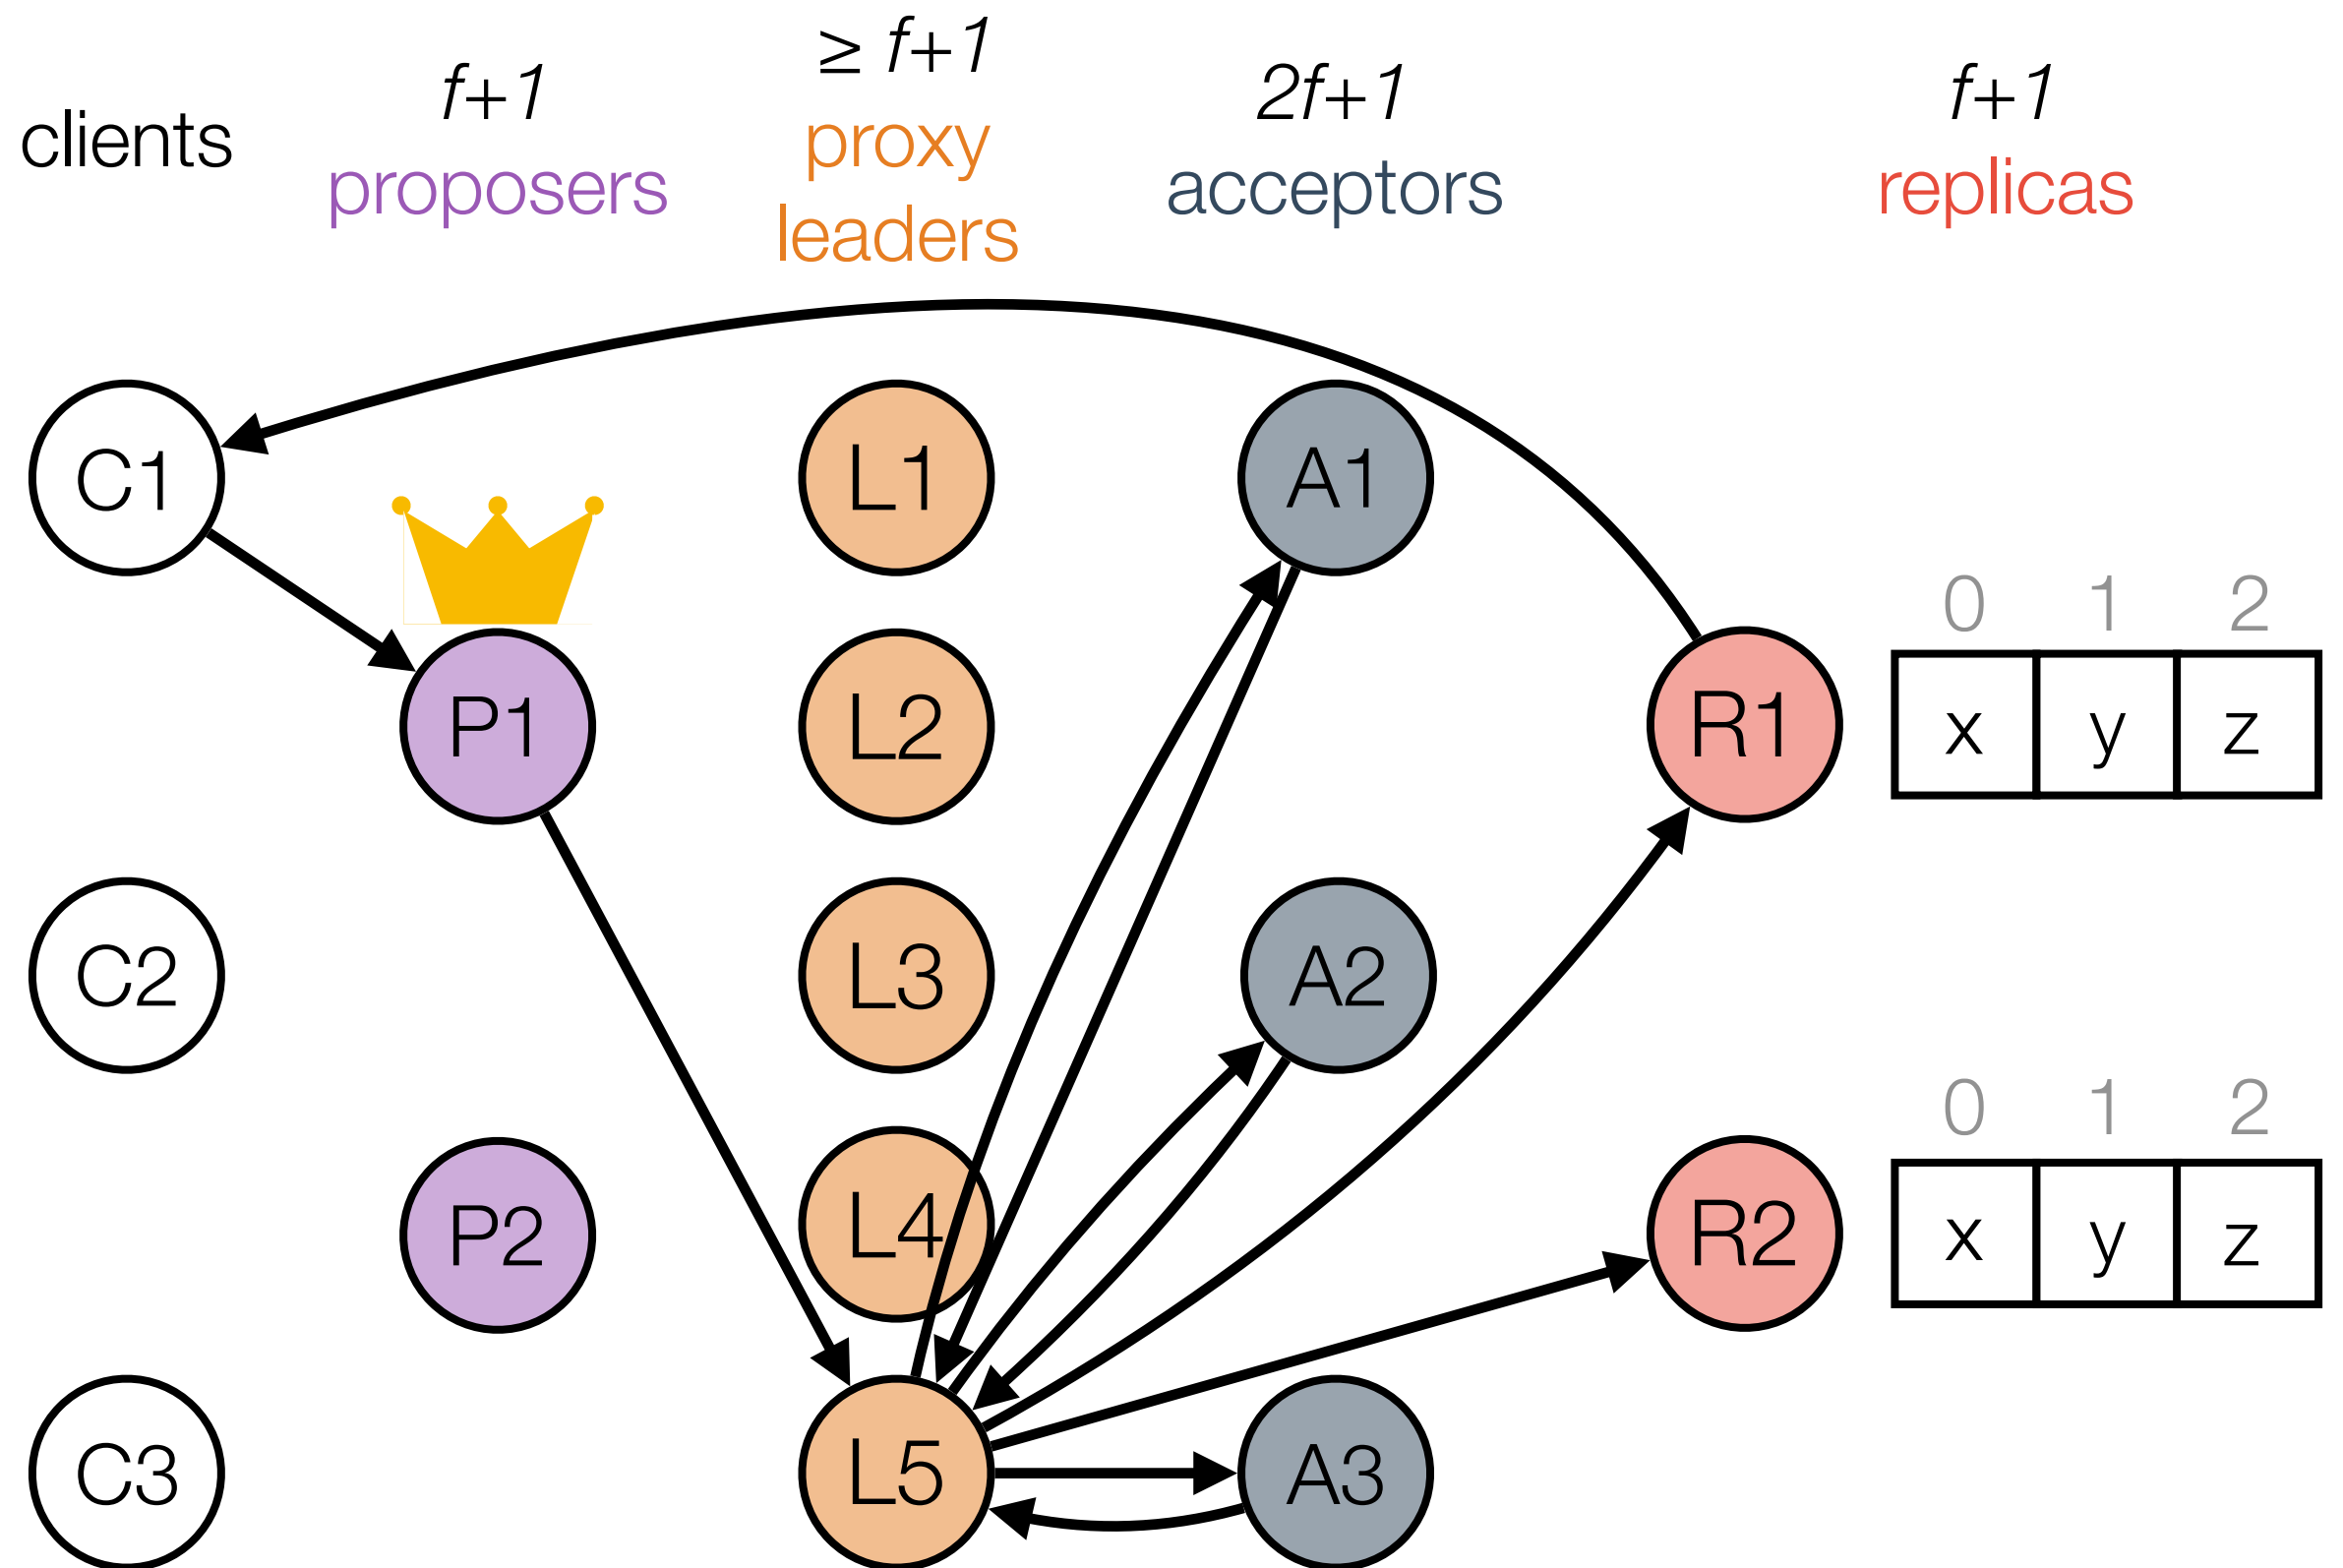


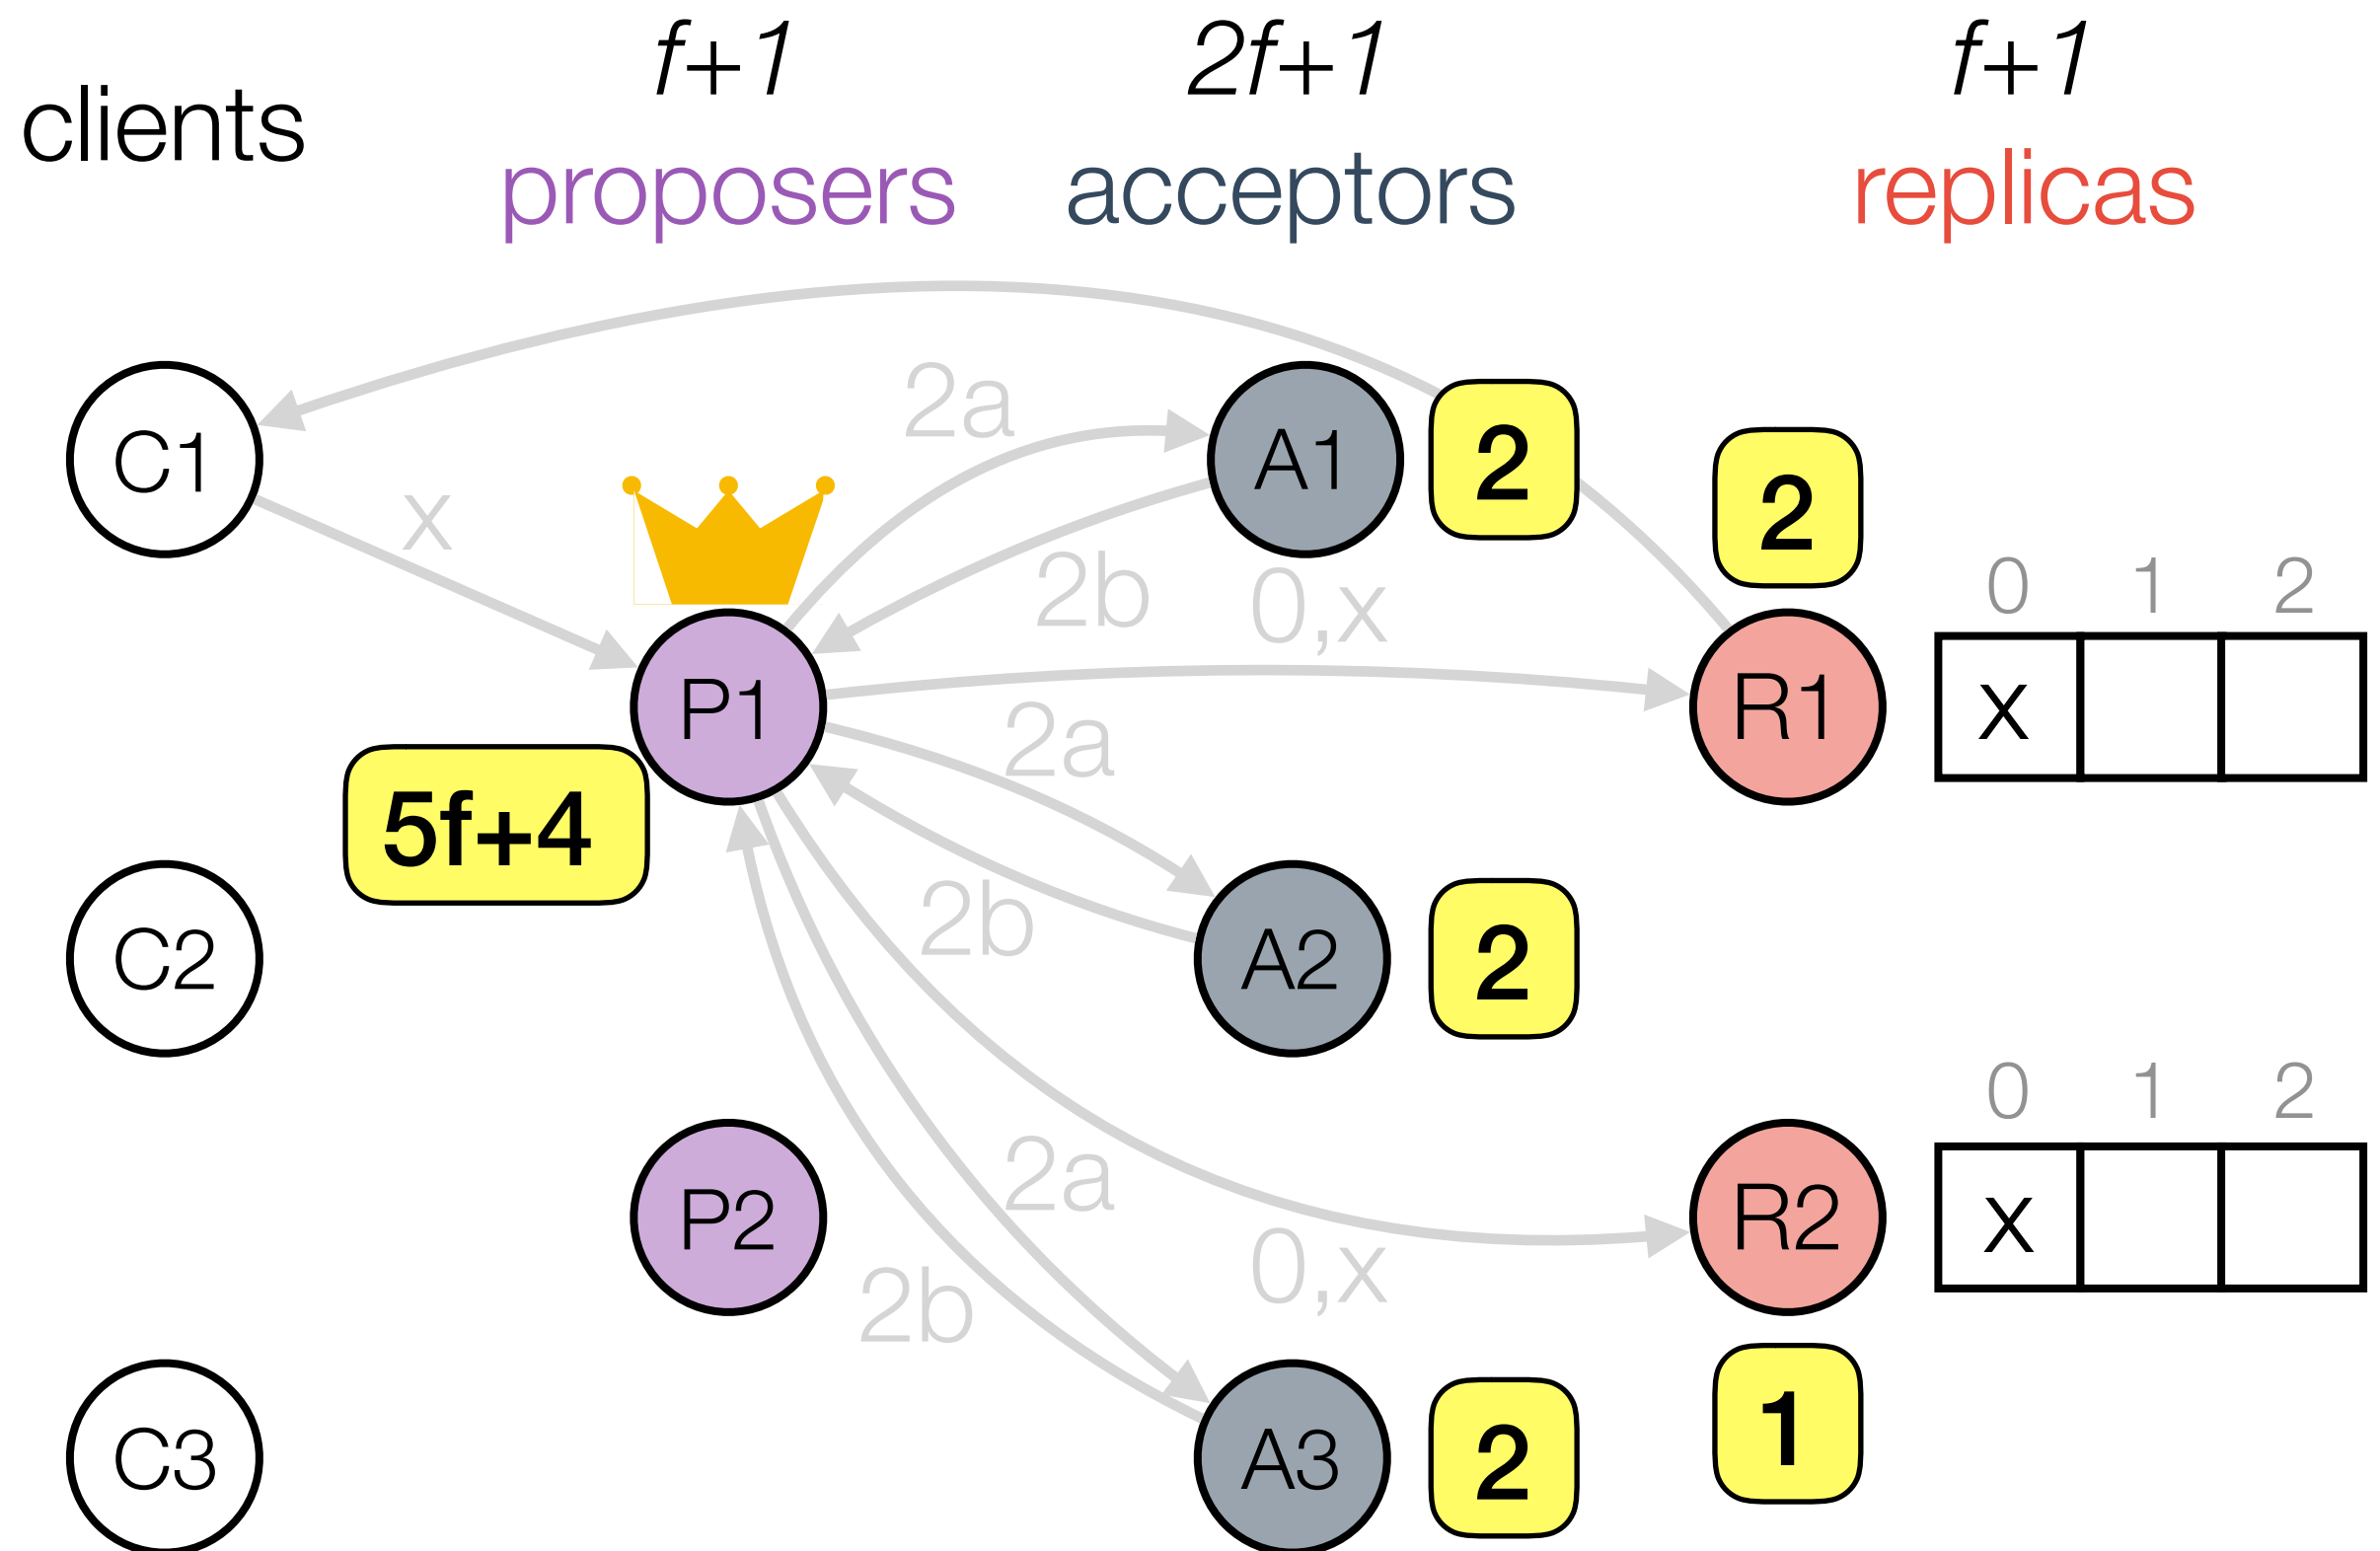


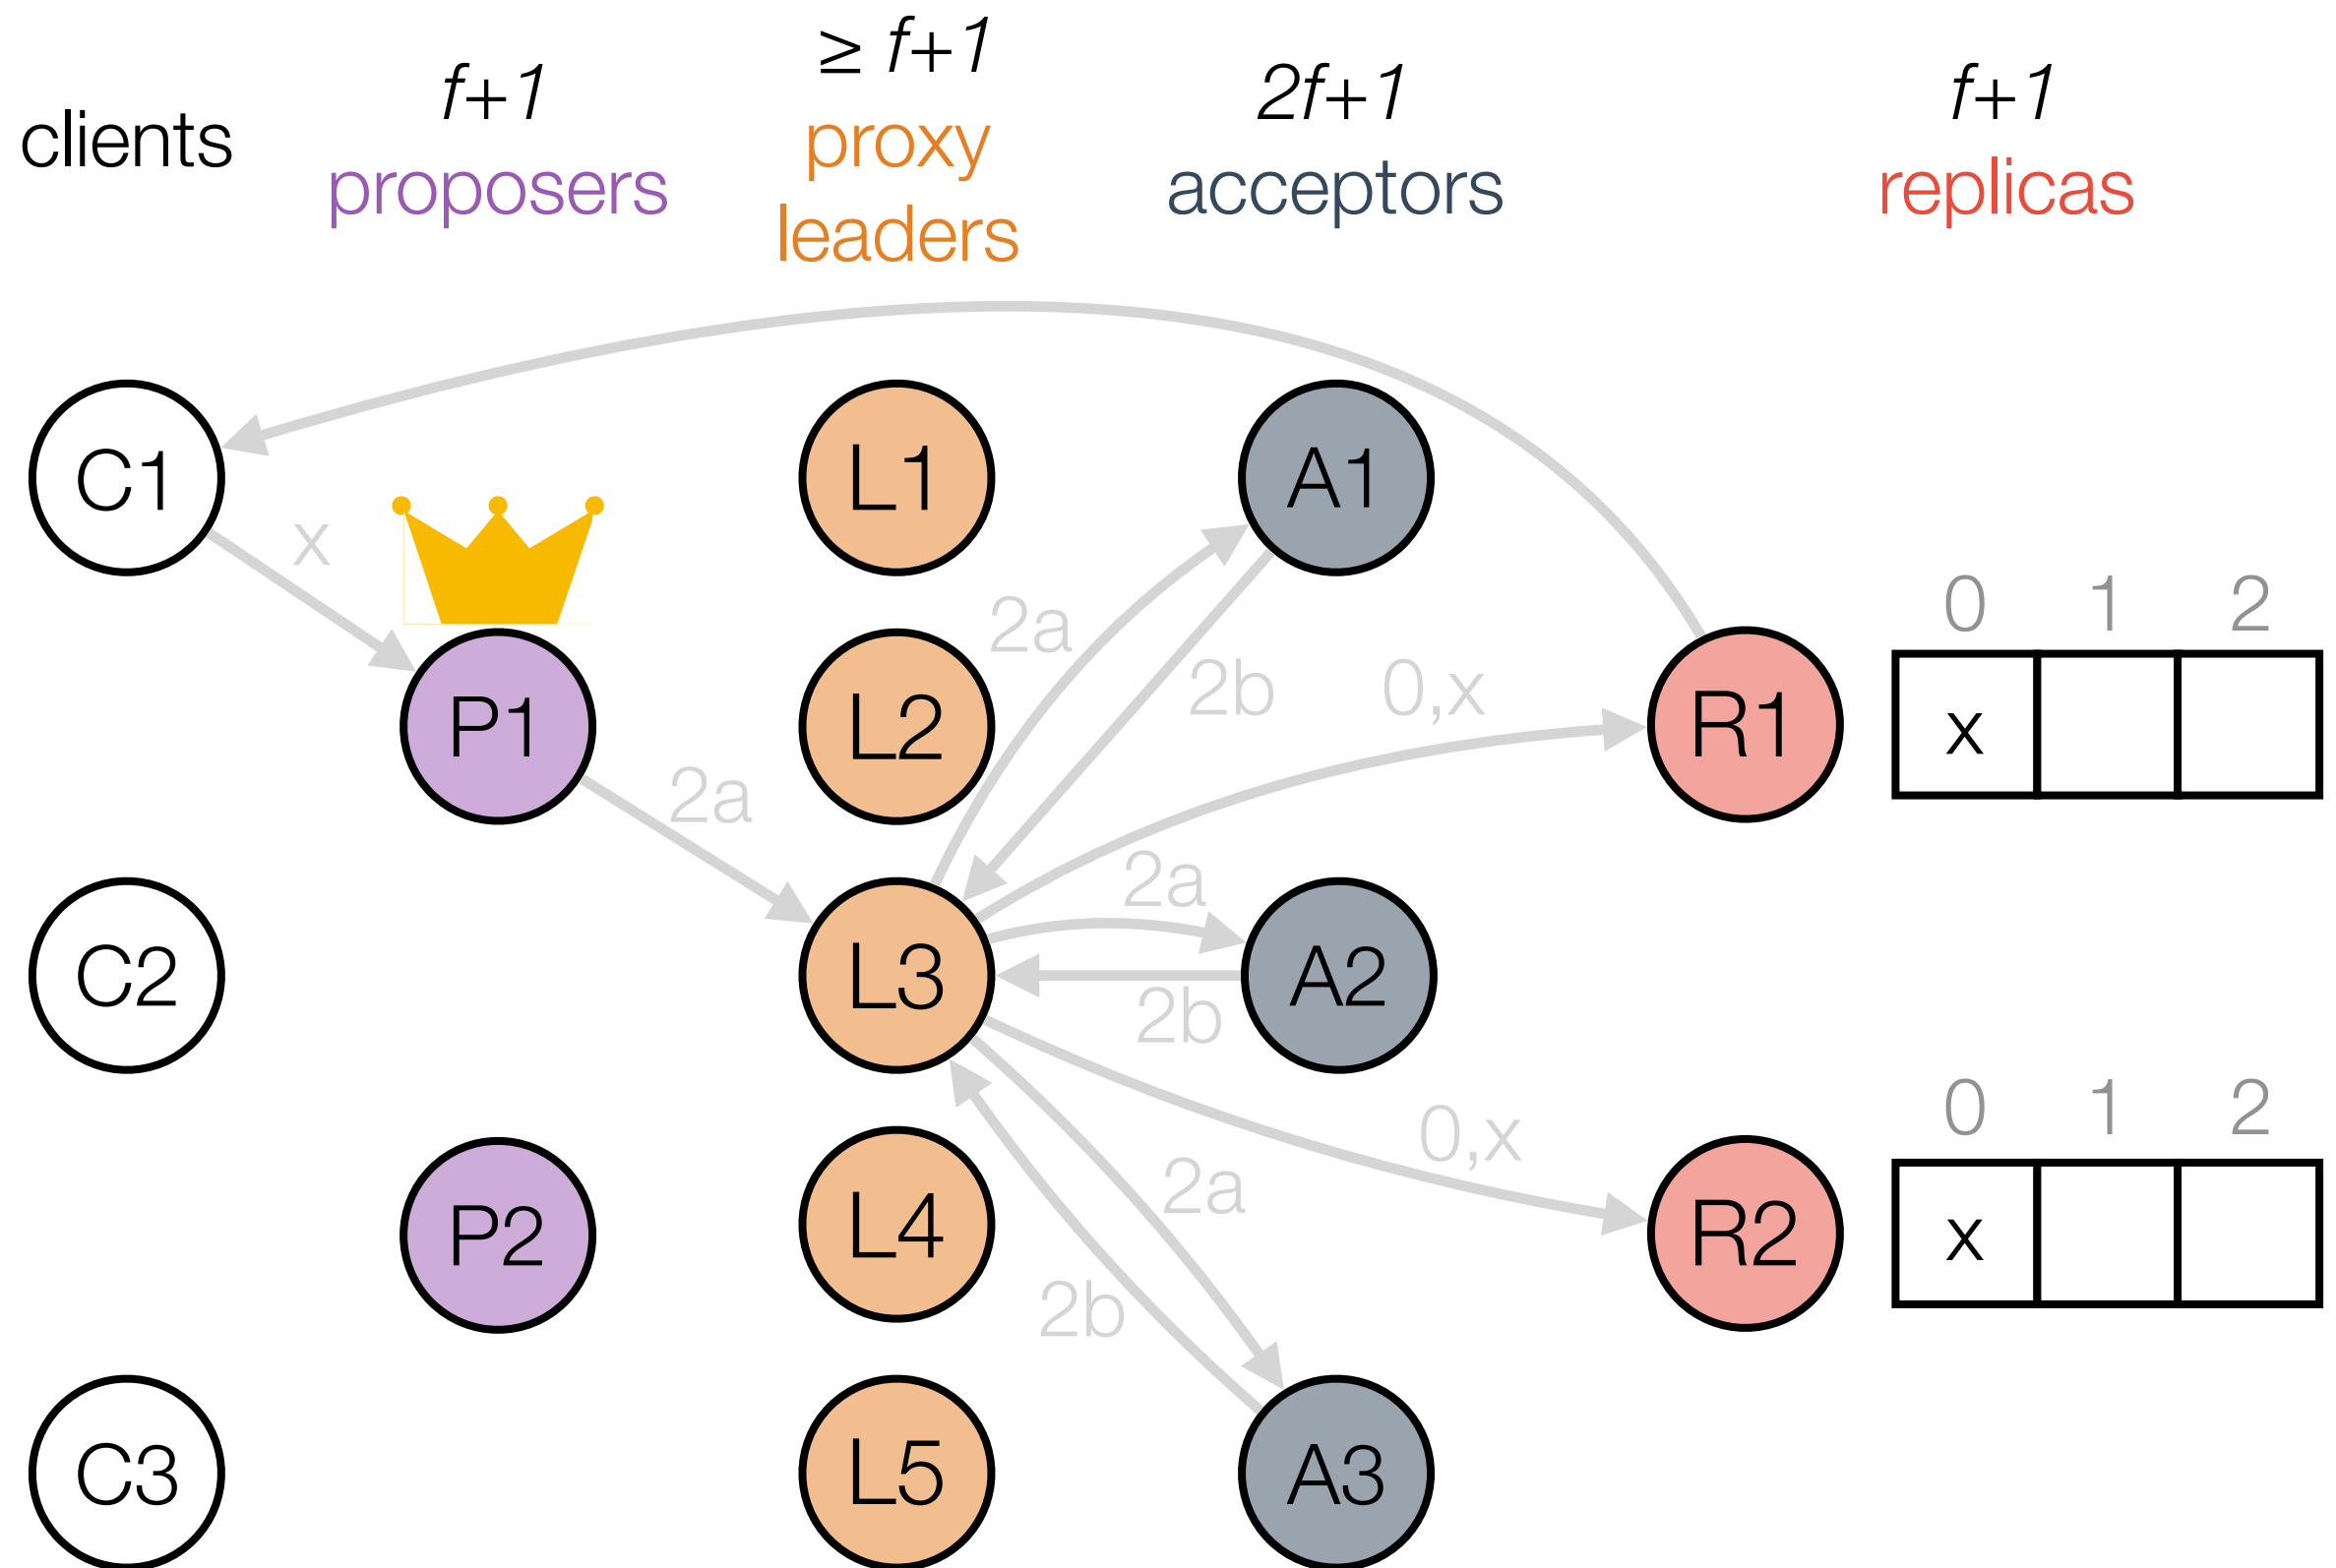
Key Insight: **Scale up**
the proxy leaders

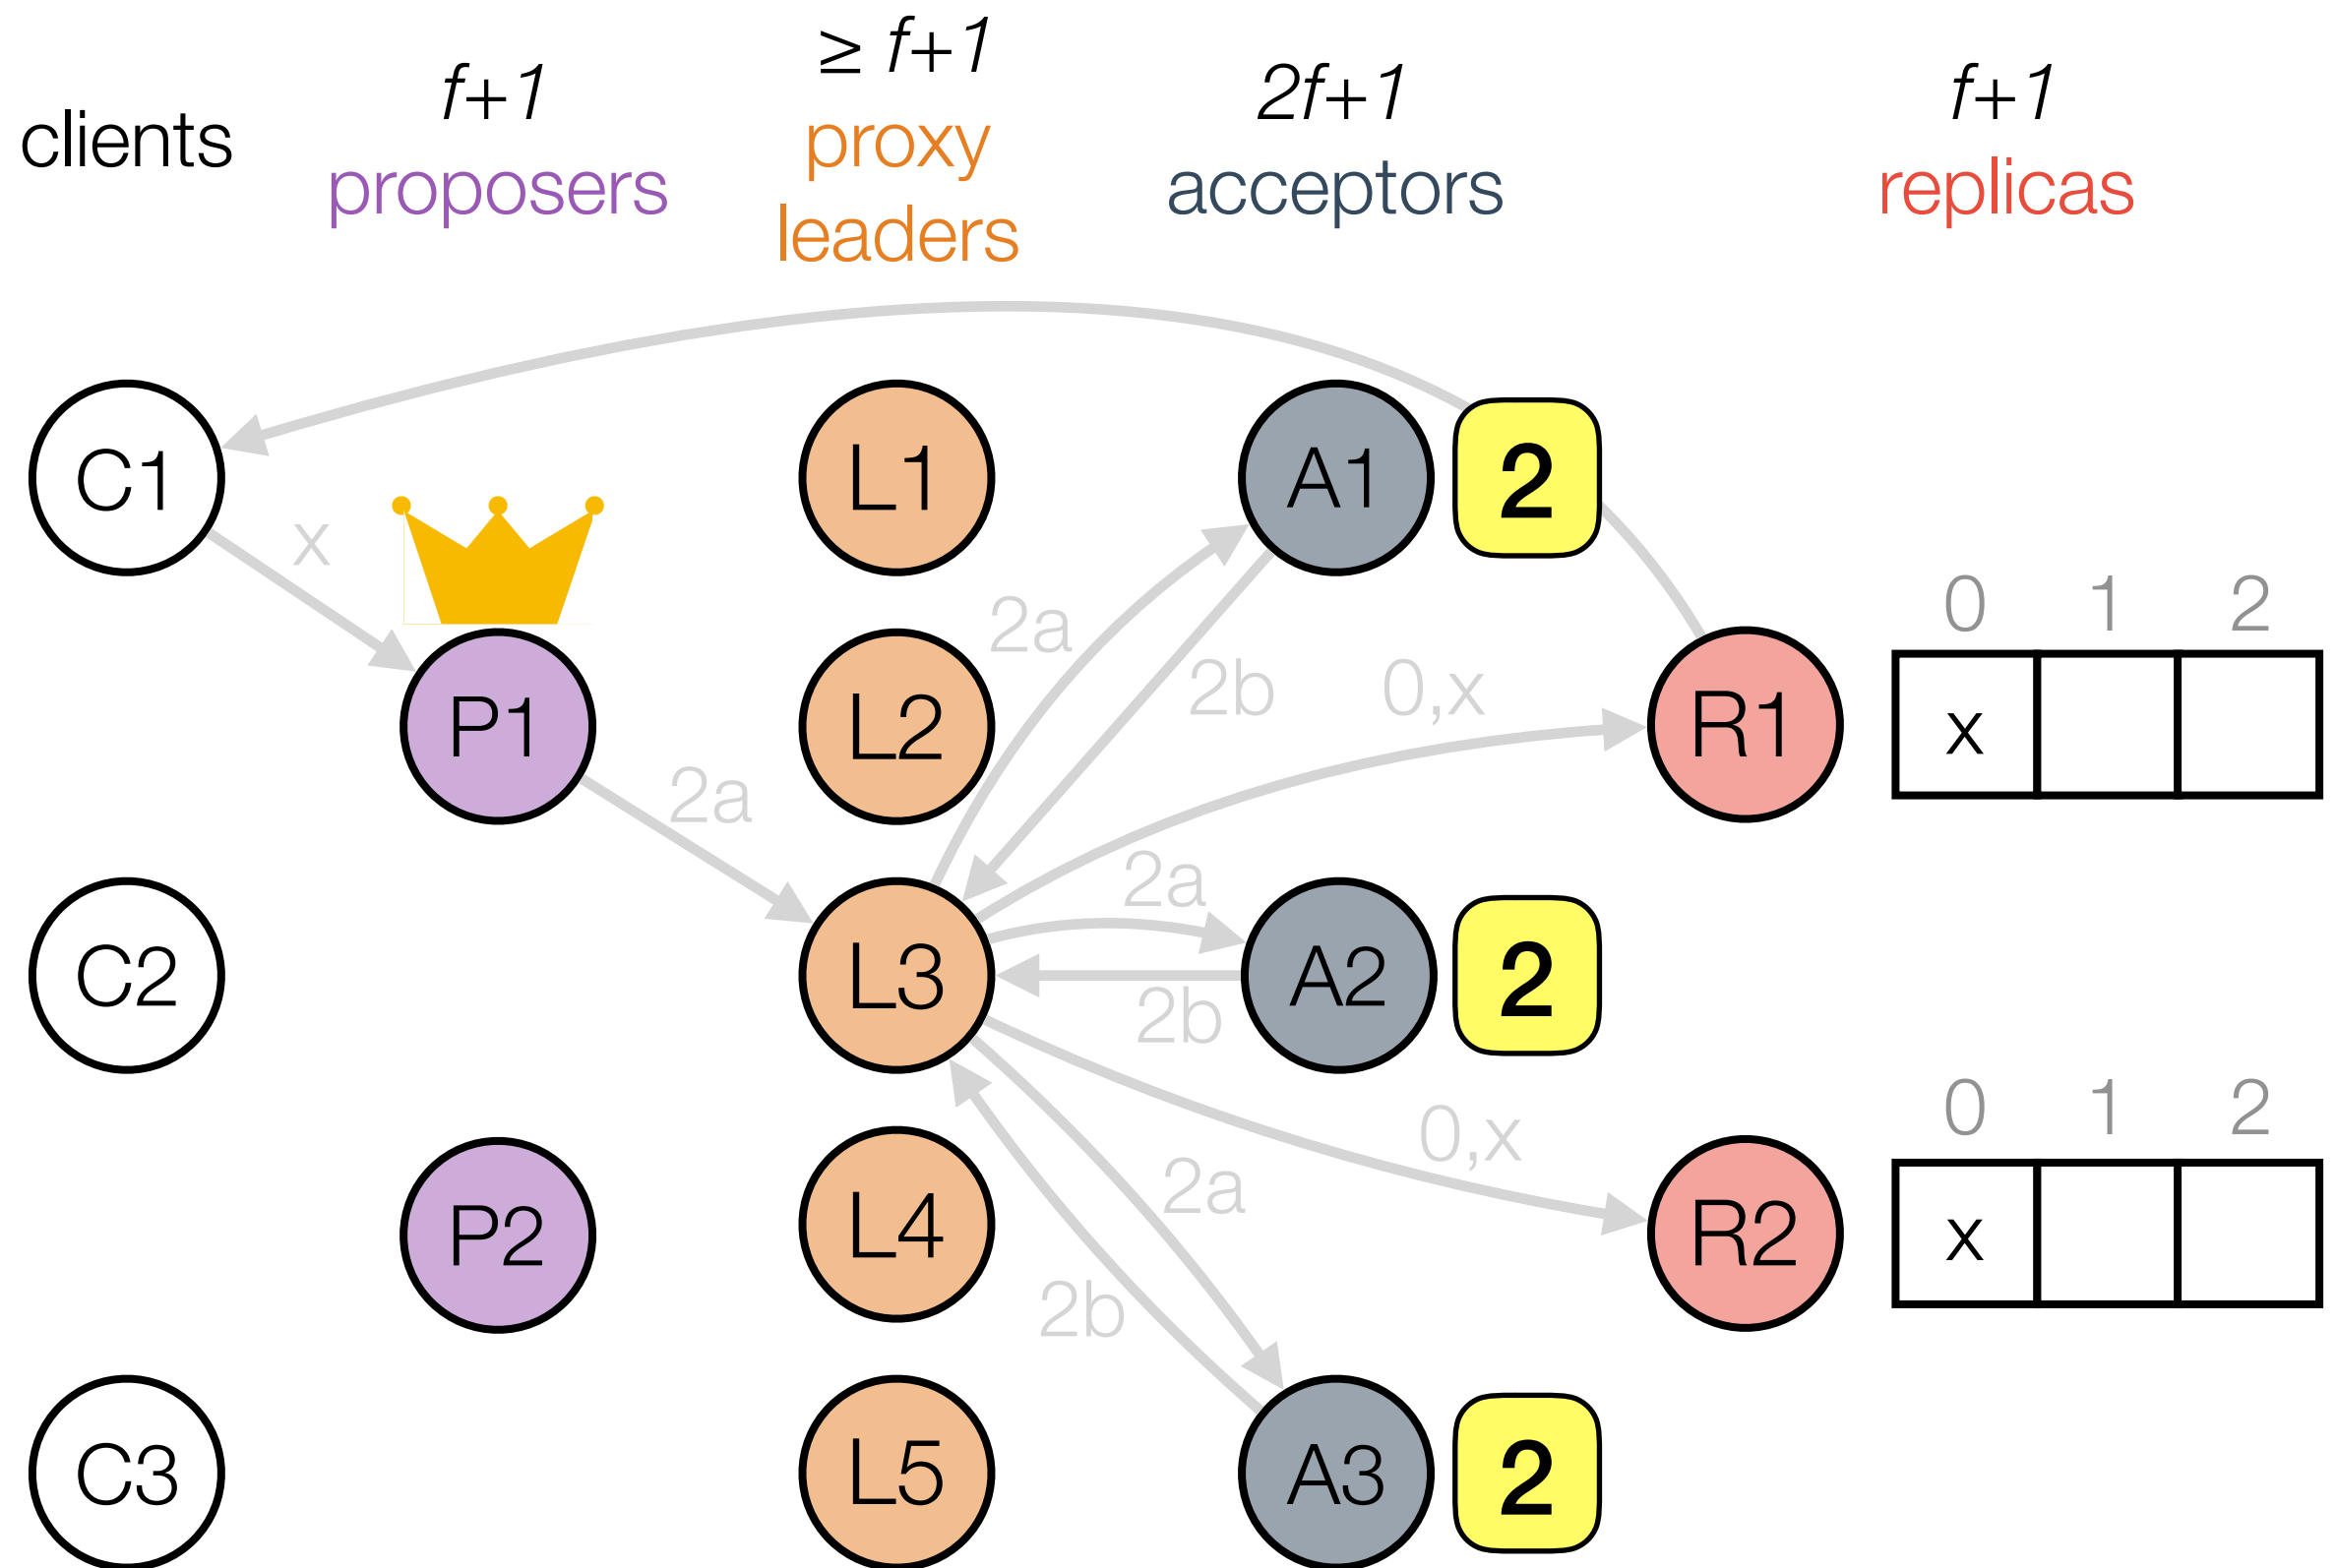


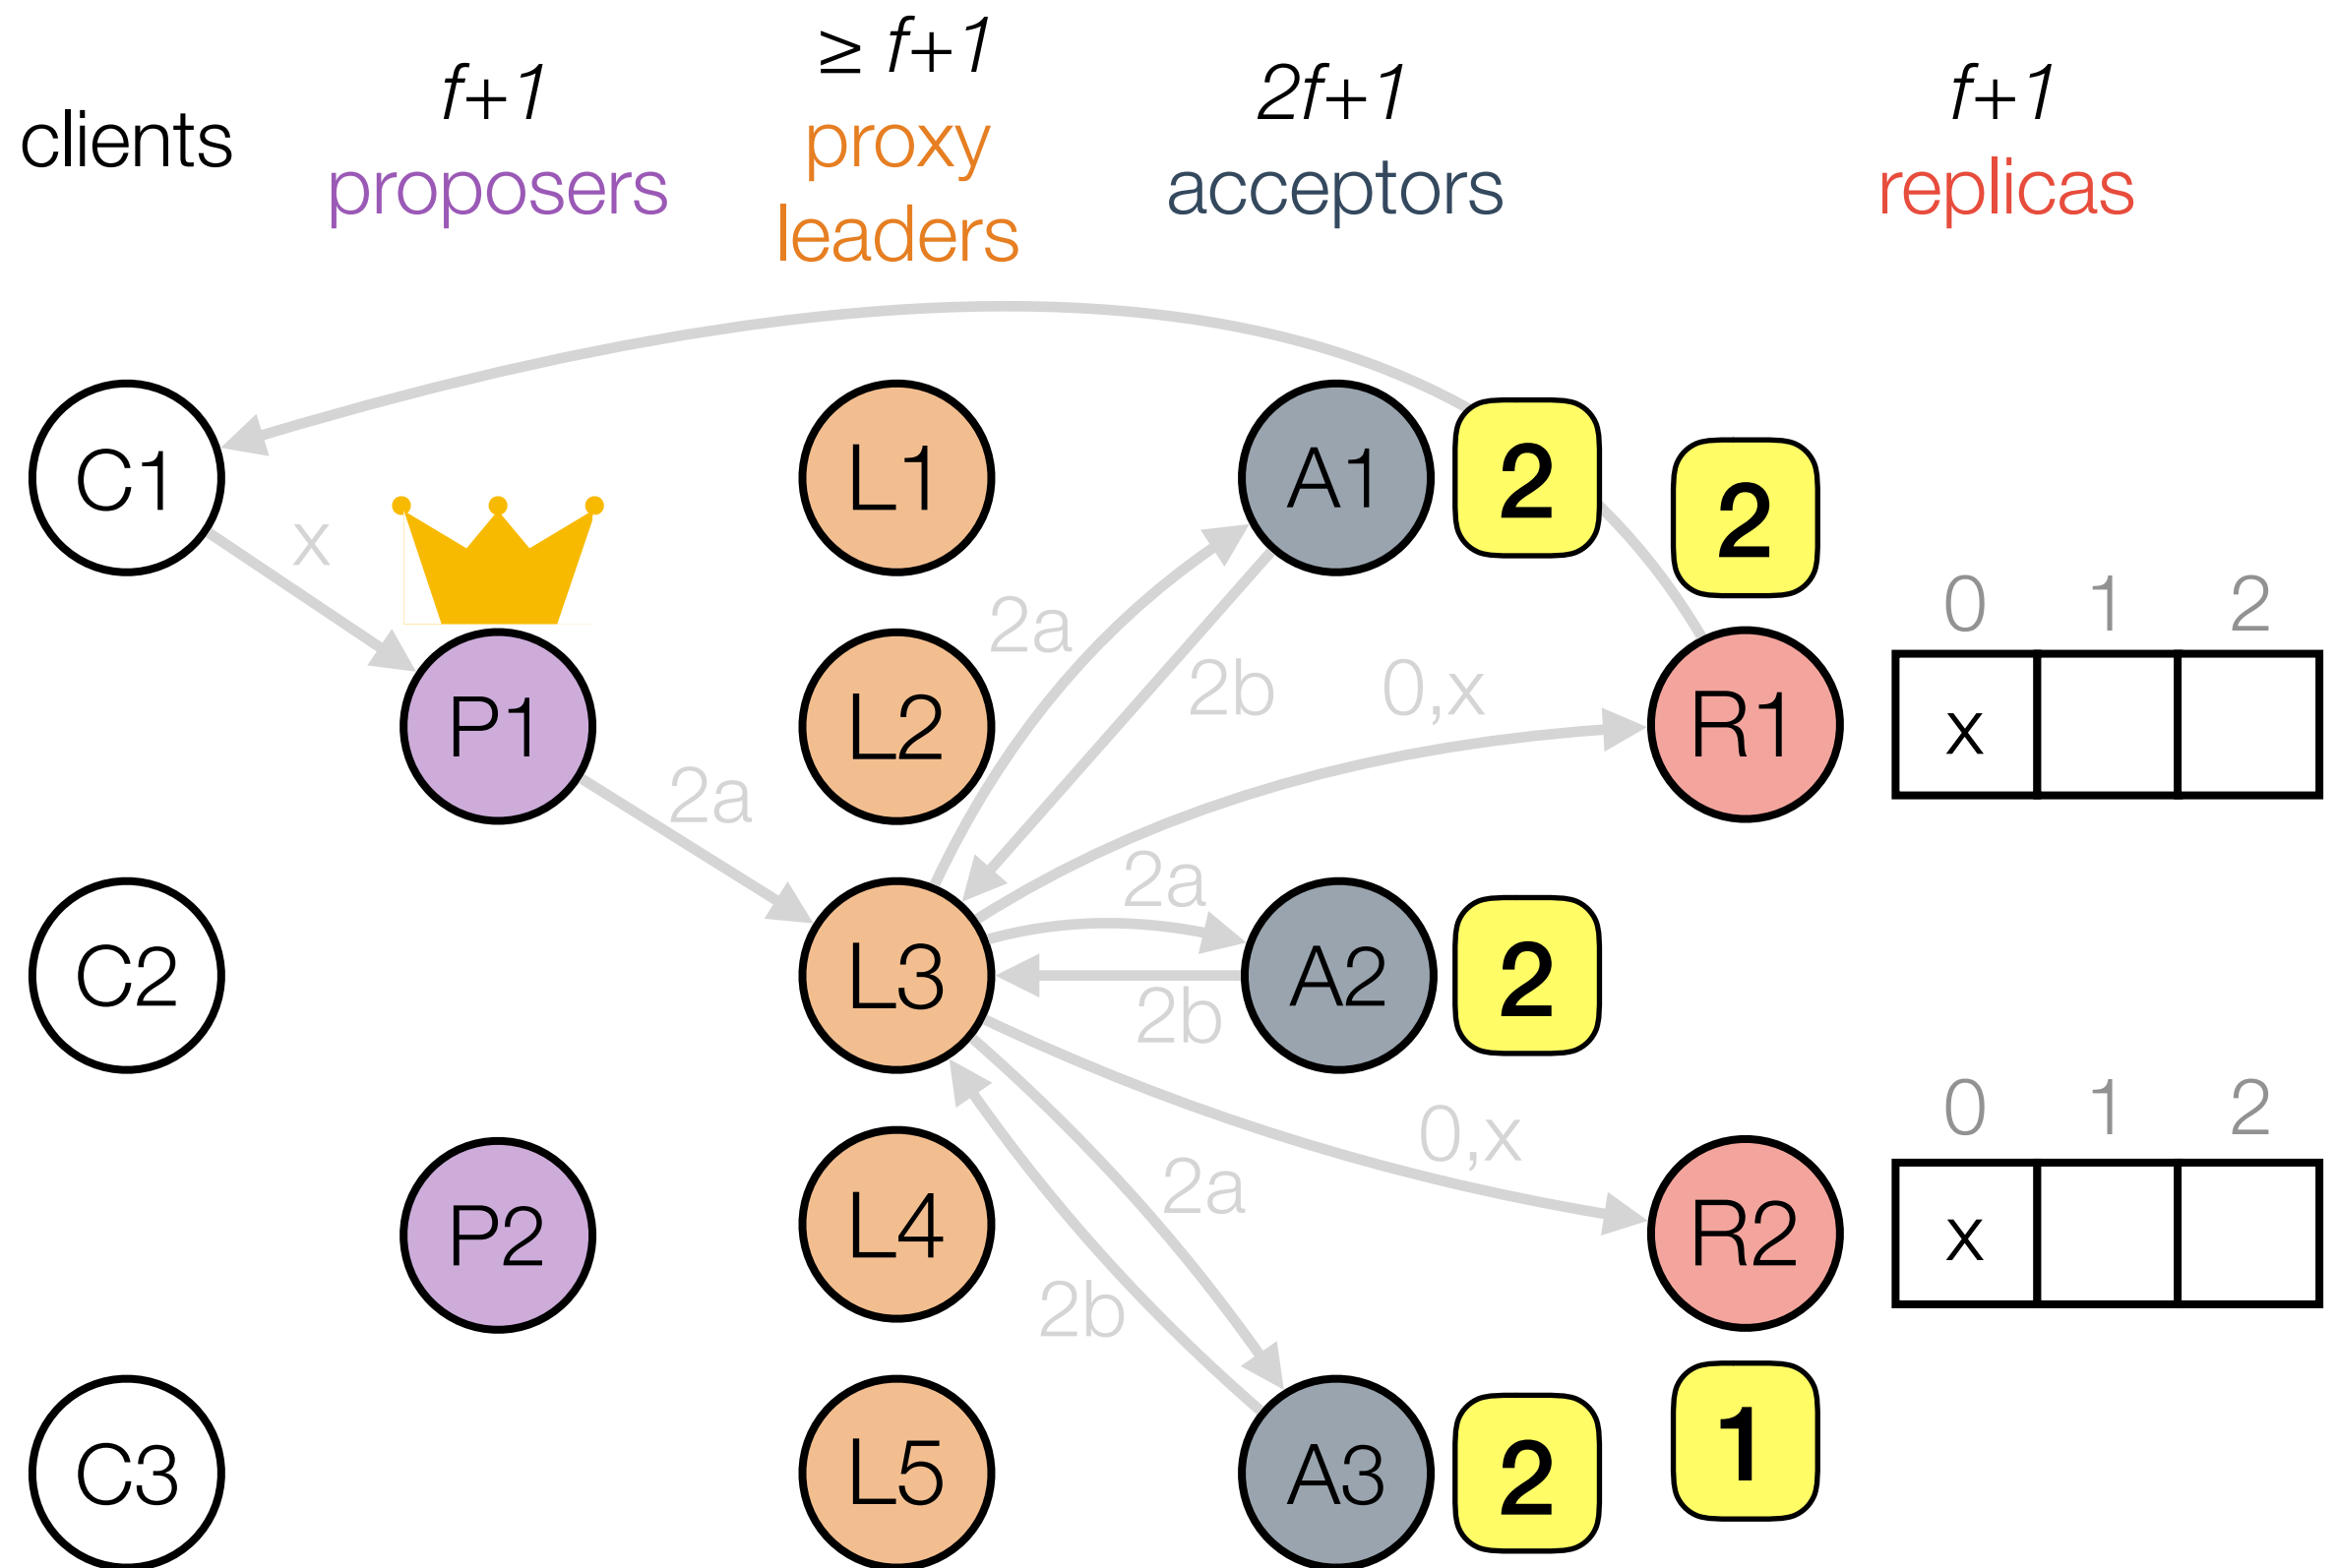


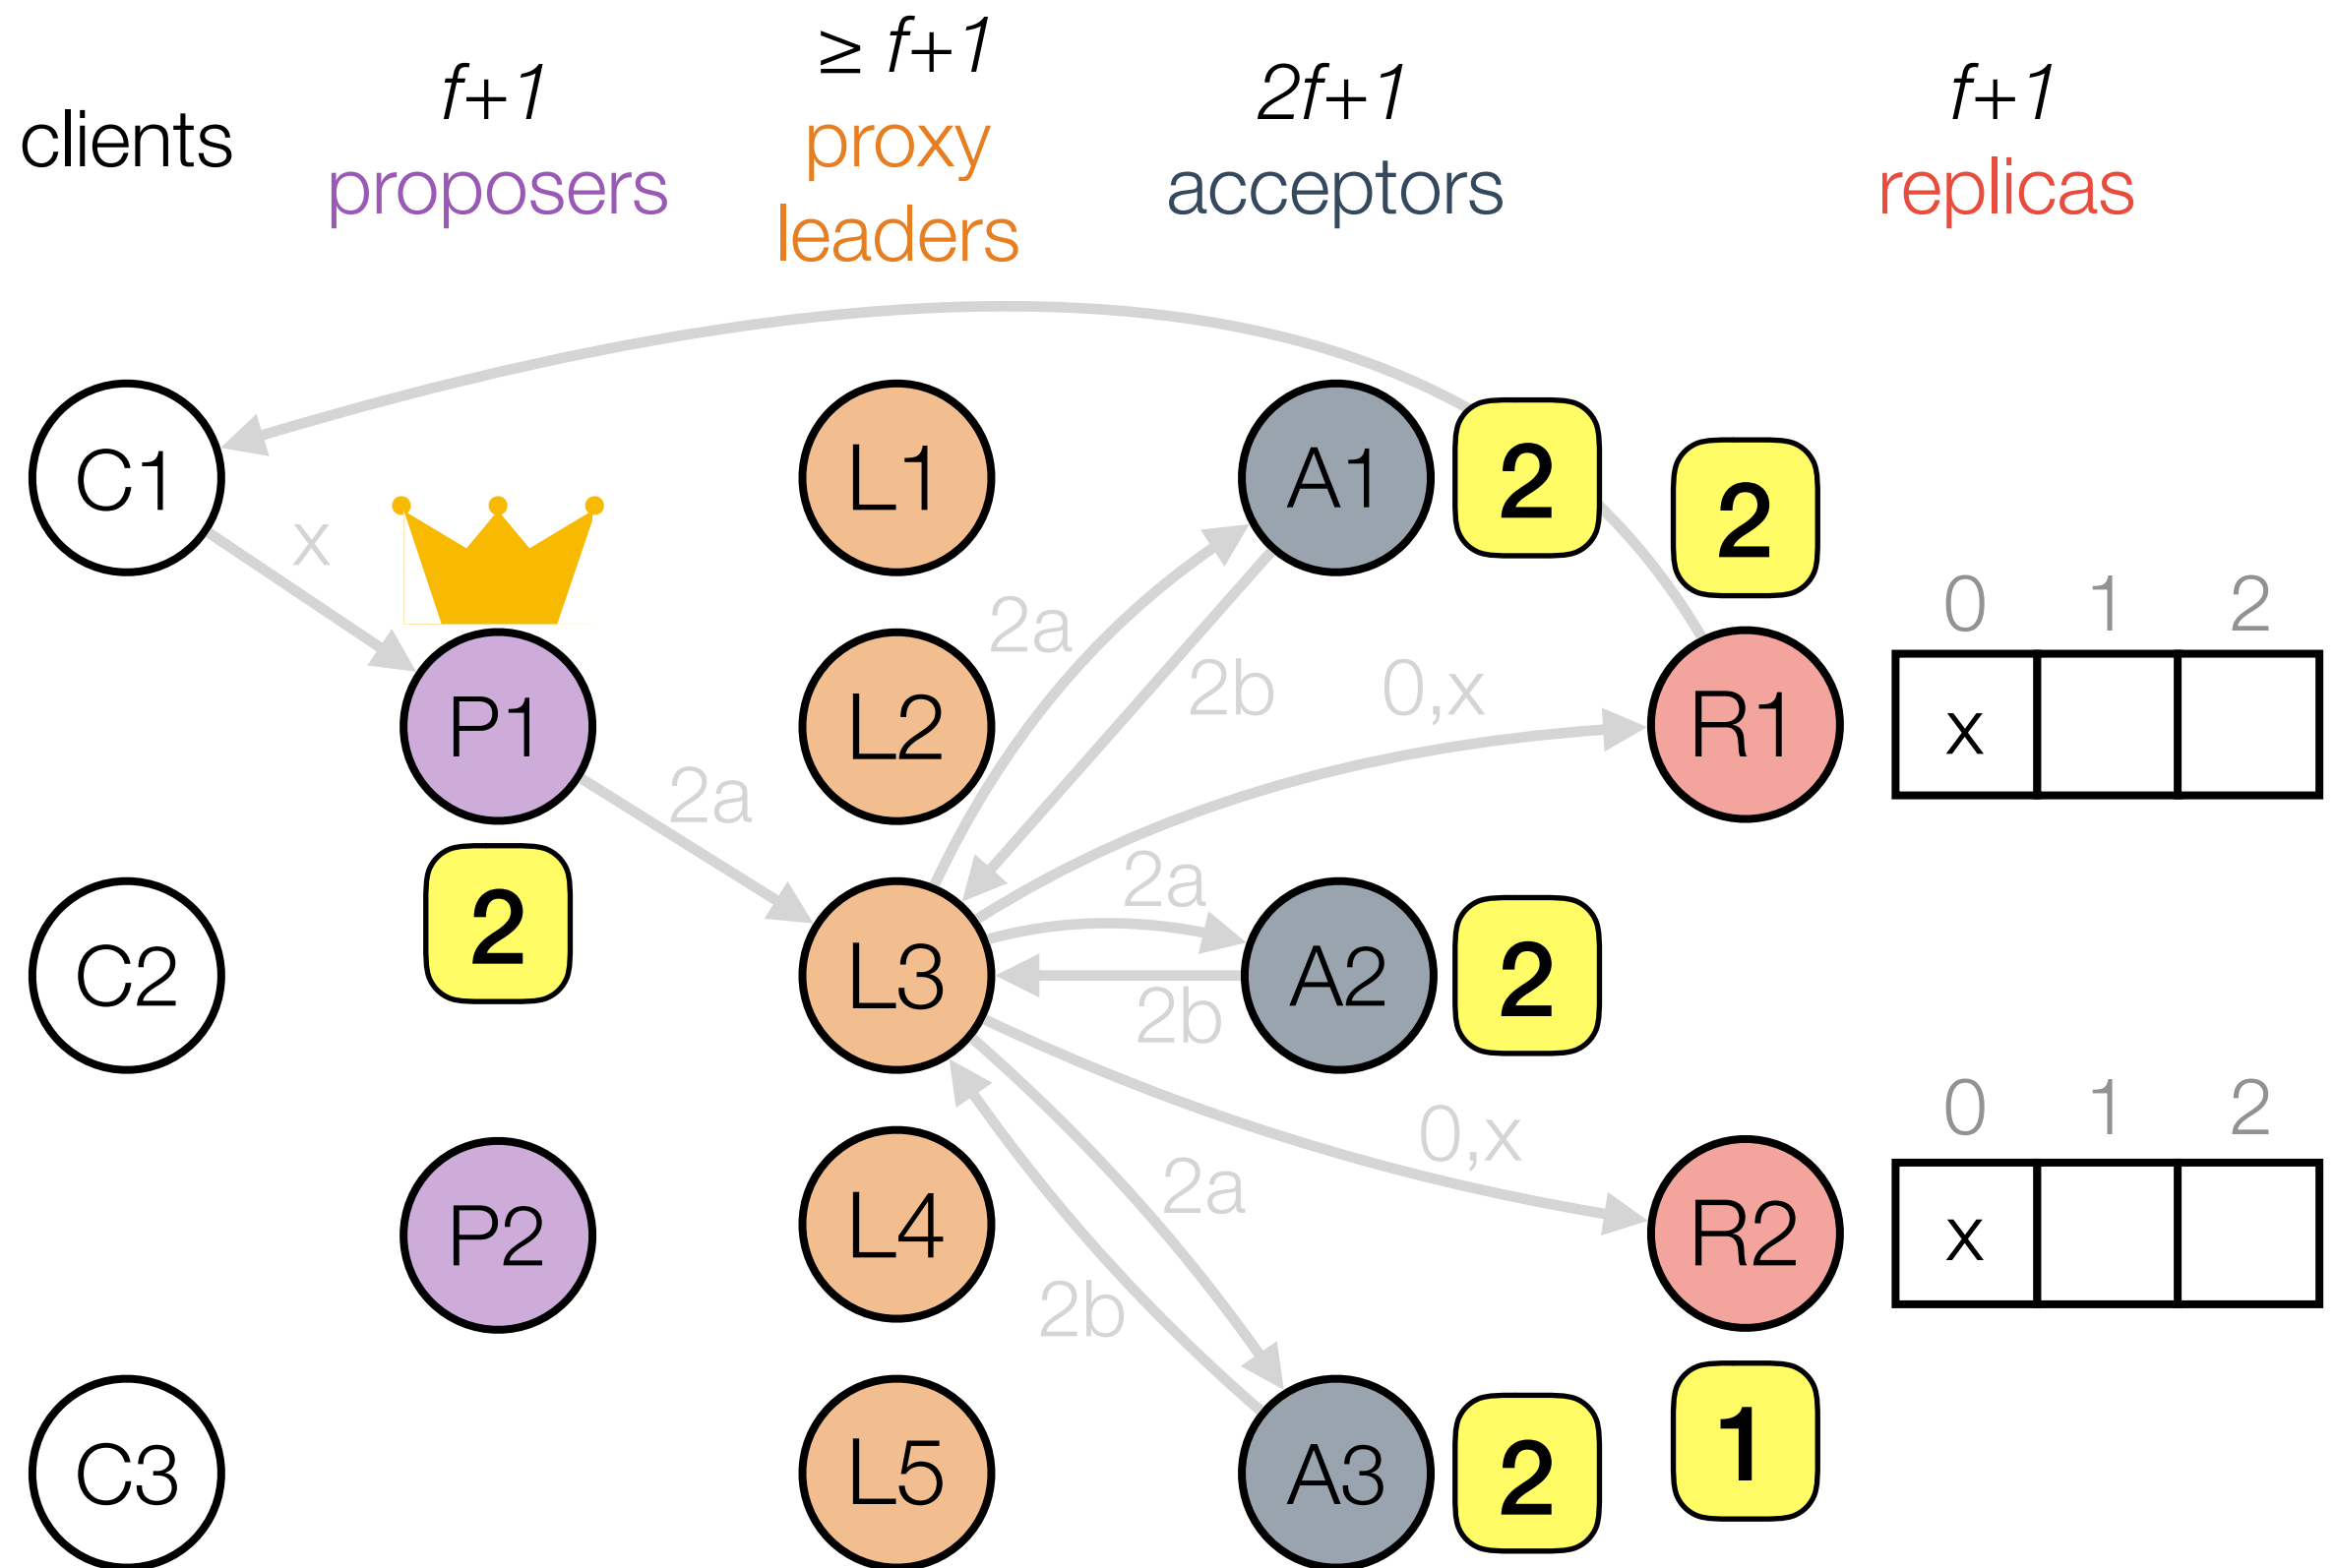


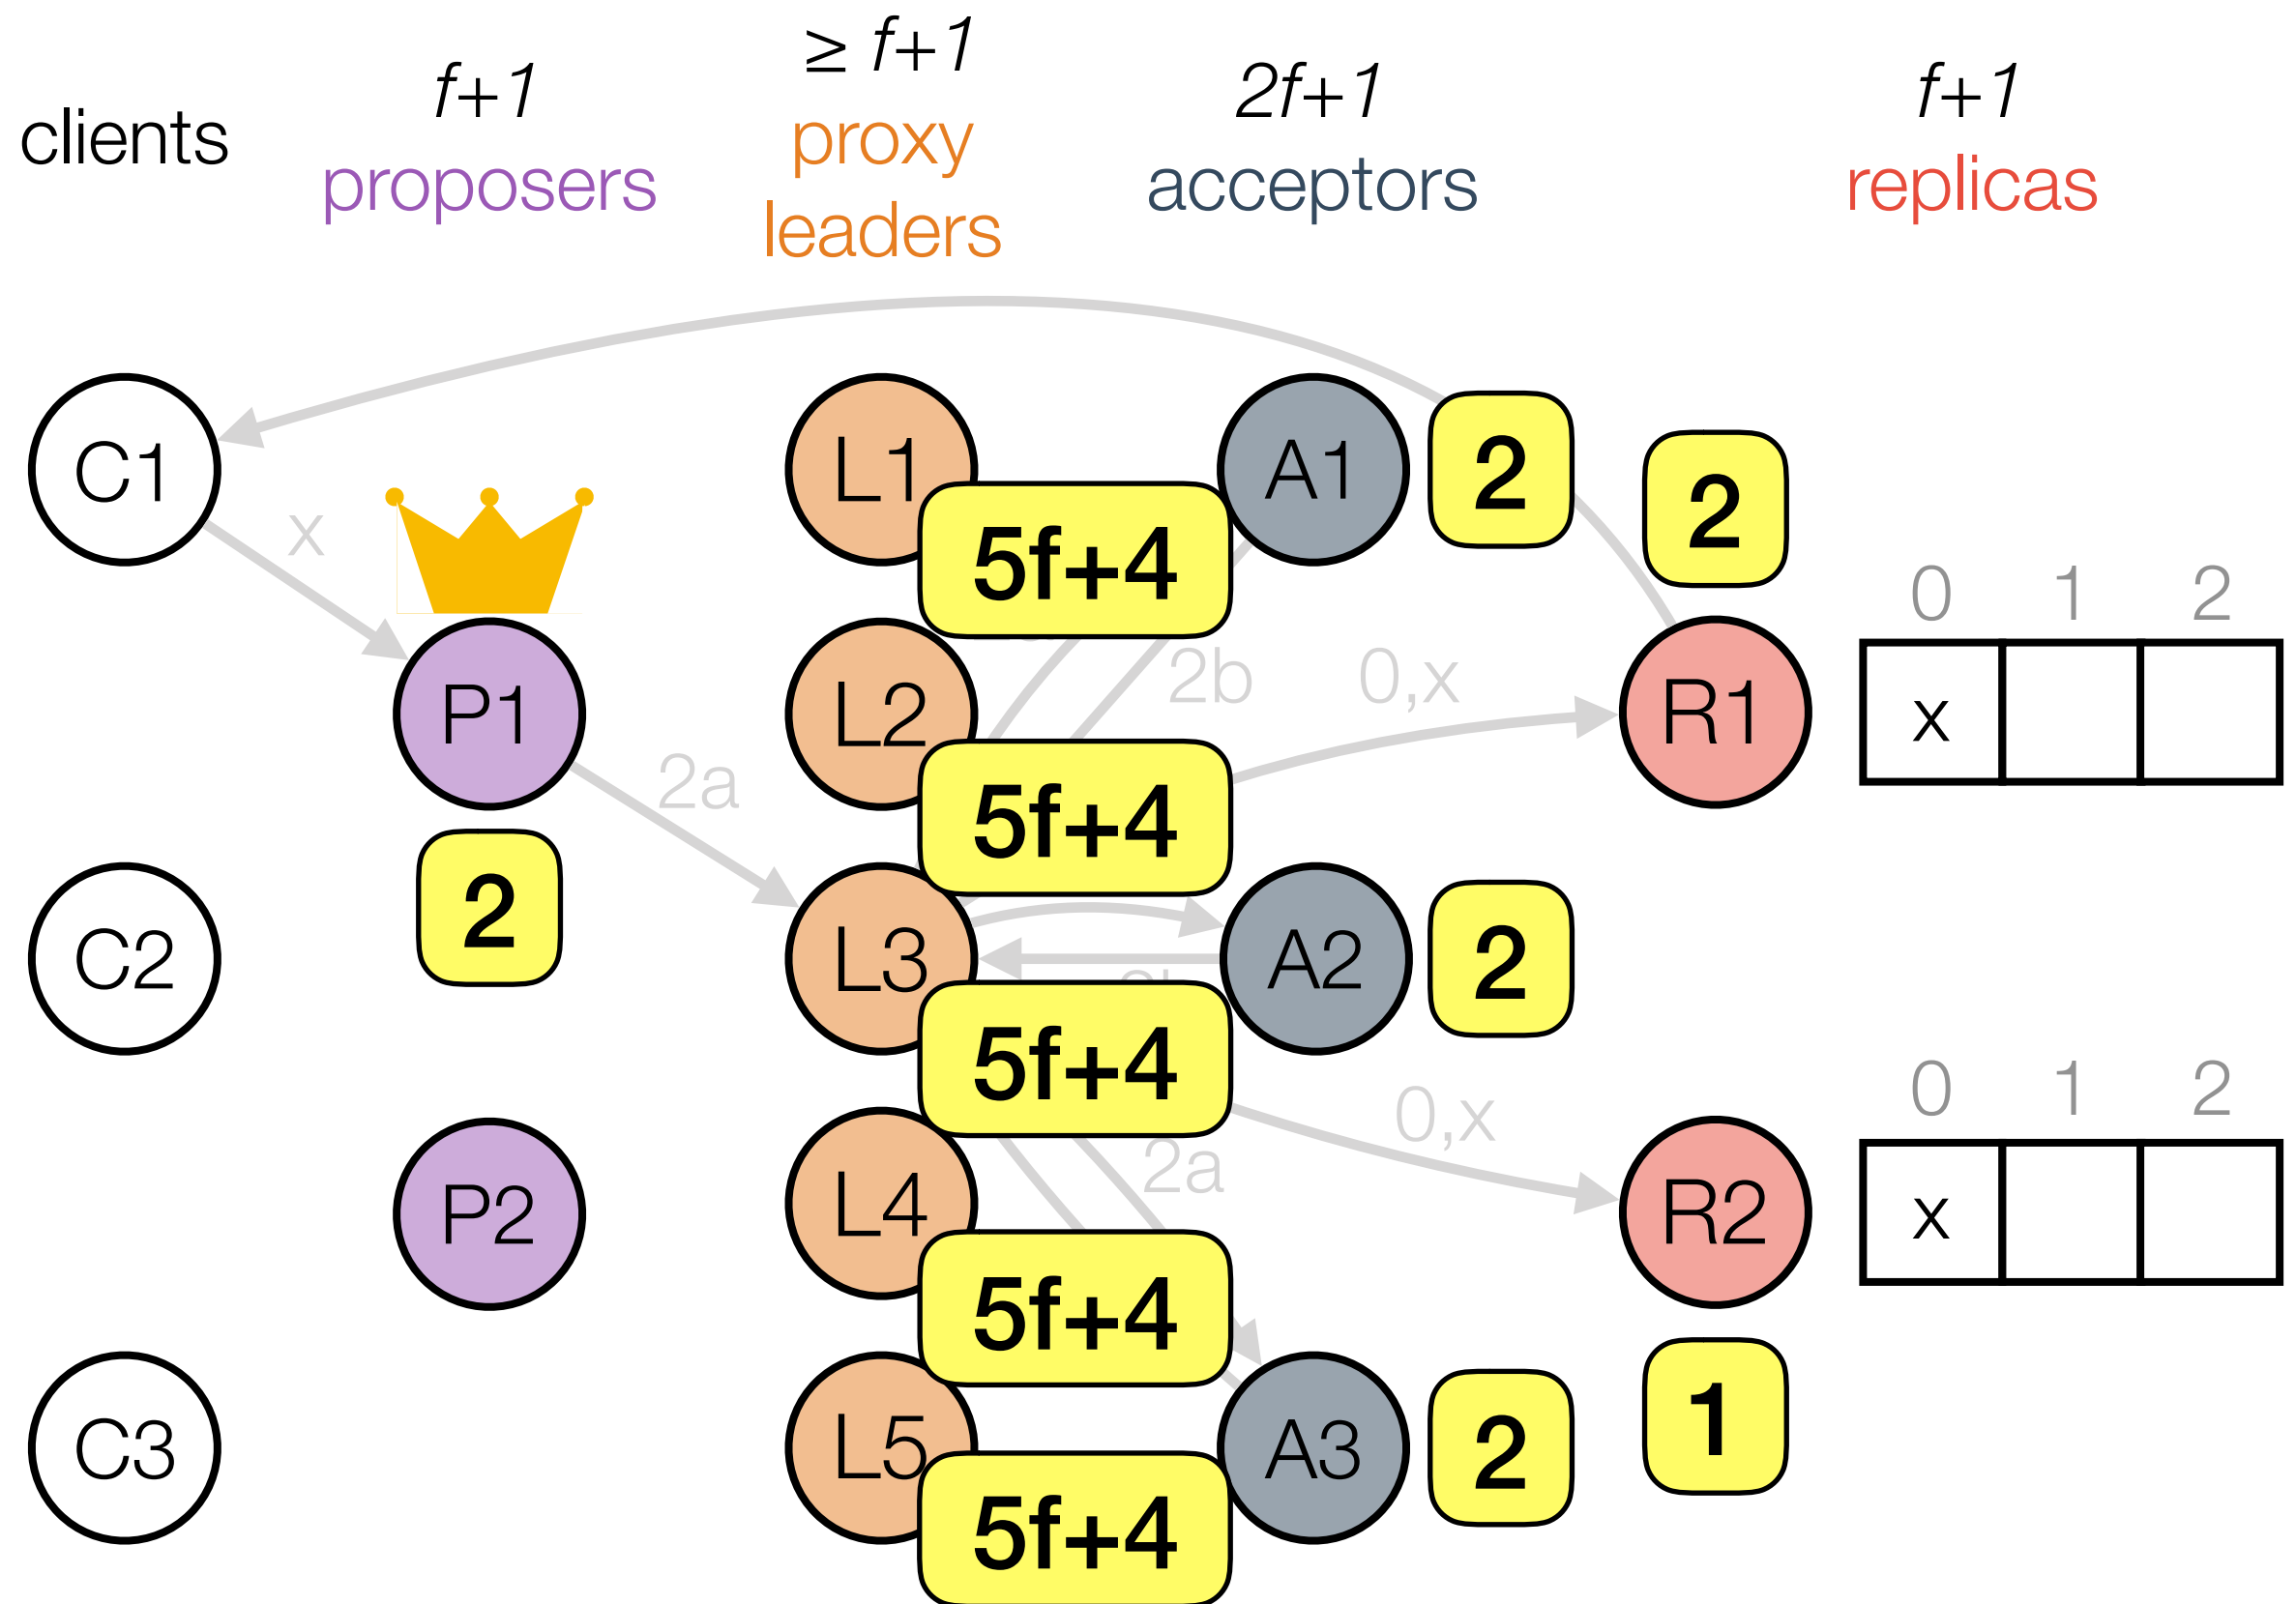




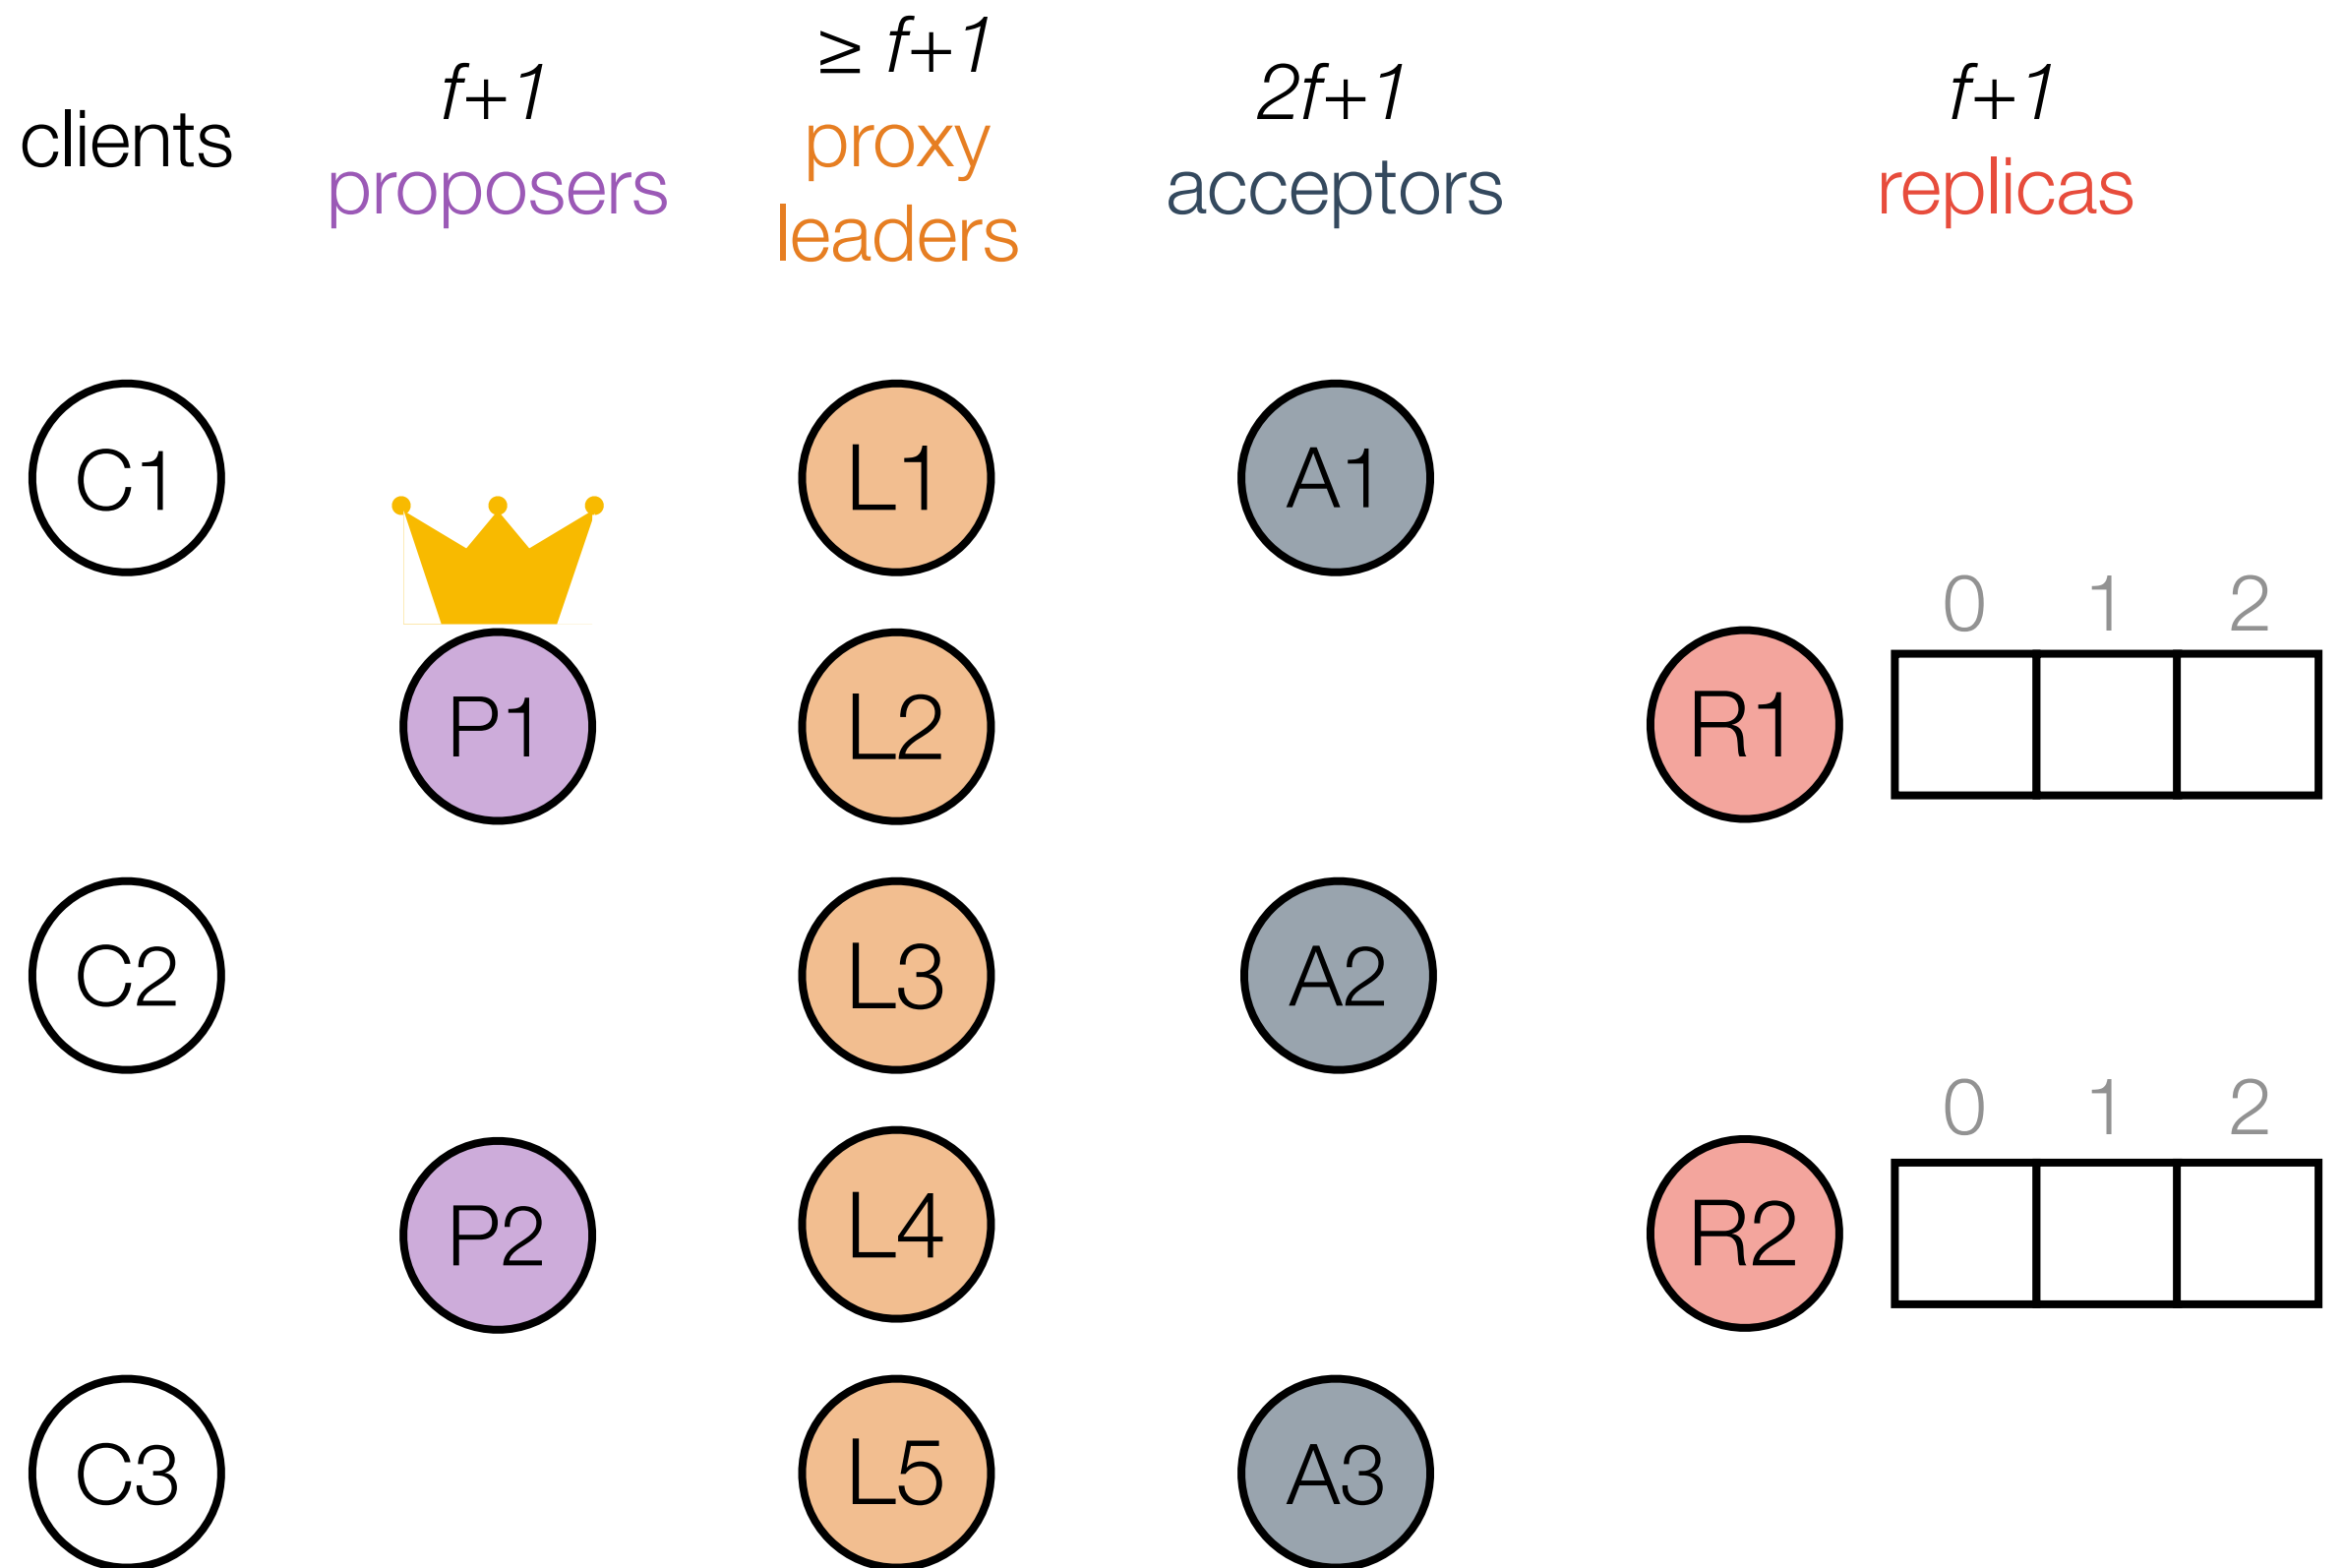









Compartmentalization:
Decouple and scale





Acceptors

0	1	2	3	4	5	6	7	
a	b	c	d	e	f			

Acceptors


0	1	2	3	4	5	6	7	
a	b	c	d	e	f			

A1


A2

A3


Acceptors

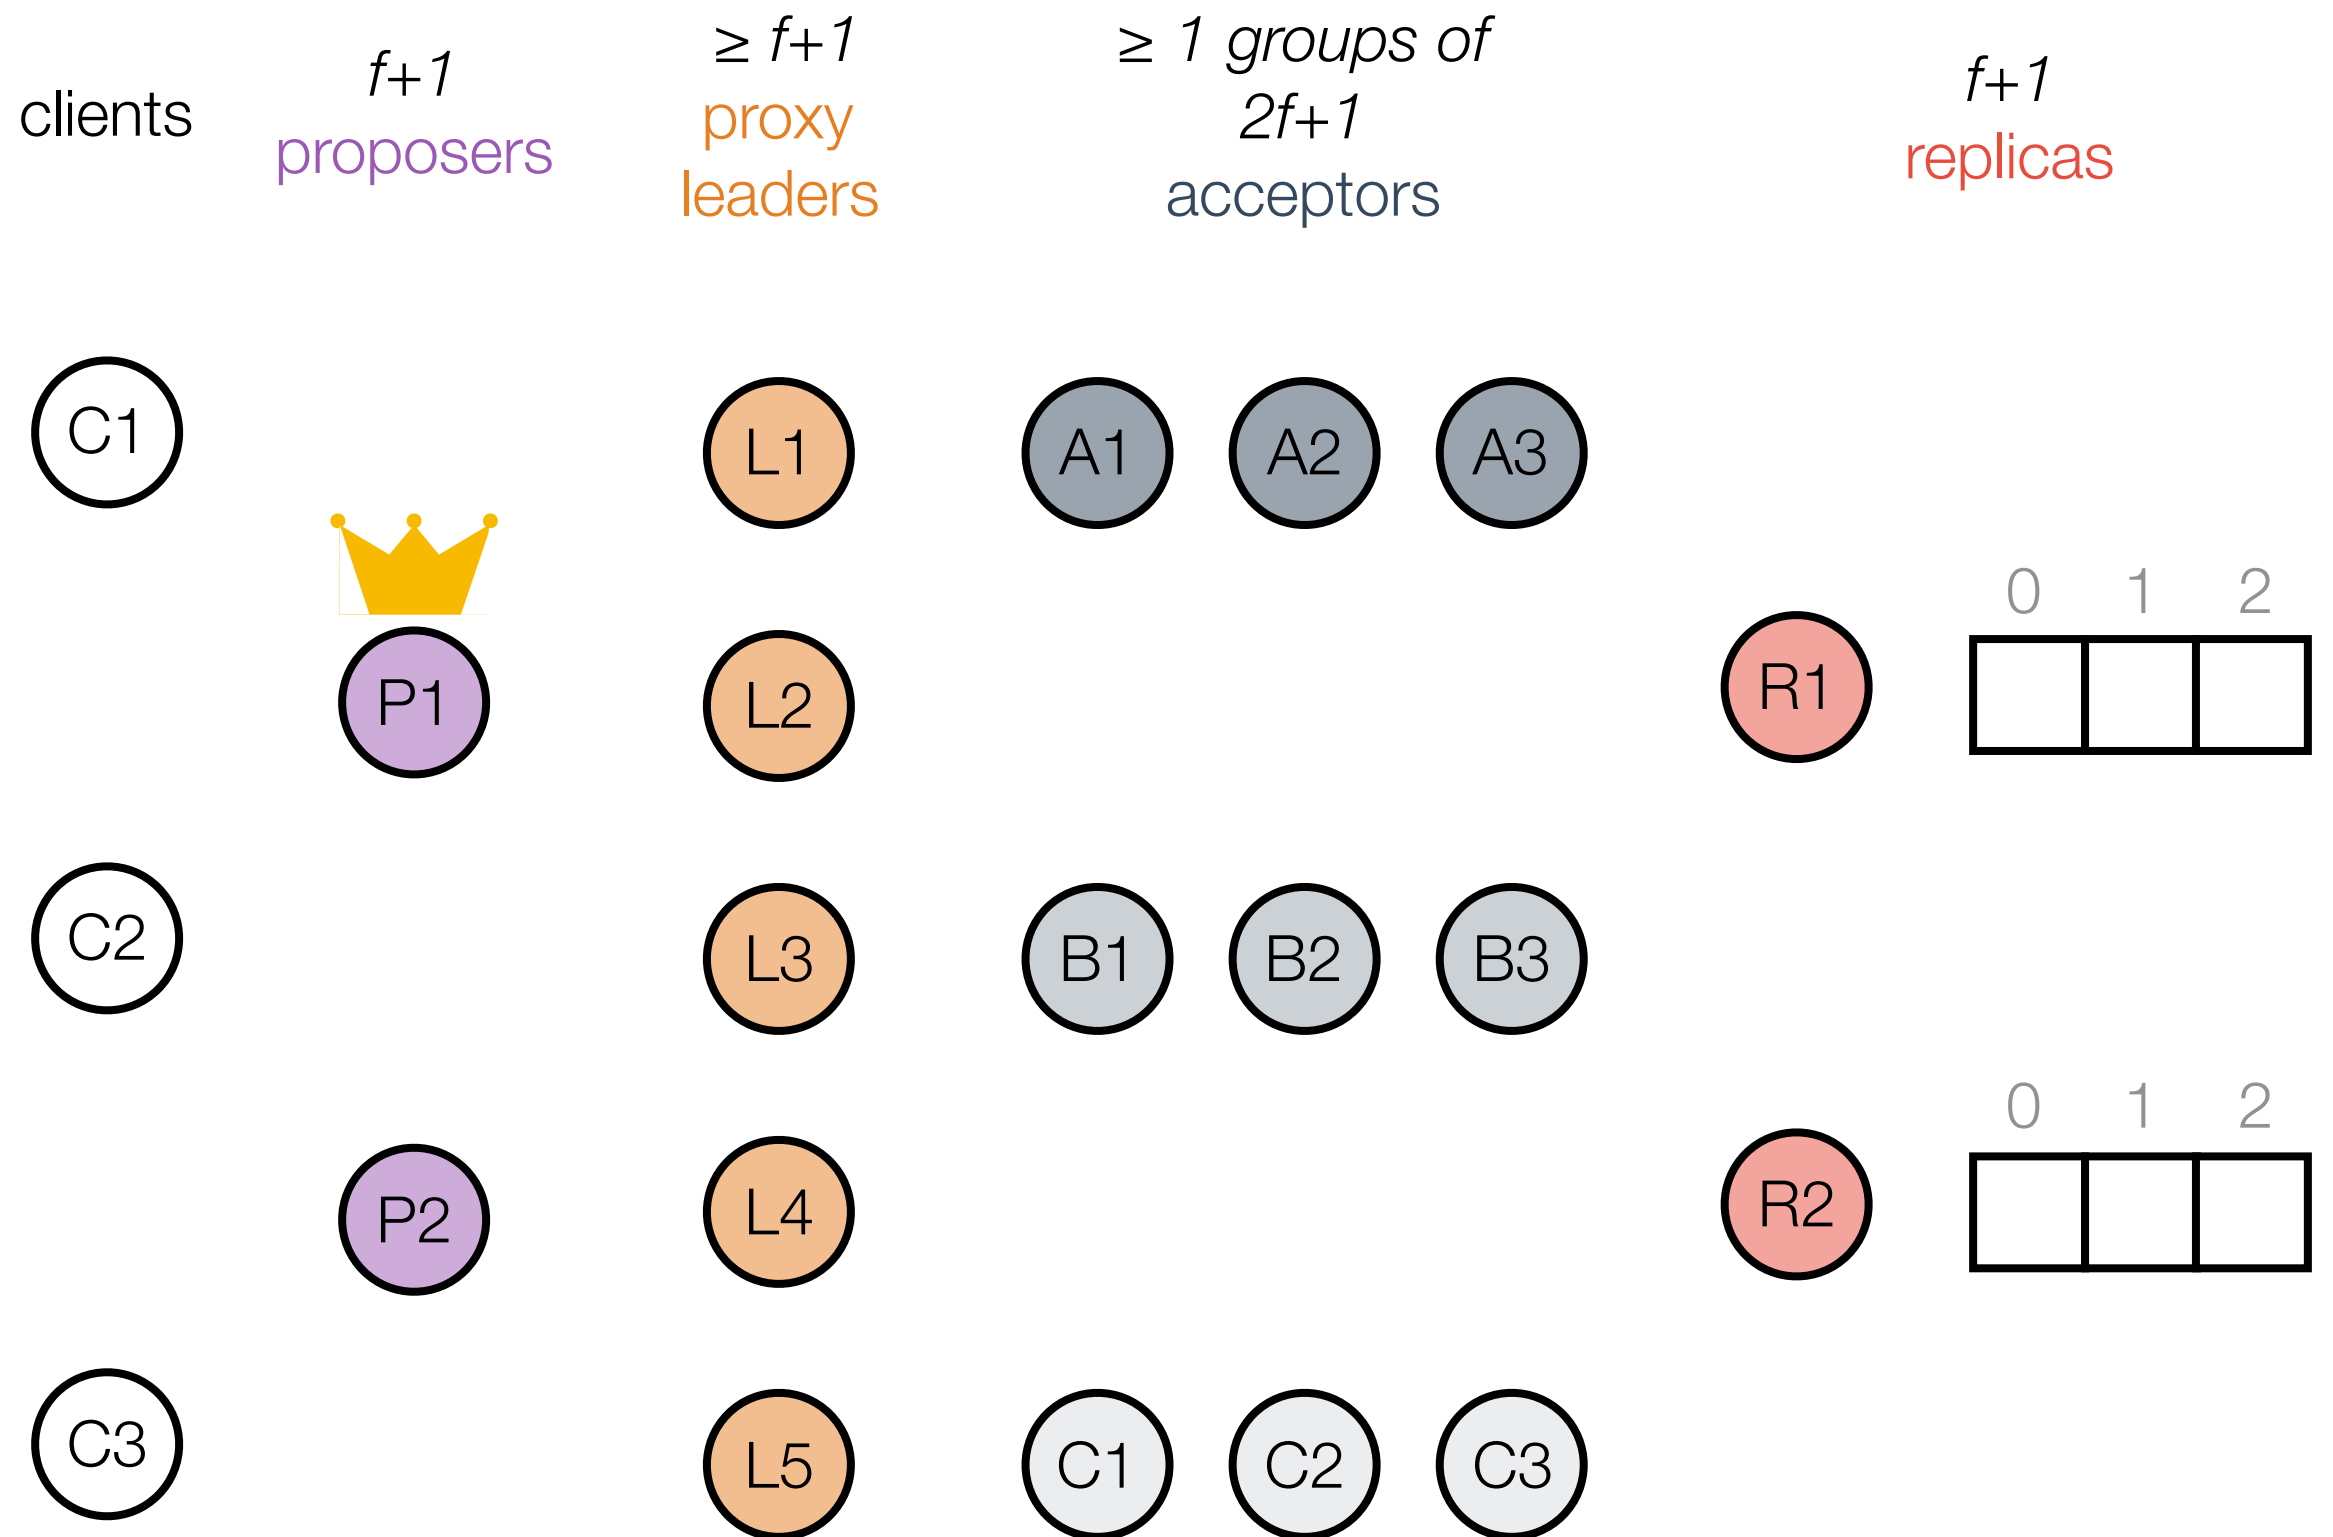


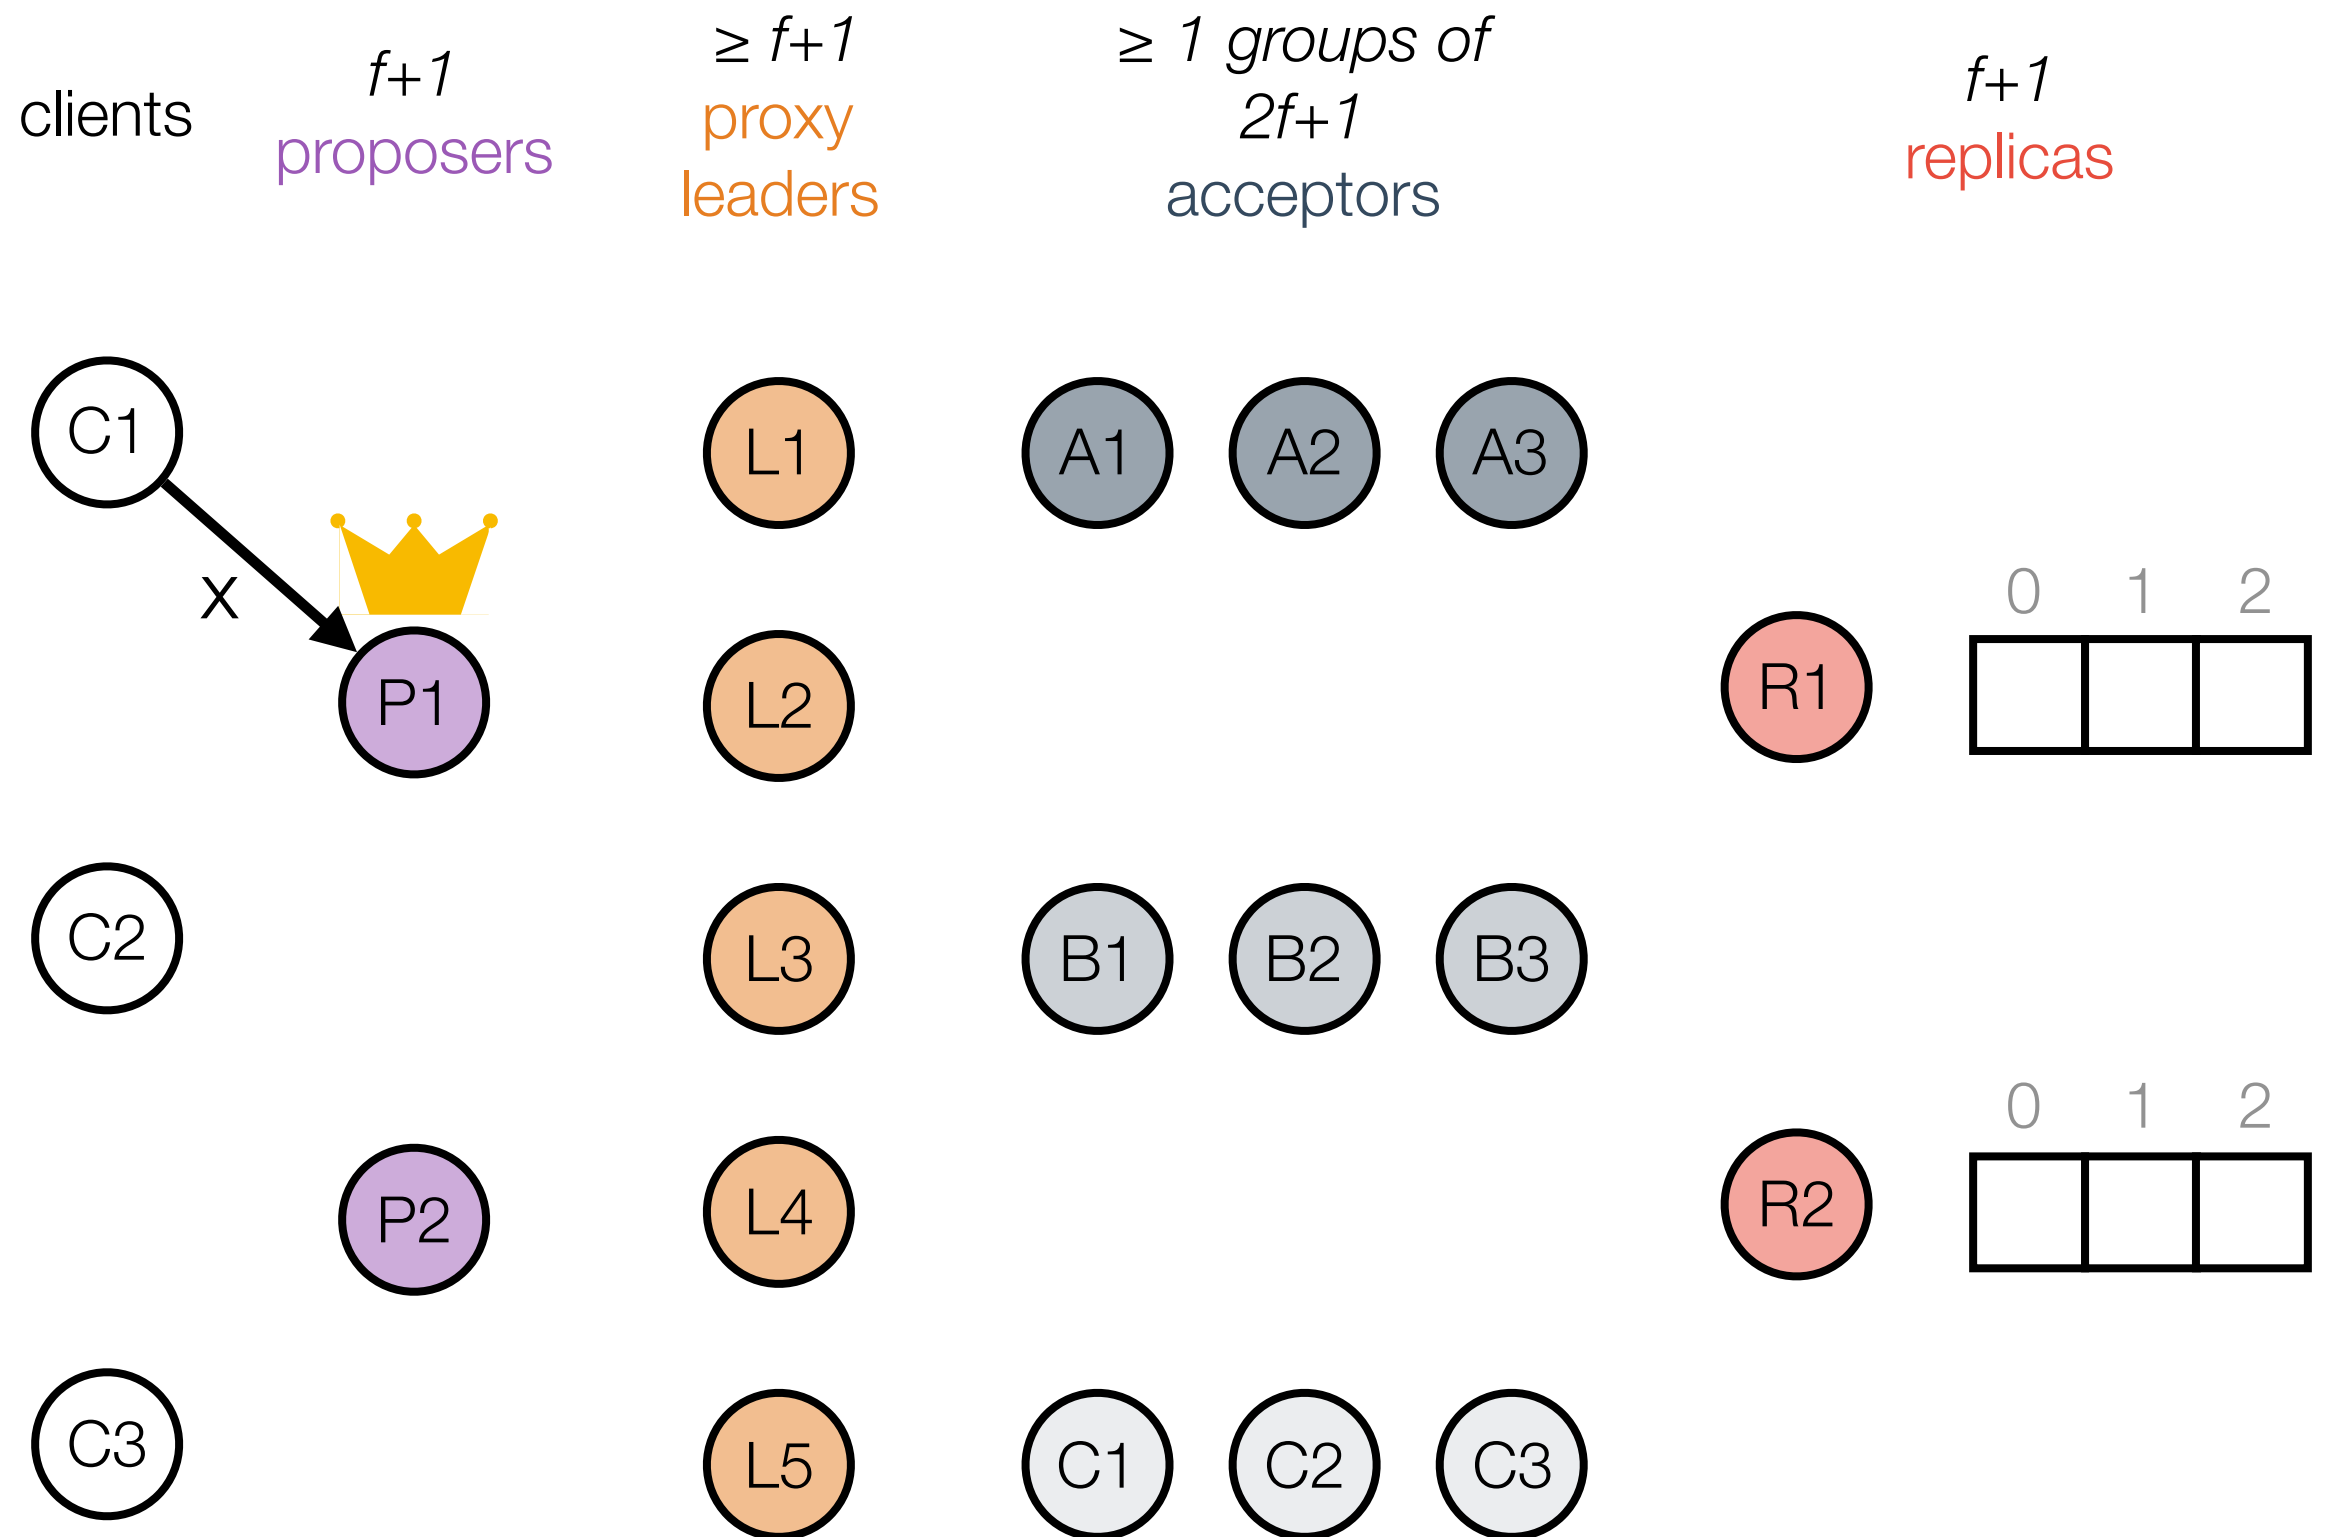
Acceptors

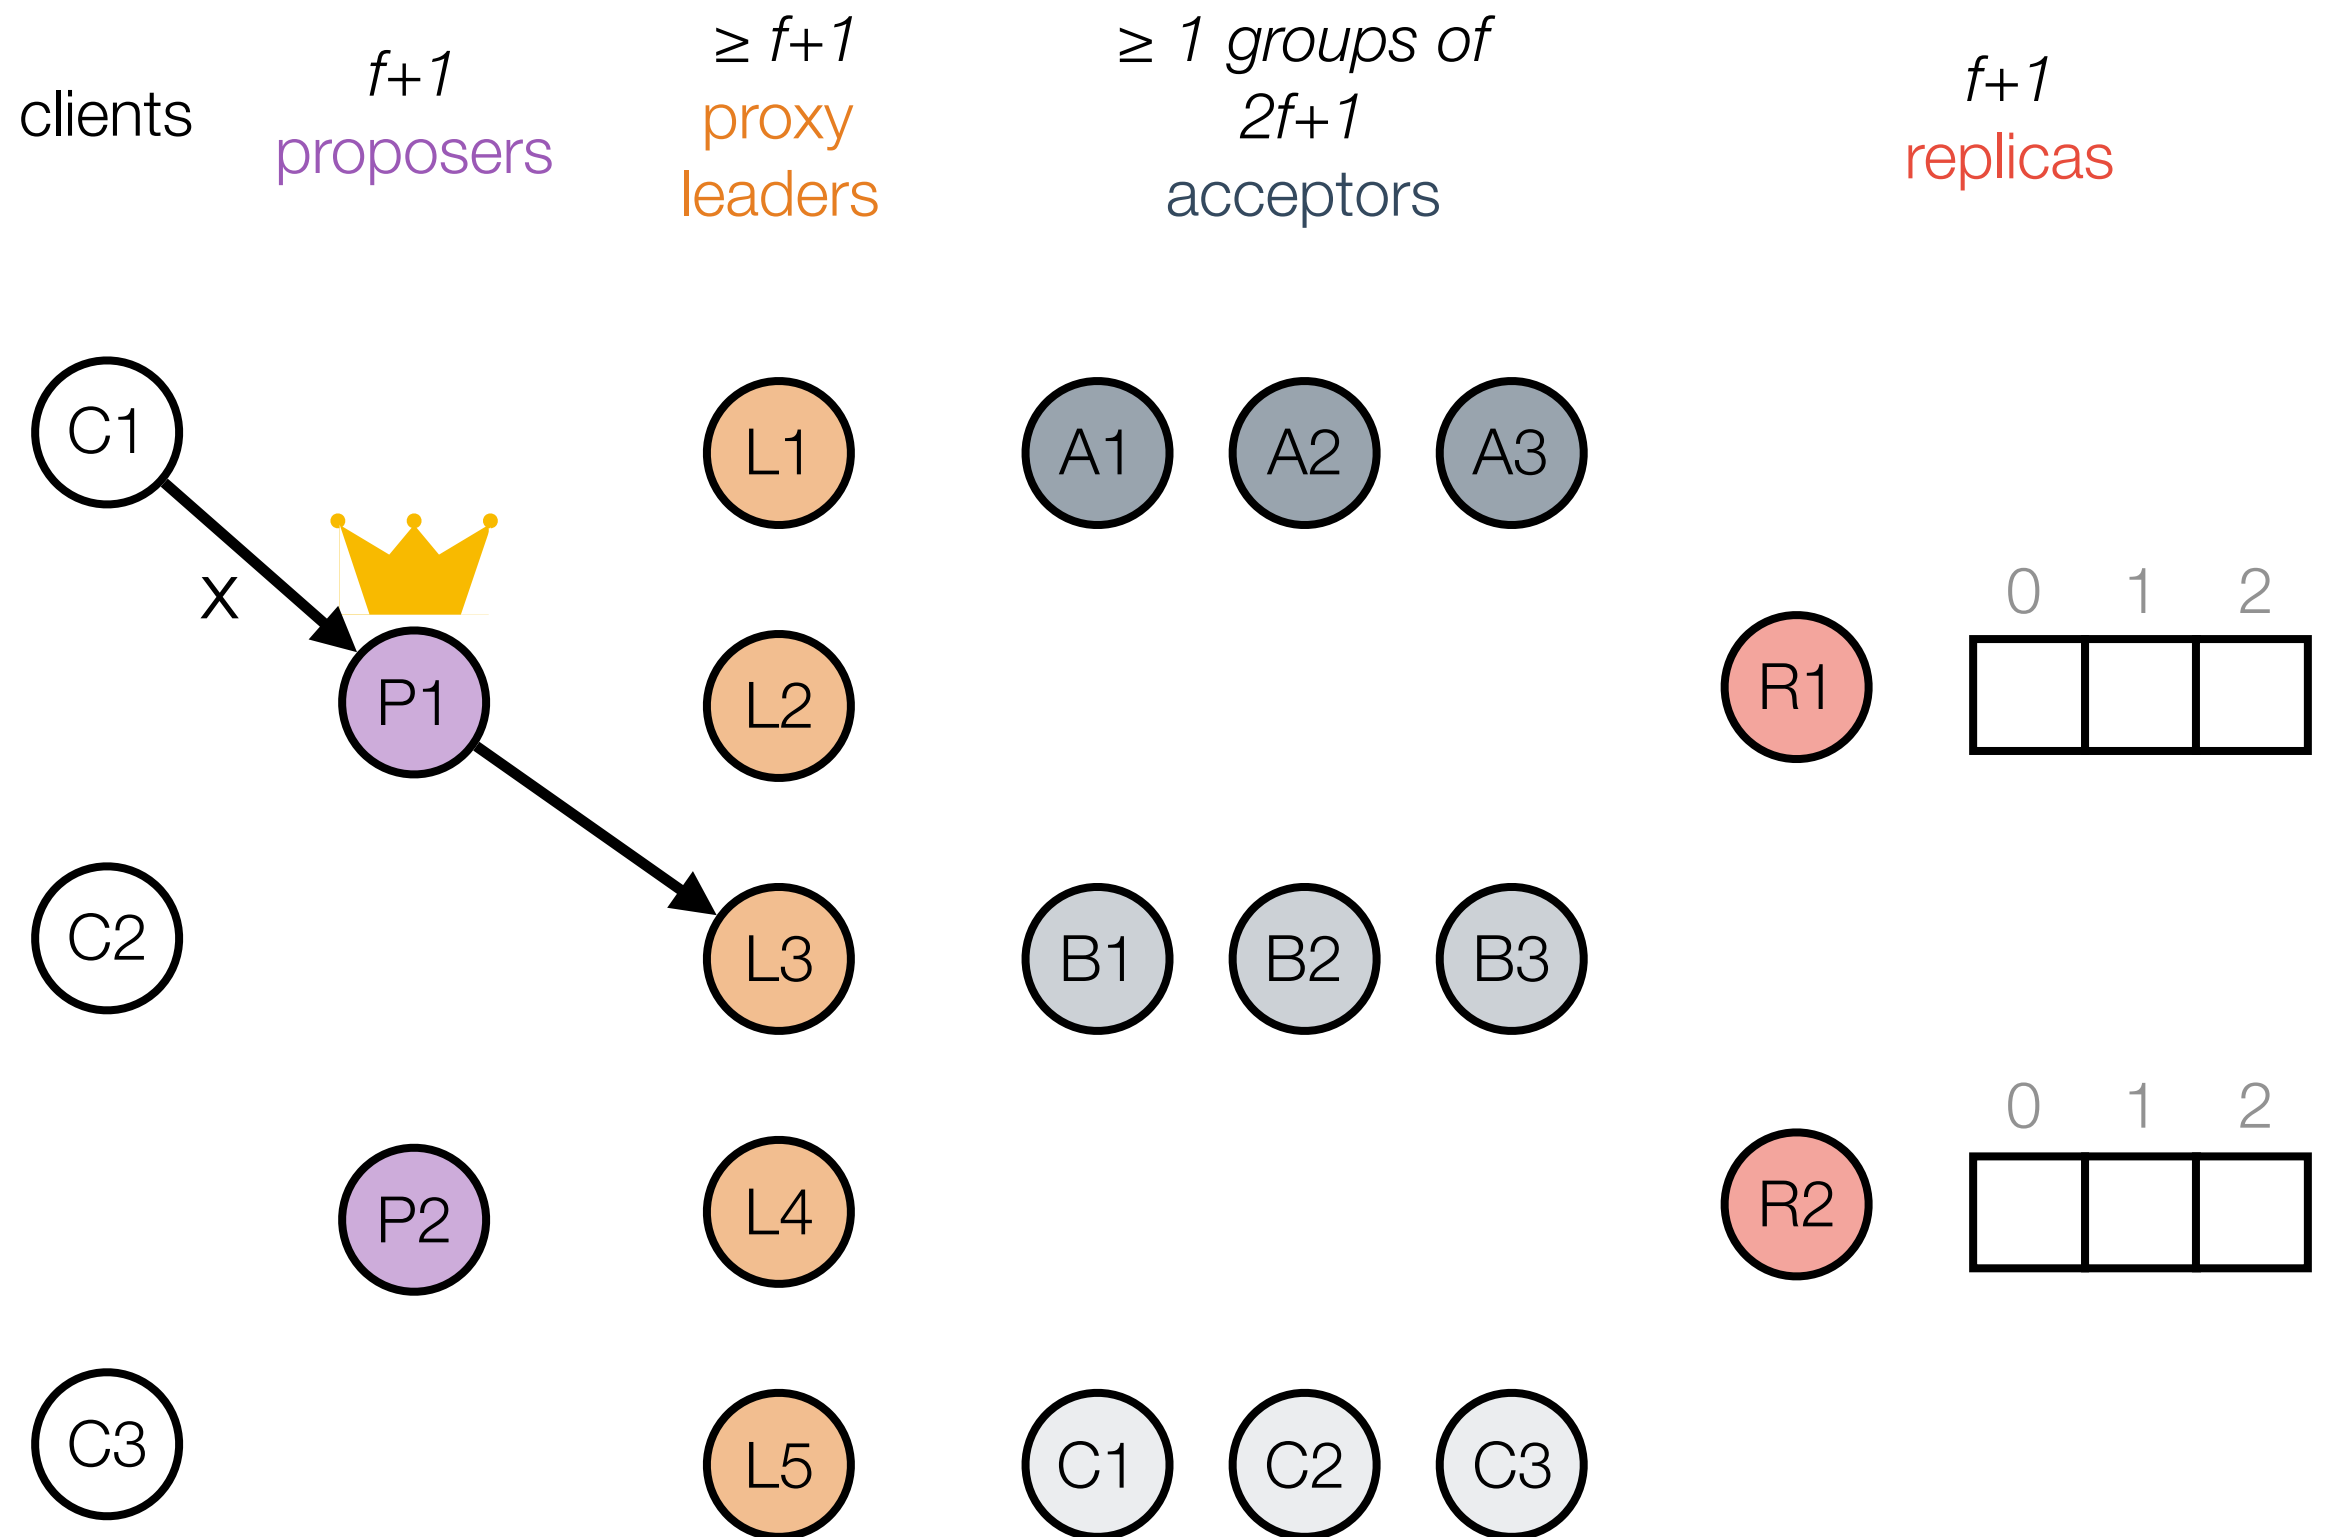


Acceptors









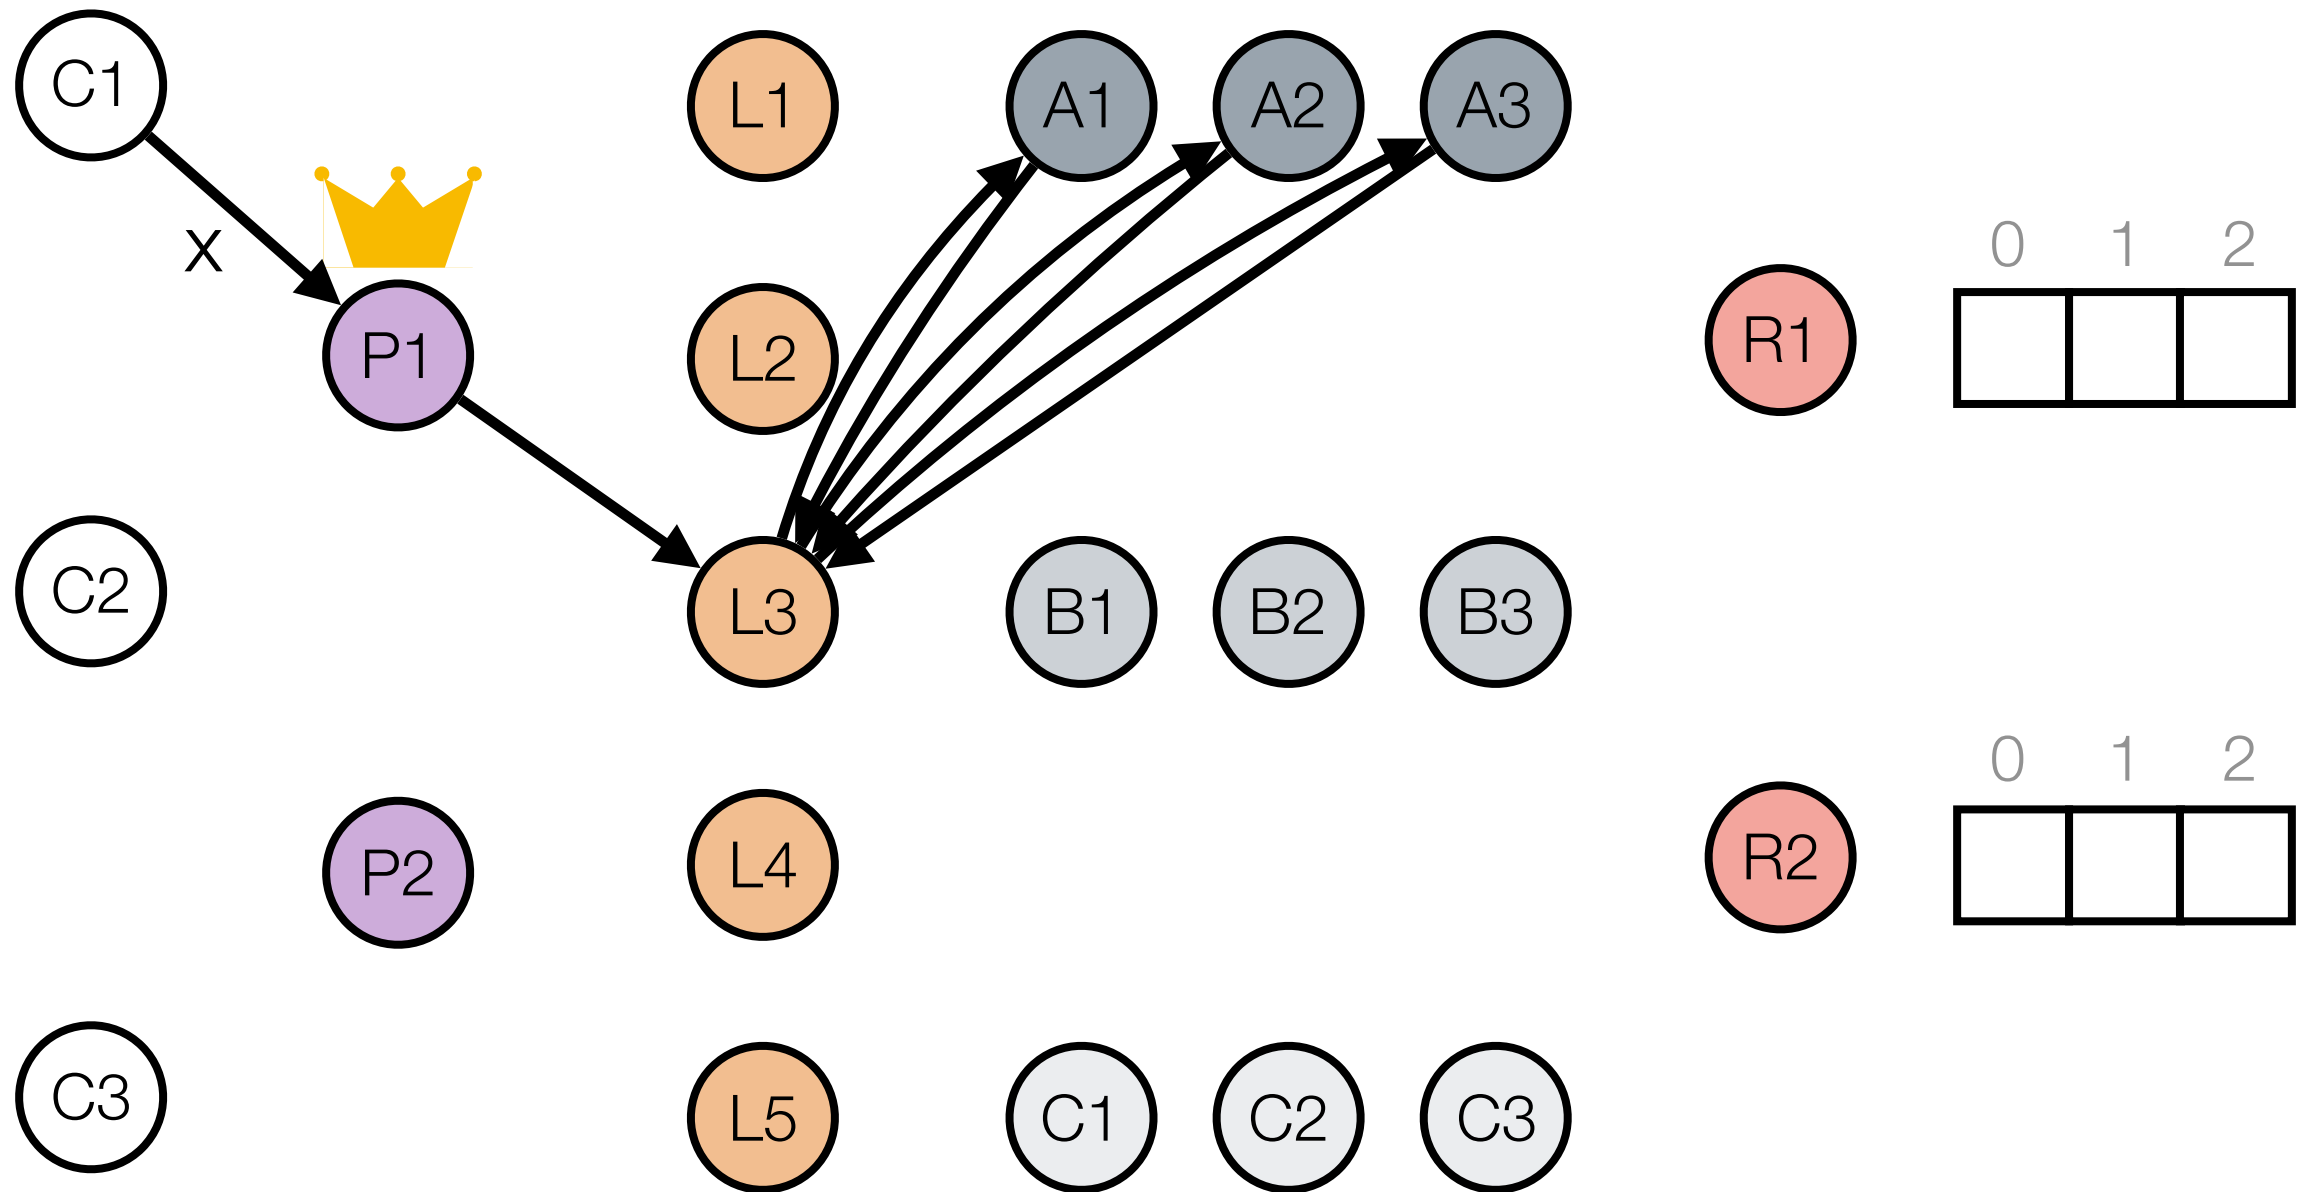
clients

$f+1$
proposers

$\geq f+1$
proxy
leaders

≥ 1 groups of
 $2f+1$
acceptors

$f+1$
replicas



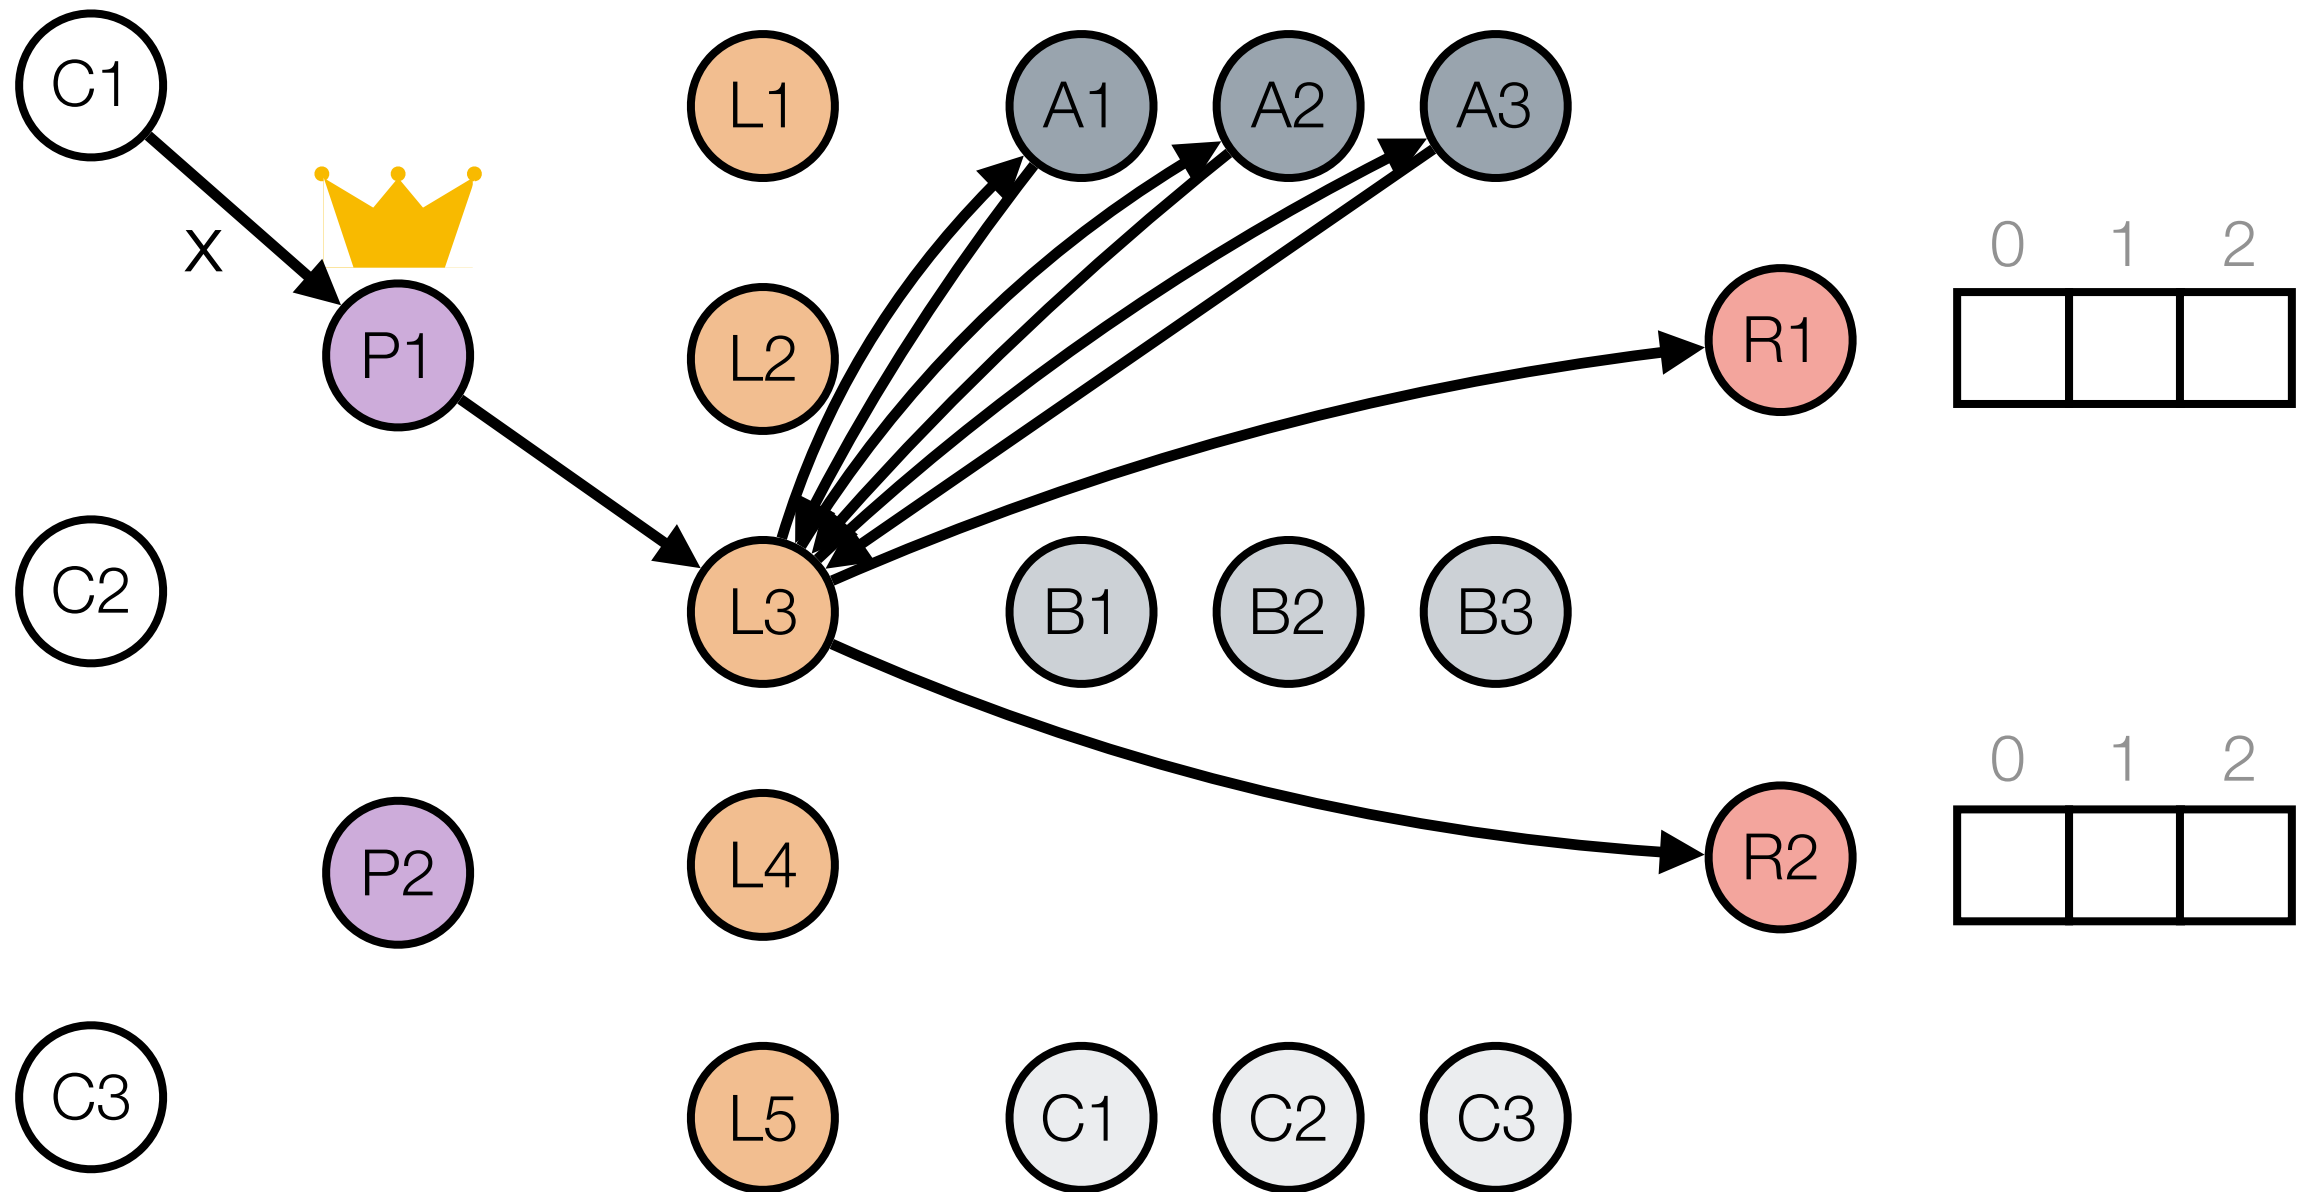
clients

$f+1$
proposers

$\geq f+1$
proxy
leaders

≥ 1 groups of
 $2f+1$
acceptors

$f+1$
replicas



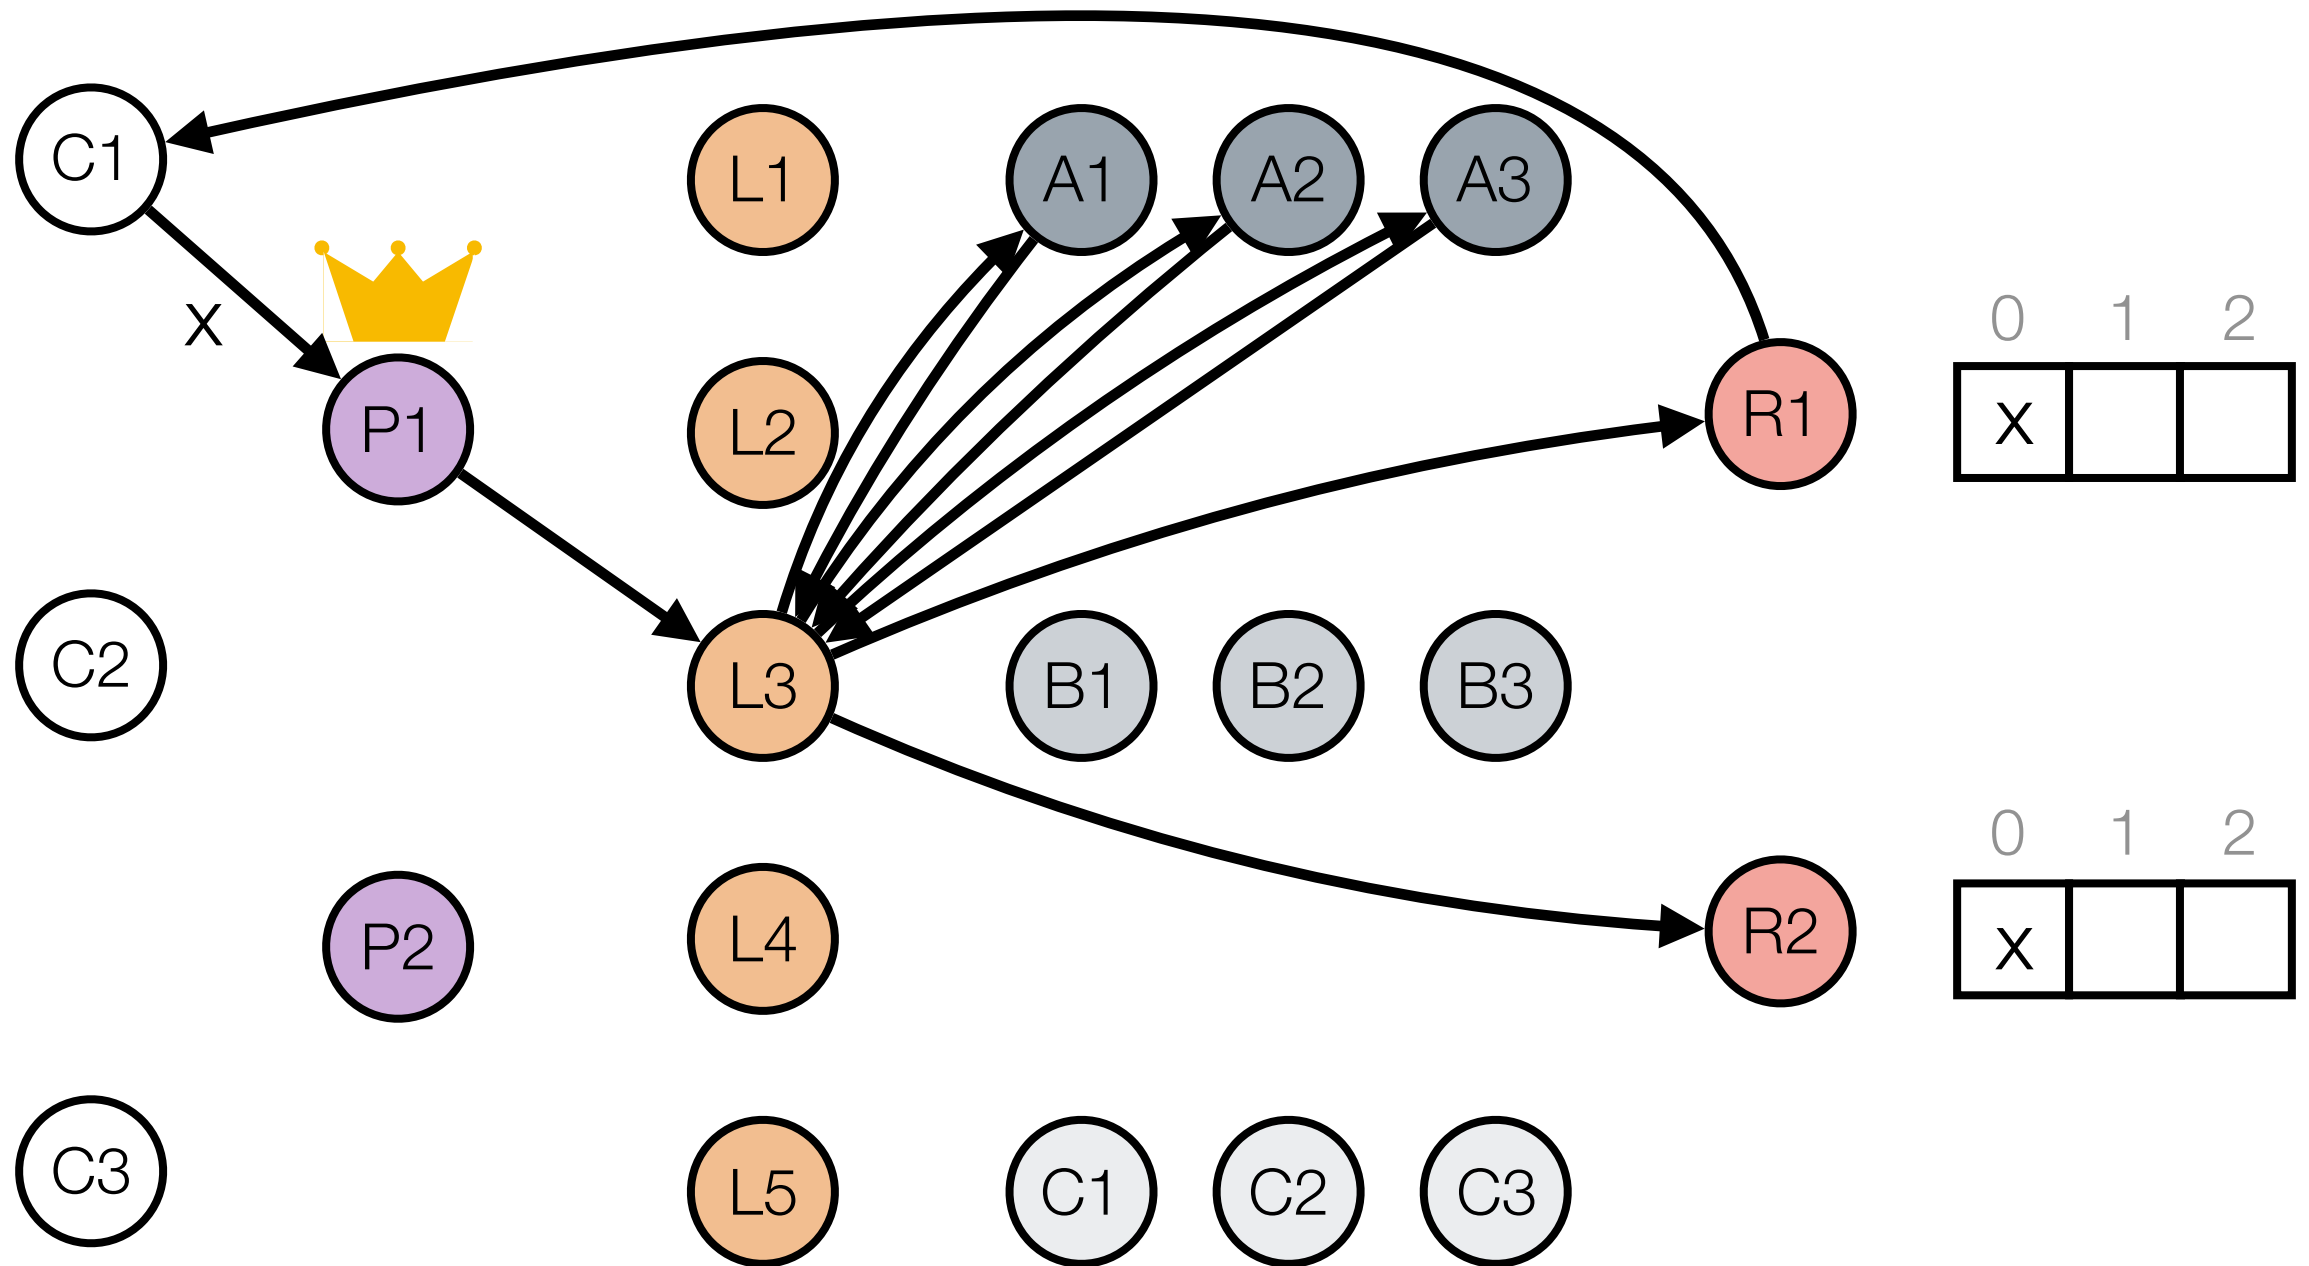
clients

$f+1$
proposers

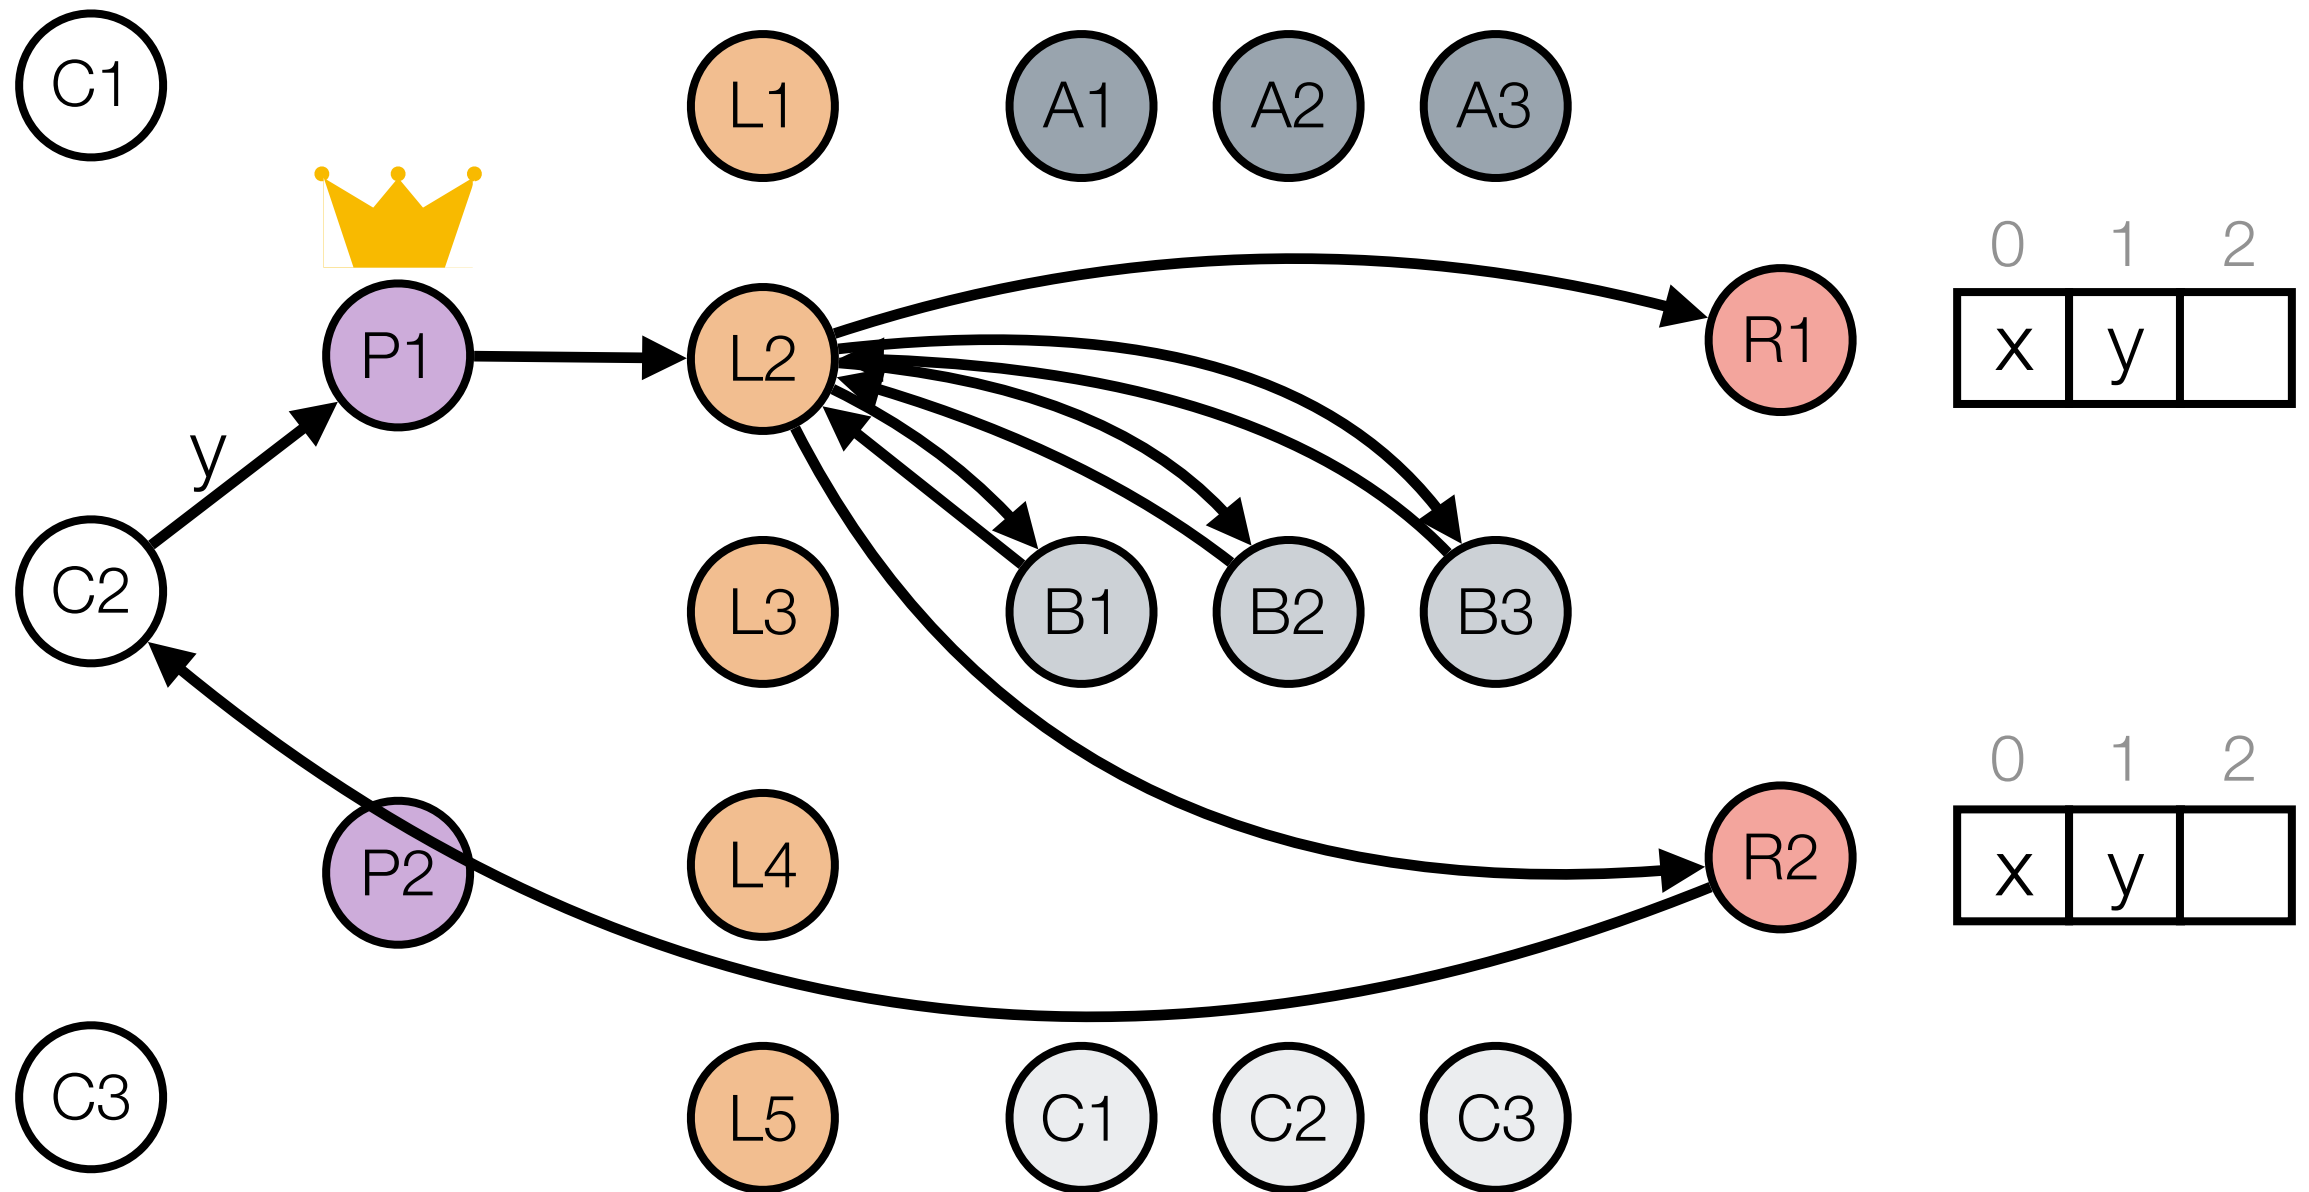
$\geq f+1$
proxy
leaders

≥ 1 groups of
 $2f+1$
acceptors

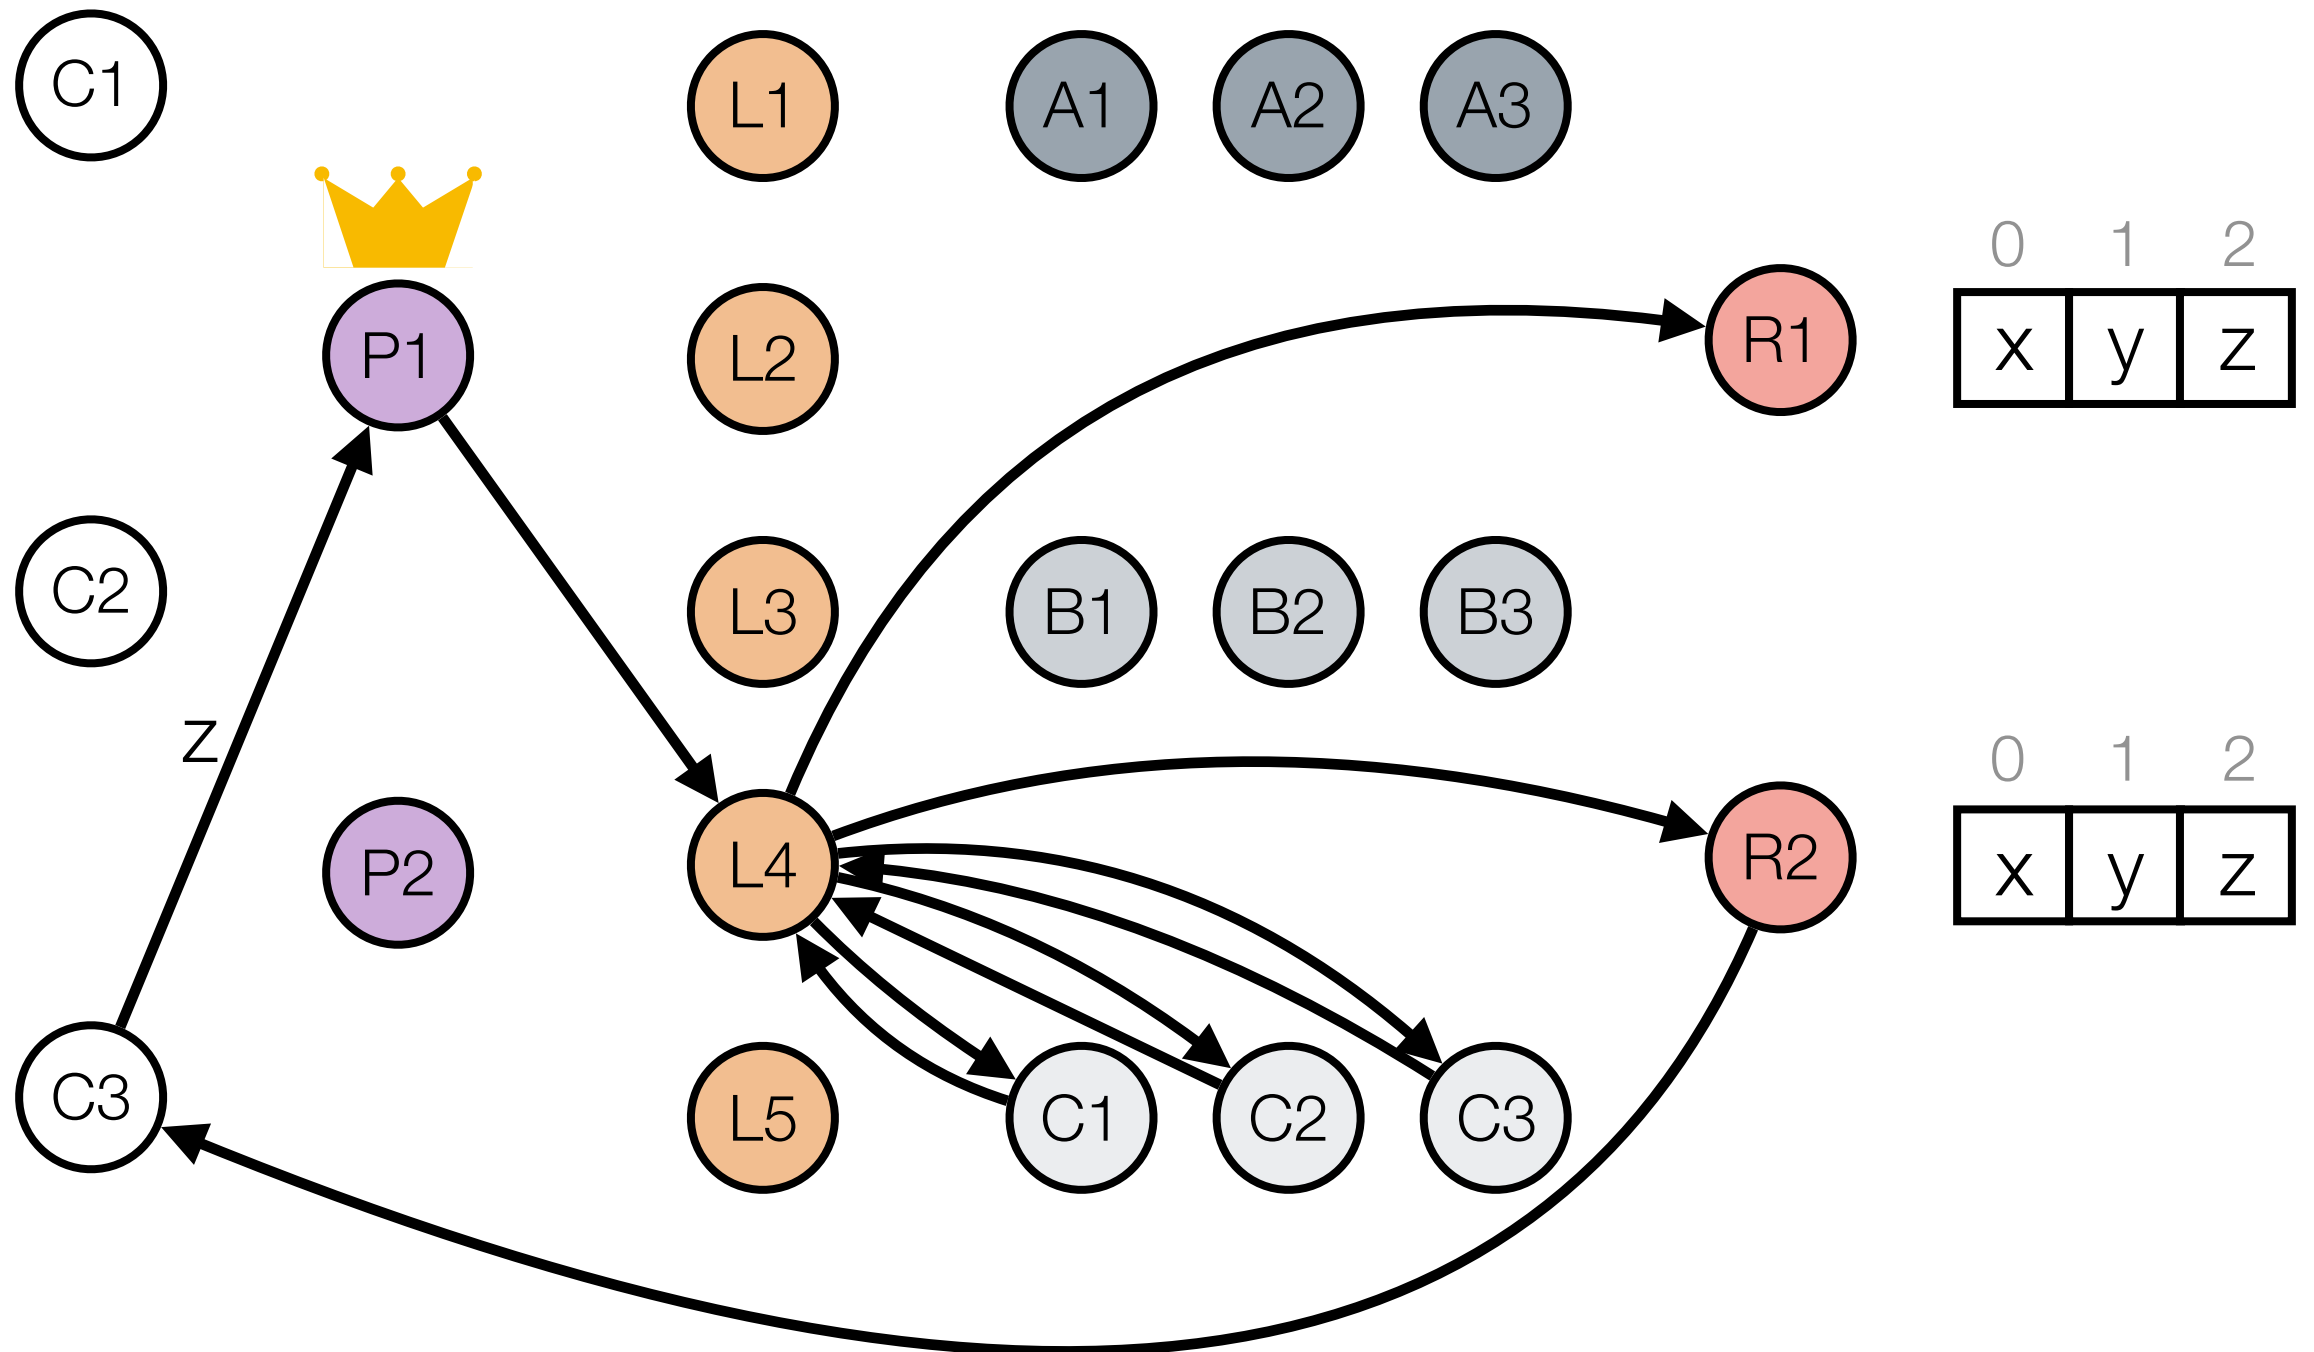
$f+1$
replicas



clients $f+1$ proposers $\geq f+1$ proxy leaders ≥ 1 groups of $2f+1$ acceptors $f+1$ replicas



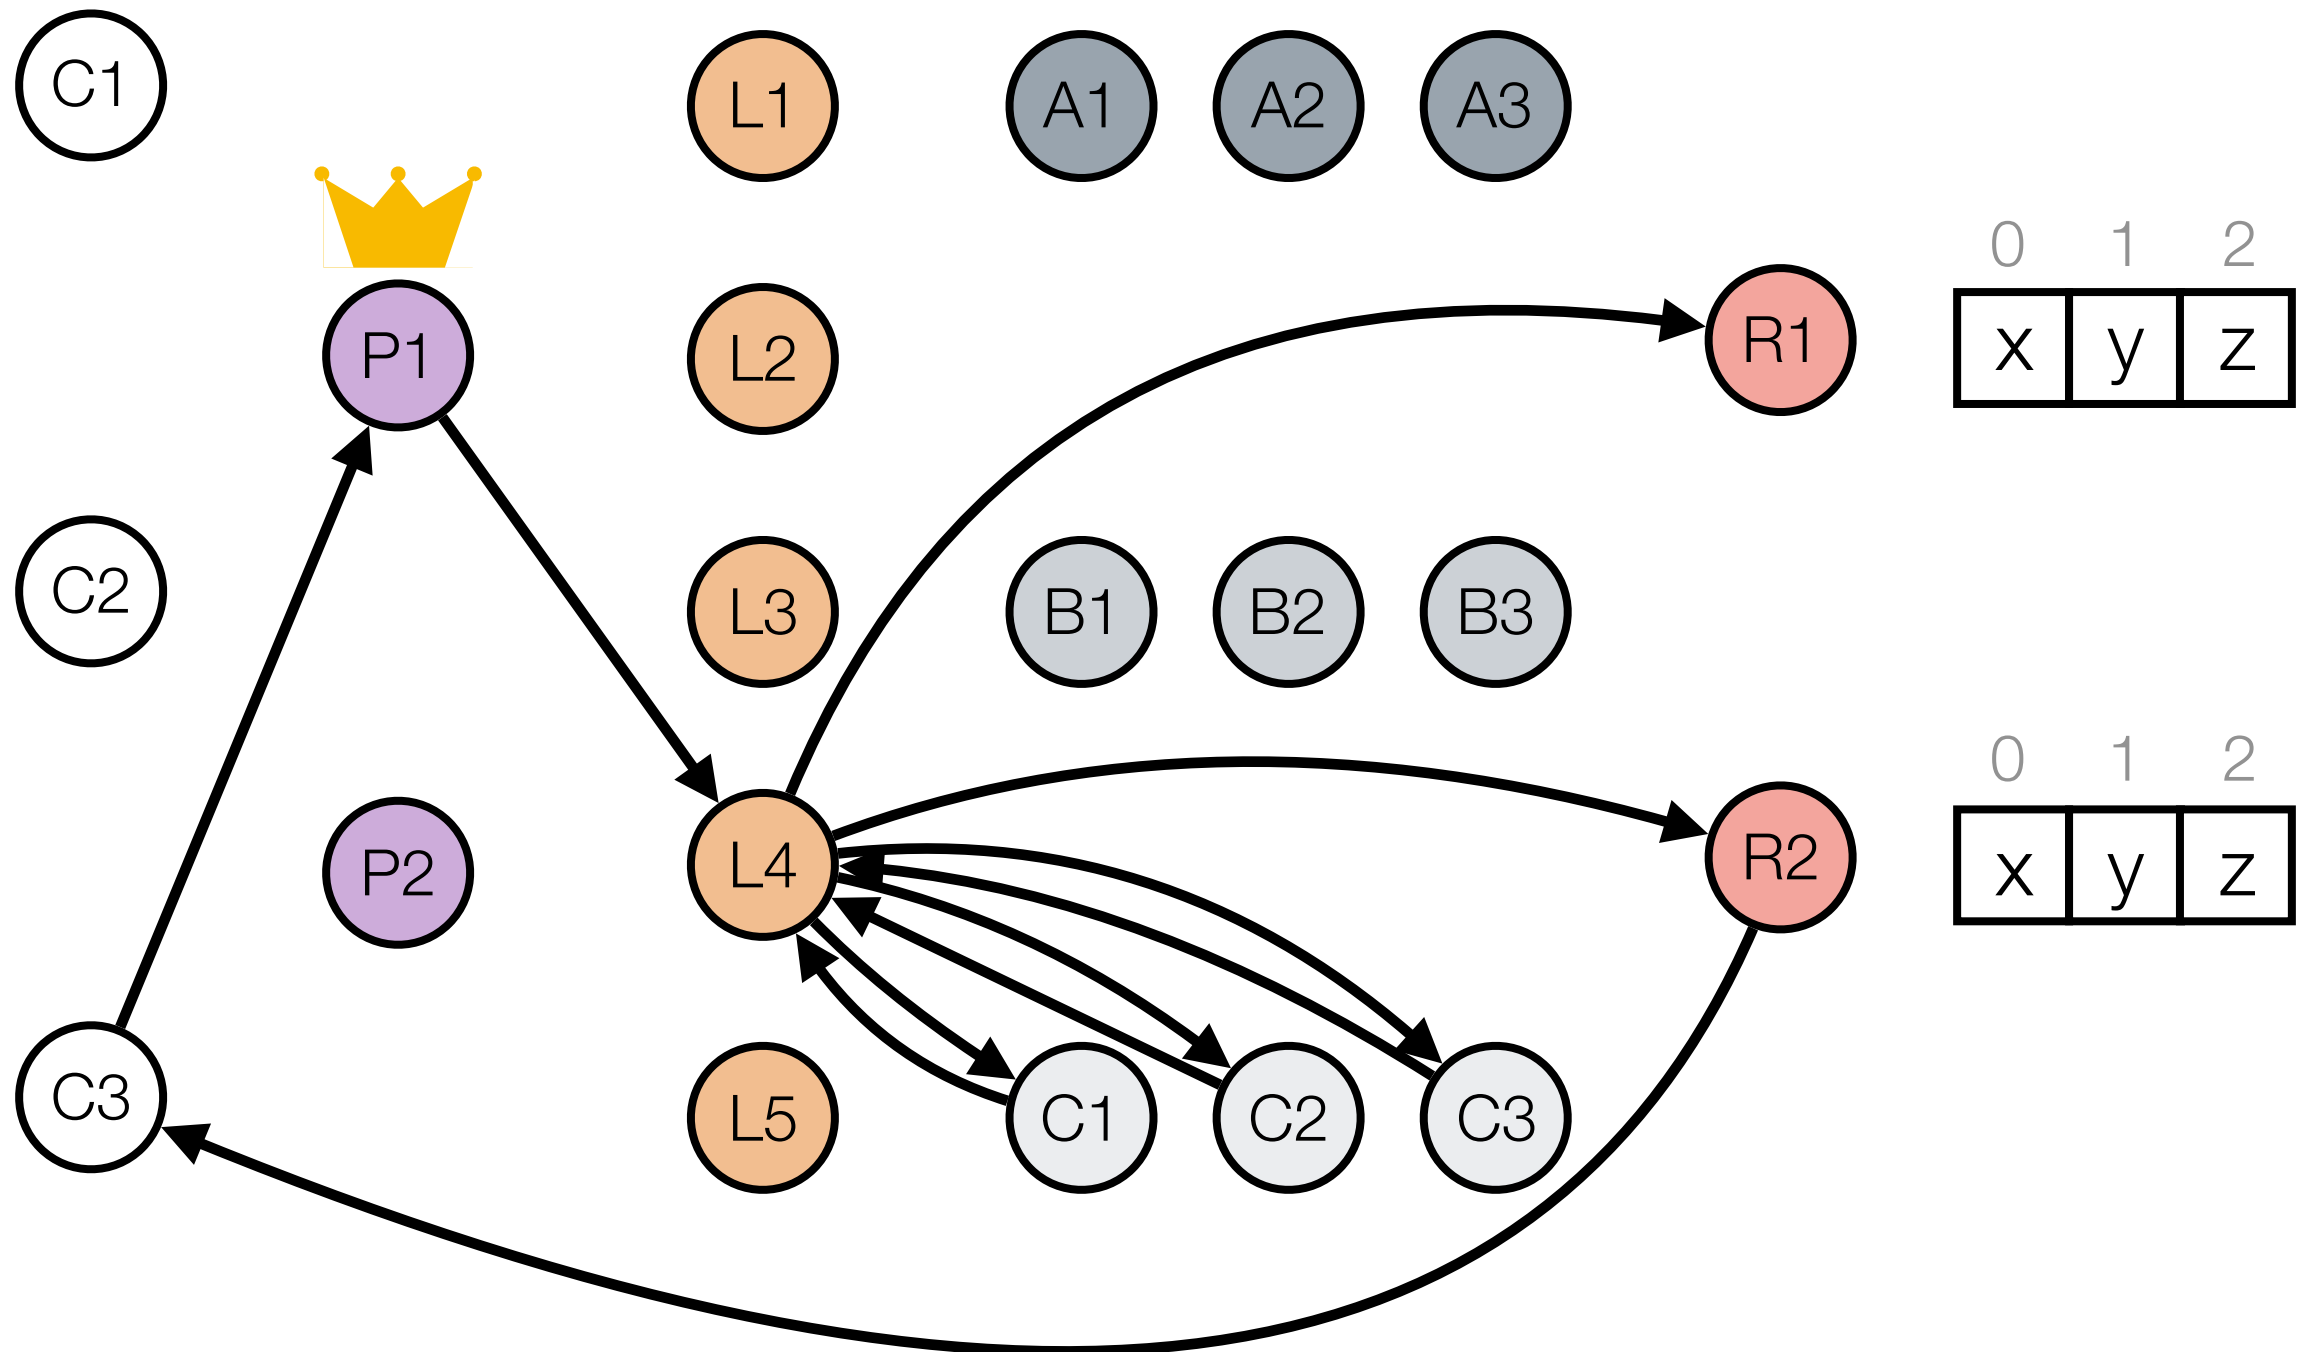
clients $f+1$ proposers $\geq f+1$ proxy leaders ≥ 1 groups of $2f+1$ acceptors $f+1$ replicas




“Using more than $2f+1$
[acceptors] for f failures is
possible but illogical
because it requires a larger
quorum size with **no**
additional benefit”

-TAPIR (SOSP 2015)


clients $f+1$ proposers $\geq f+1$ proxy leaders ≥ 1 groups of $2f+1$ acceptors $f+1$ replicas



clients $f+1$ $f+1$ $\geq f+1$ ≥ 1 groups of $f+1$ $f+1$
 batchers proposers proxy $2f+1$ replicas unbatchers
 leaders acceptors

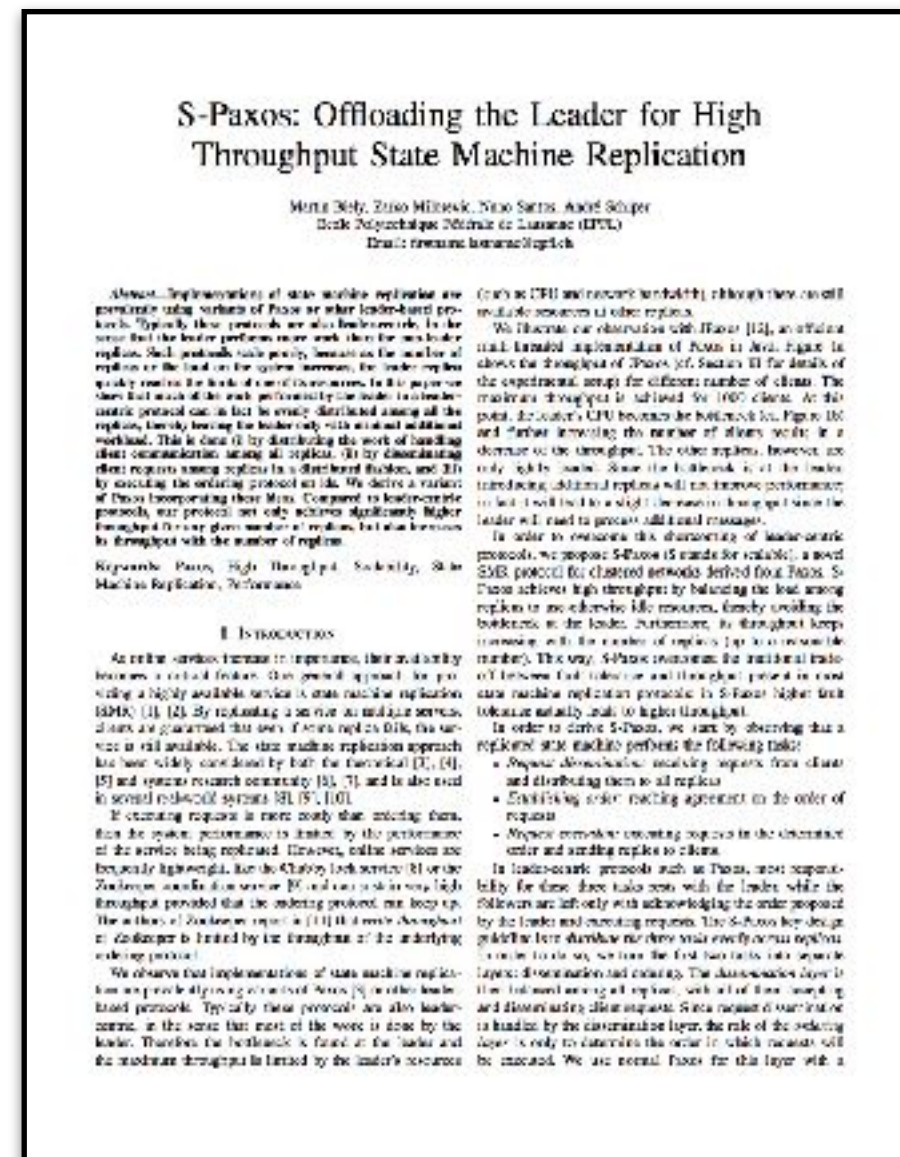
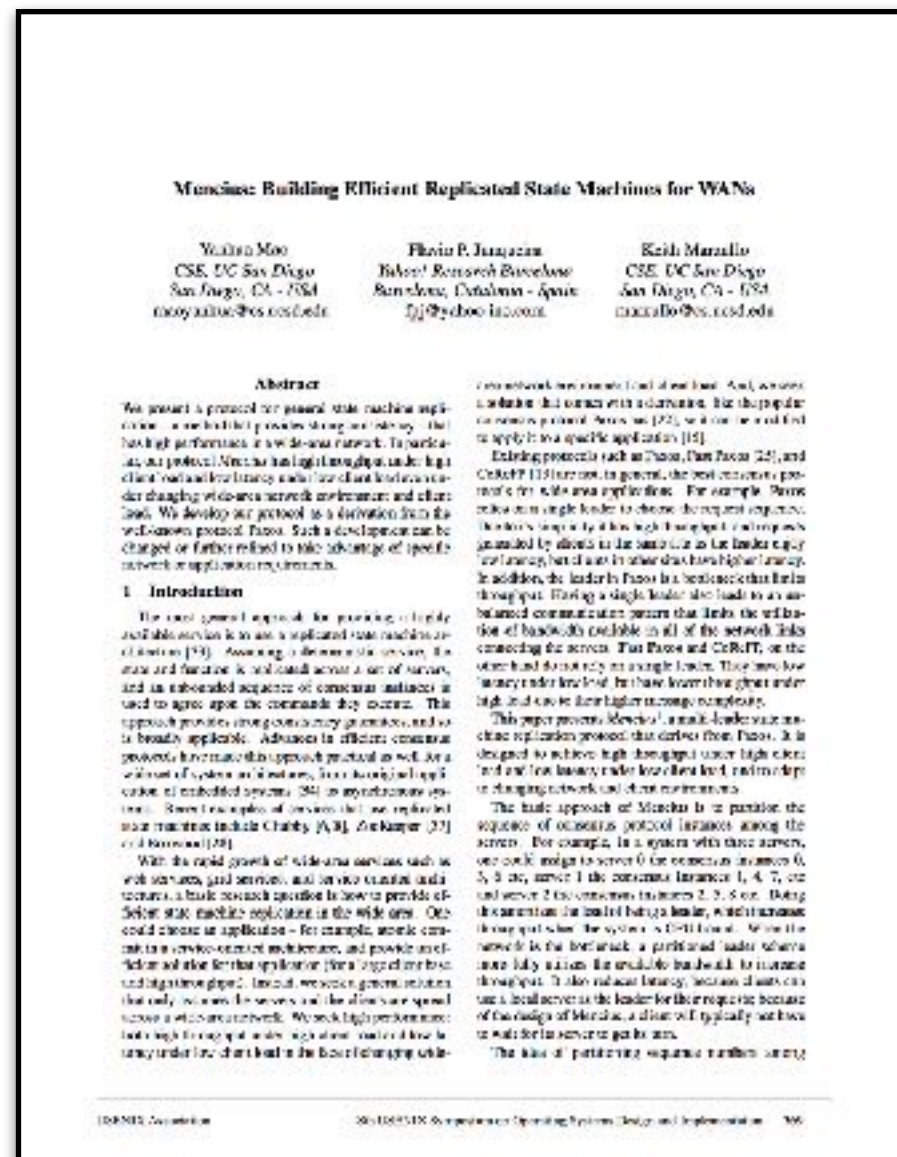



Mencius: round robin partition the log between multiple leaders



Mencius: round robin partition the log between multiple leaders

S-Paxos: decouple data flow from control flow





Some protocols
become the
bottlenecks they try
to avoid!

mwhittaker.github.io

mwhittaker.github.io

mwhittaker.github.io/frankenpaxos

mwhittaker.github.io

mwhittaker.github.io/frankenpaxos

github.com/mwhittaker/frankenpaxos